09/05/2023

Accelerating iterative relational algebra operations

oneAPI Hackathon: CUDA to SYCL Migration

UNIVERSITY OF ILLINOIS CHICAGO

Team Members

Ahmedur Rahman Shovon

Ph.D. Student (CS)
University of Illinois Chicago
Email: ashov@uic.edu

Thomas Gilray

Assistant Professor (CS)
University of Alabama at Birmingham
Email: gilray@uab.edu

Sidharth Kumar

Assistant Professor (CS)
University of Illinois Chicago
Email: sidharth@uic.edu

Transitive Closure Computation using Iterative Relational Algebra

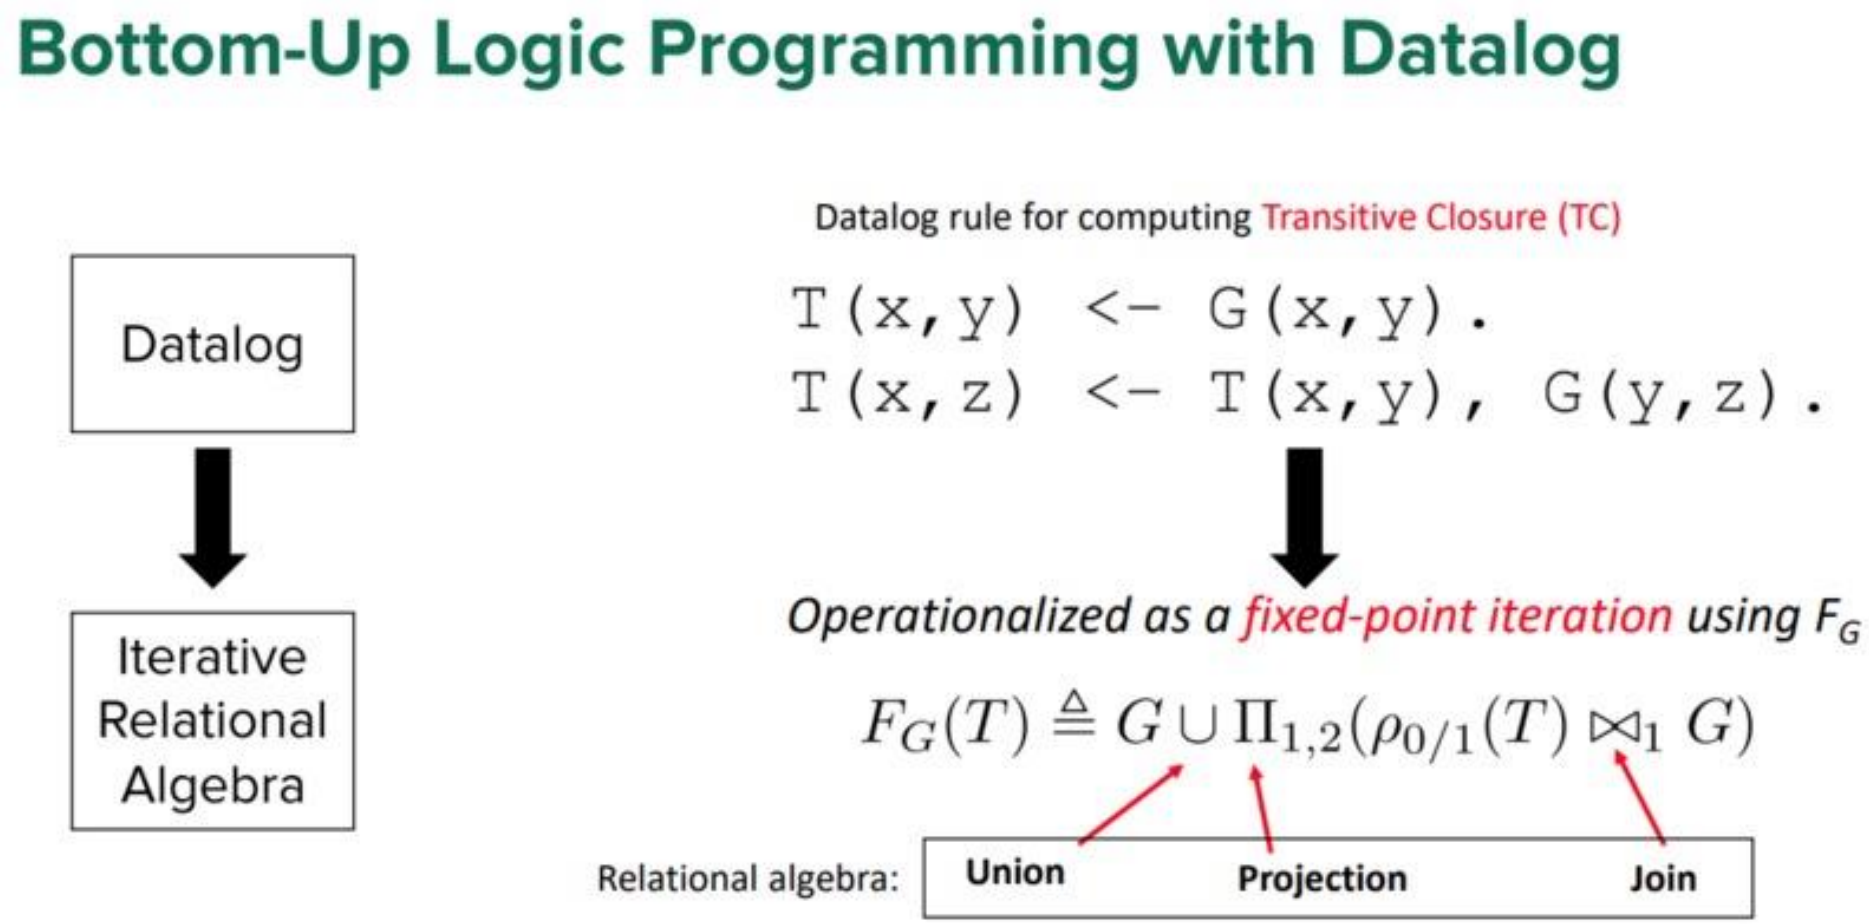

Inspiration

- Iterative relational algebra (RA kernels in a fixed-point loop) enables bottom-up logic programming languages such as Datalog which can be implemented using relational algebra primitives (e.g., projections, reorderings, and joins)
- While much has explored standalone RA operations on the GPU, relatively less work focuses on iterative RA, which exposes new challenges (e.g., deduplication and memory management)

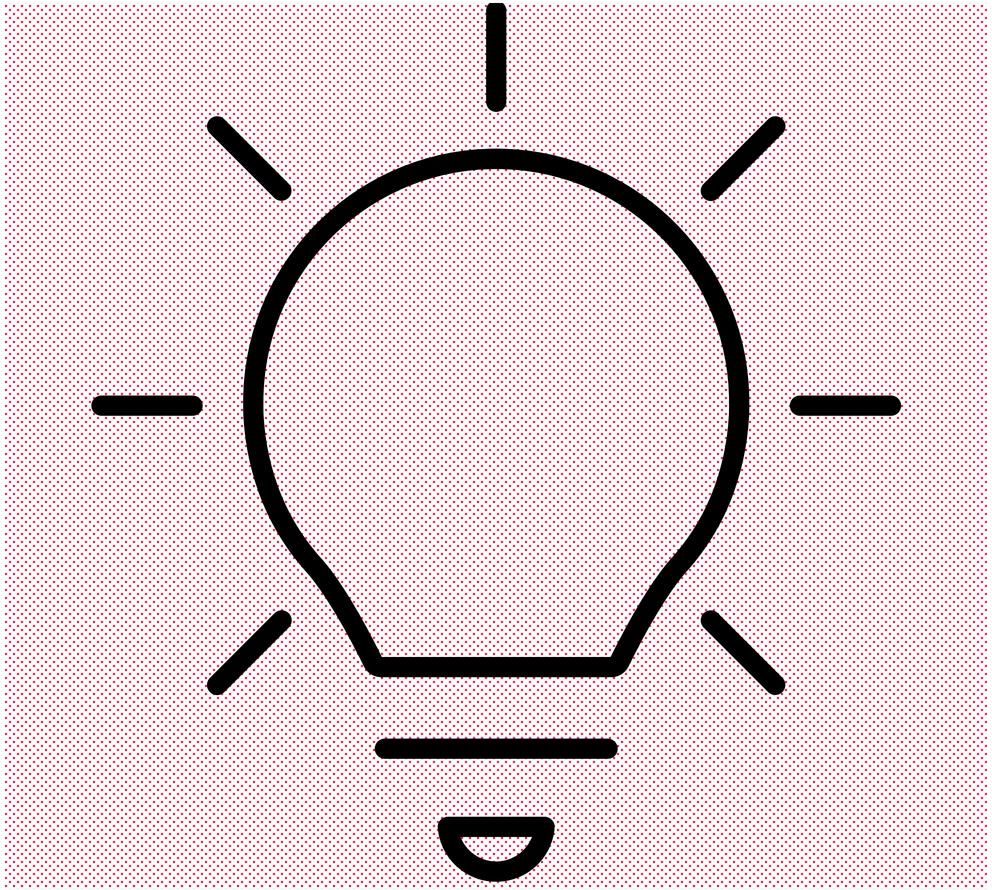

What it does

- Developed a GPU-based hash-join implementation, leveraging
  - a novel open-addressing-based hash table implementation
  - operator fusing to optimize memory access
  - two variant implementations of deduplication
- Implemented transitive closure using our hash-join-based CUDA library and compared its performance against cuDF (GPU-based) and Souffle (CPU-based)

Environment and Datasets

- Benchmarked our experiments on the ThetaGPU supercomputer of Argonne National lab using a single nVidia A100 GPU
- CUDA kernel size: 3456 X 512
- CUDA version: 11.4
- Souffle version: 2.3 with 128 threads
- Datasets: Stanford large network dataset collection, SuiteSparse matrix collection, and road network real datasets collection

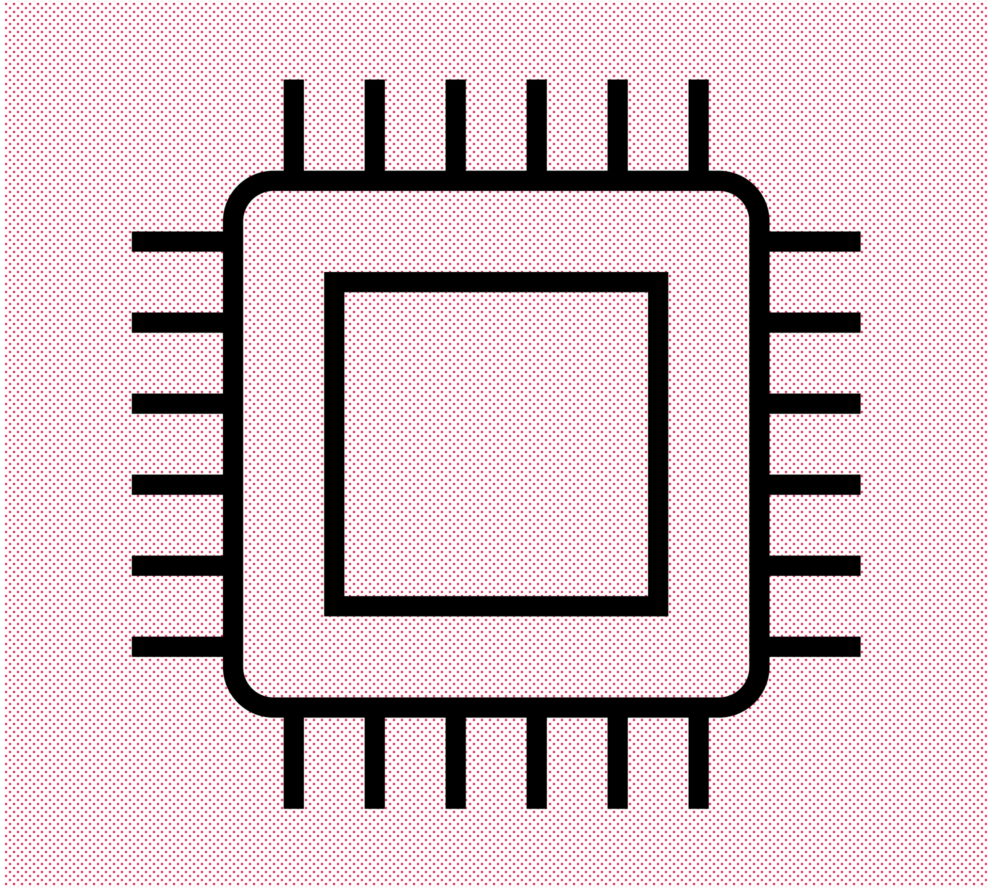

Challenges we ran into

- Unlike C++, CUDA lacks efficient data structures, limiting our implementation's capabilities
- The available VRAM of a single GPU imposes constraints on our implementation's scalability
- Debugging kernel code posed significant challenges, turning it into a nightmarish experience

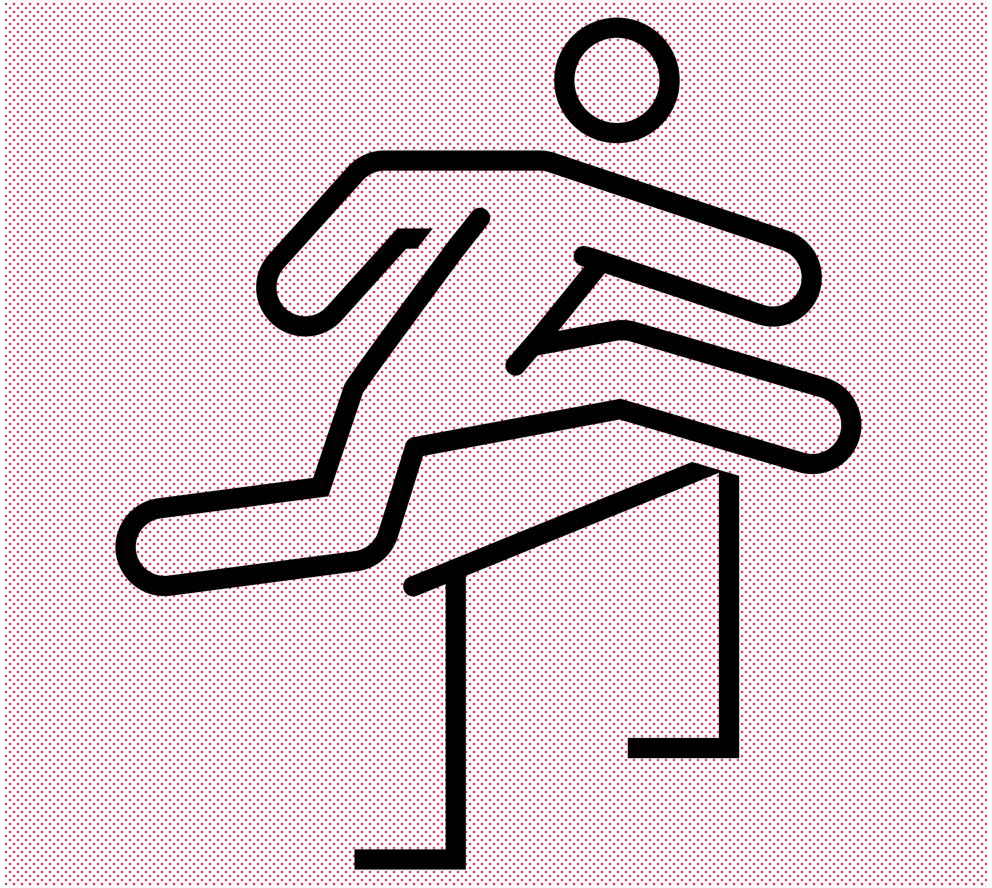

What we learned

- Efficient memory management is crucial for successful CUDA implementations
- Handling the indeterministic result size per iteration in Iterative RA operations requires careful consideration
- While low-level GPU code allows optimization opportunities, it demands considerable time and effort compared to using off-the-shelf libraries like cuDF

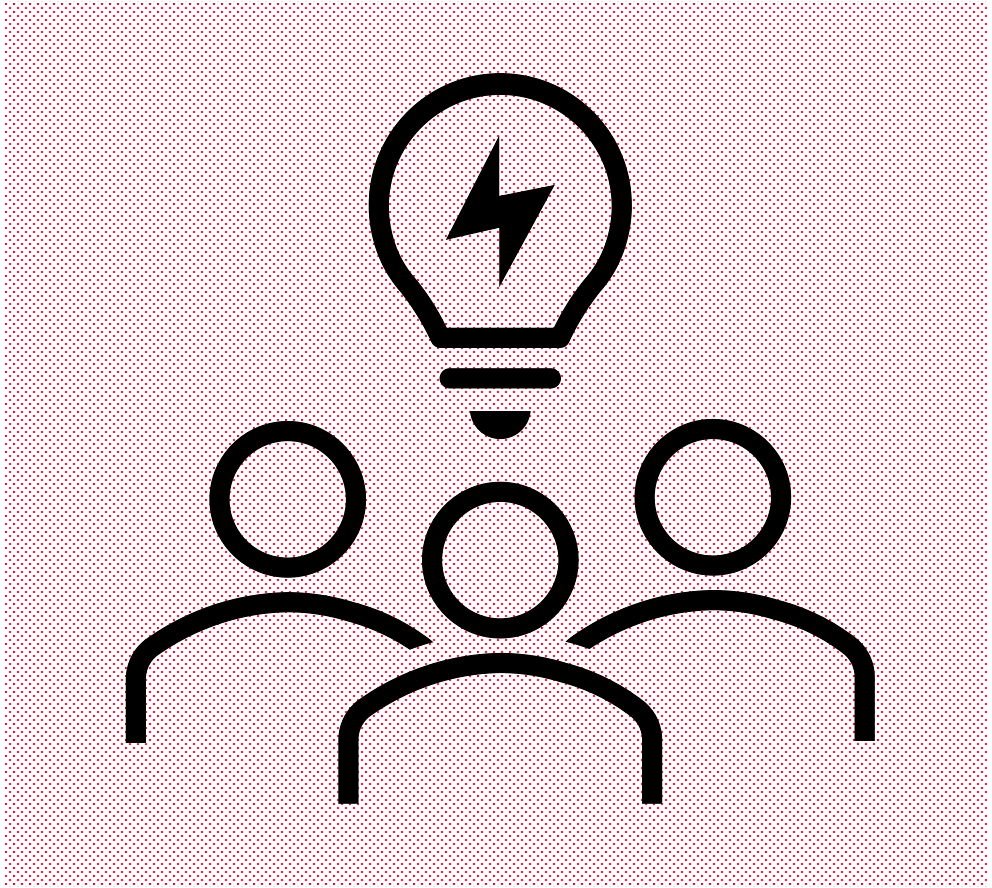

Accomplishments

- Our hash-join-based transitive closure computation shows favorable results against both cuDF and Souffle, with gains up to 10.8x against cuDF and 3.9x against Souffle

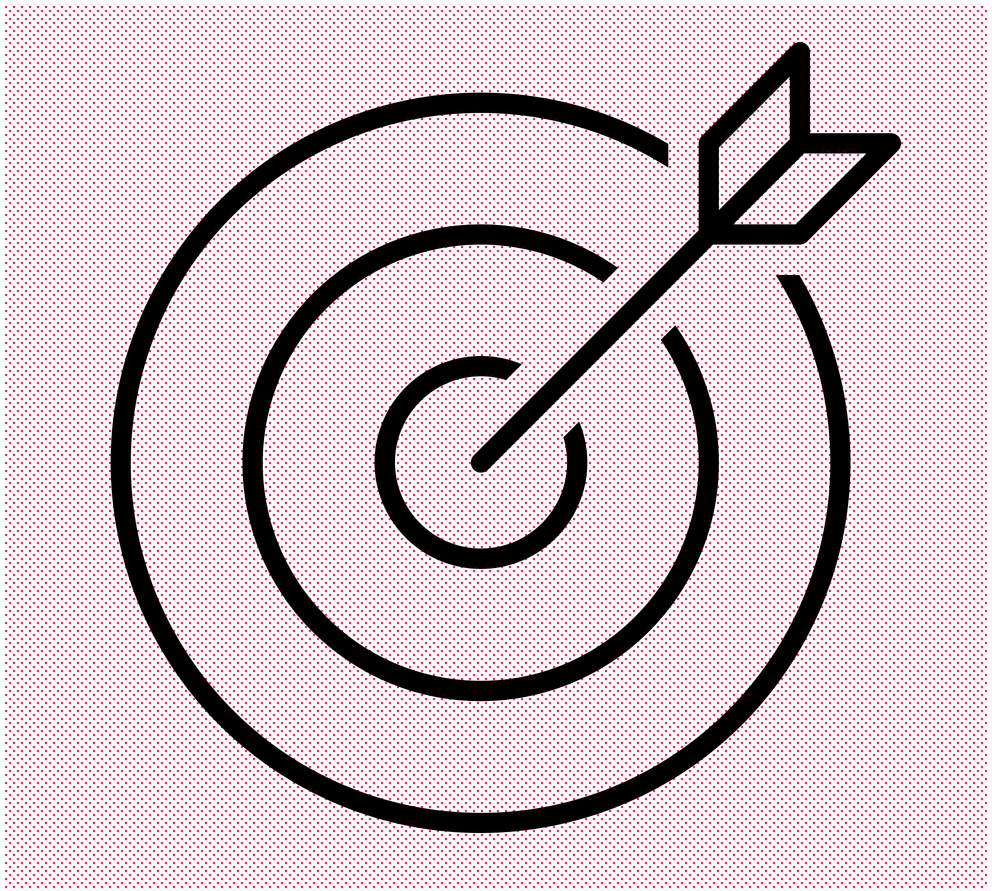

Publications

- Shovon, A. R., Gilray, T., Micinski, K., & Kumar, S. (2023). Towards Iterative Relational Algebra on the {GPU}. In 2023 USENIX Annual Technical Conference (USENIX ATC 23) (pp. 1009-1016).
- Shovon, A. R., Dyken, L. R., Green, O., Gilray, T., & Kumar, S. (2022, November). Accelerating Datalog applications with cuDF. In 2022 IEEE/ACM Workshop on Irregular Applications: Architectures and Algorithms (IA3) (pp. 41-45). IEEE.

Project Repository

- Transitive closure computation using CUDA: https://github.com/harp-lab/usenixatc23
- Transitive closure computation using SYCL: https://github.com/arsho/tc

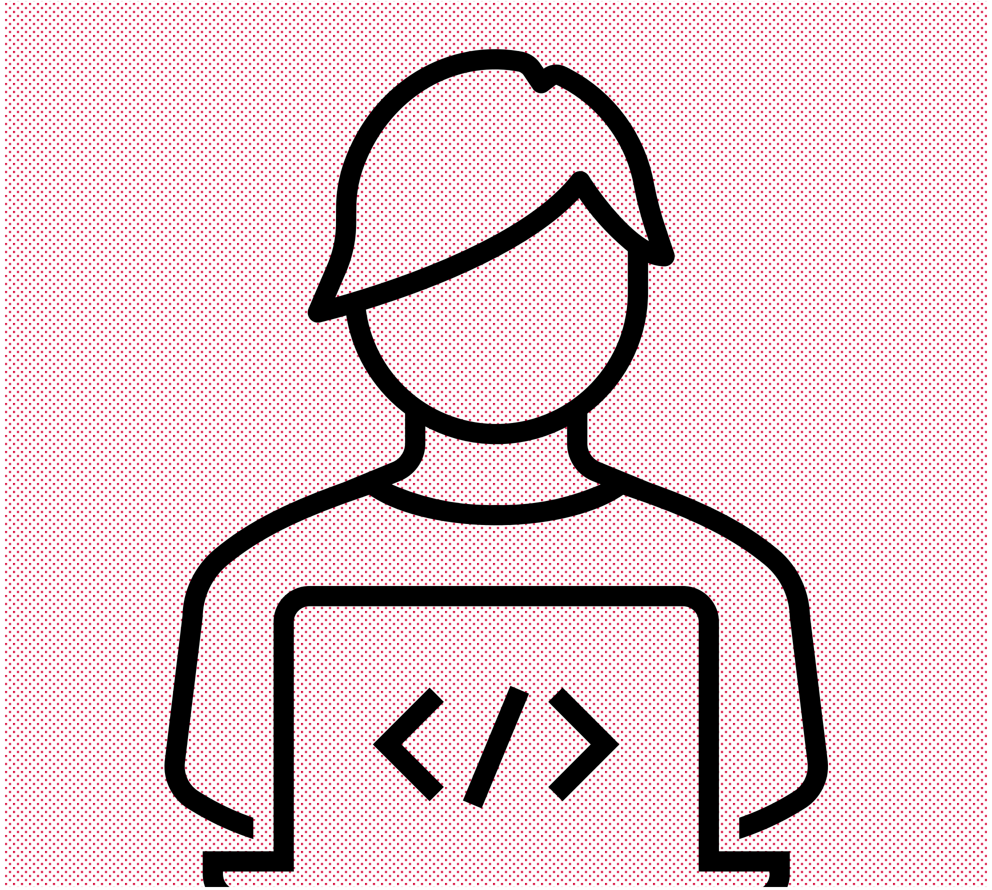

Porting TC computation CUDA implementation

Clean	Clean CUDA Project
Install	Install SYCLomatic
Convert	Convert CUDA code to SYCL
Check	Check the SYCL code and modify if necessary
Execute	Run in Intel Dev Cloud

Clean CUDA Project

- Our CUDA code has multiple files and a Makefile that has auxilary commands
- To port the CUDA project to SYCL first we cleaned the CUDA project
- We made one file that has CUDA code, simplified the Makefile, and kept one test dataset
- The folder *sycl_implementation* has the single cuda file

```
.
├── cuda_implementation
│   ├── data_7035.txt
│   ├── error_handler.cu
│   ├── kernels.cu
│   ├── Makefile
│   ├── tc_cuda.cu
│   └── utils.cu
├── LICENSE
├── README.md
└── sycl_implementation
    ├── data_7035.txt
    ├── Makefile
    ├── tc_cuda.cu
    └── tc_sycl
```

Install SYCLomatic

- Open a terminal and download SYCLomatic release:
  cd ~/
  mkdir syclomatic
  cd syclomatic
  wget https://github.com/oneapi-src/SYCLomatic/releases/download/20230725/linux_release.tgz
  tar -xvf linux_release.tgz
- Add the bin path to .zshrc:
  export PATH="~/syclomatic/bin:$PATH"
- Check *c2s* version:
  c2s --version

SYCLomatic release: https://github.com/oneapi-src/SYCLomatic/tags

Convert CUDA code to SYCL

- Convert the CUDA code to SYCL and create a directory to store the SYCL code:
  intercept-build make
  c2s -p compile_commands.json --out-root tc_sycl
- Copy sample dataset to the SYCL code directory and create a compressed file:
  cp data_7035.txt tc_sycl
  tar -cvf tc_sycl.tgz tc_sycl

SYCLomatic doc: https://github.com/oneapi-src/SYCLomatic

Check SYCL Converted Code

- We had error in converted SYCL code using SYCLomatic *20230725* release. In converted SYCL code, we needed to replace *std::reducet* to *std::reduce*
- This error did not appear when we used SYCLomatic *20230830* release

SYCLomatic doc: https://github.com/oneapi-src/SYCLomatic

Execute in Intel Dev Cloud

- Upload the SYCL code folder to Intel Dev Cloud:
  scp tc_sycl.tgz idc:~/
- Connect to Intel Dev Cloud, start an interactive session, load the modules:
  ssh idc
  srun --pty bash
  source /opt/intel/oneapi/setvars.sh
- Execute the SYCL code:
  tar -xvf tc_sycl.tgz
  cd tc_sycl
  *icpx -fsycl *.cpp*

Intel Dev Cloud instances: https://scheduler.cloud.intel.com/#/systems

Feedback

- Examples on dynamic data structure implementations using SYCL would be helpful
- Automated generation of additional comments on ported code can explain the converted code

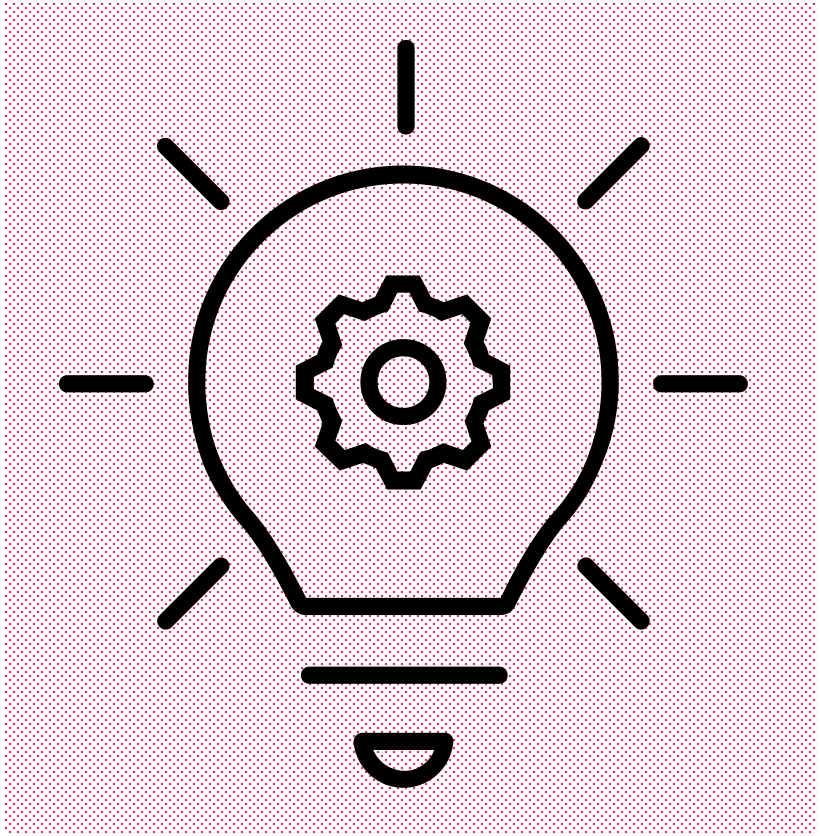

Status

- We ported the CUDA transitive closure computation code to SYCL but it raised several errors on ported code; specially on Thrust CUDA code
- After closely investigate the errors, we found most of the errors can be mitigated using standard SYCL functions, no need to use Thrust APIs in those scenarios
- As SYCL supports many standard data structures, we decided to implement SYCL implementation from scratch

Future Direction

- Implement transitive closure computation using SYCL from scratch
- Compare the TC computation with CUDA, cuDF, and Souffle on single GPU
- Extend the computation to use multiGPU environment on intel GPU targeting the Aurora supercomputer

Thank You