



01/26/2024

Declarative Analytics on Heterogeneous Exascale Systems

Ahmedur Rahman Shovon

WCP Committee:

Dr. Michael Papka (Chair)

Dr. Sidharth Kumar (Advisor)

Dr. Pedram Rooshenas



Declarative Analytics on Heterogeneous Exascale Systems

Users expresses **what** to achieve with the data rather than **how** to accomplish it

User

UserID	UserName	UserEmail	Country
101	Alice	alice@example.com	USA
102	Bob	bob@example.com	USA
103	Eve	eve@example.com	Australia

WHAT

SELECT UserID FROM User WHERE Country = 'USA';

HOW

Advanced approach: Logic programming (Datalog)

- Makrynioti, N., & Vassalos, V. (2019). Declarative data analytics: A survey. *IEEE Transactions on Knowledge and Data Engineering*, 33(6), 2392-2411.
- Salesforce. (2024). Click, Not Code: The Benefits of Declarative Programming vs. Imperative Programming retrieved from <https://www.salesforce.com/products/platform/best-practices/declarative-programming-vs-imperative-programming/> on 01/24/2026

Declarative Analytics on Heterogeneous Exascale Systems

The screenshot shows a DevPro JOURNAL article. At the top, there are banners for Axium (simplifying development, estate monitoring and security) and ingenico. The main headline is "Low-Code/No-Code: Why Declarative Approaches are Winning the Future of AI". Below the headline is a subtext: "Declarative machine learning is going to be the preferred way most organizations operationalize task-specific AI to solve business problems." The author is Devret Rishi, and the date is November 28, 2023. The article features a large graphic titled "LOW CODE PLATFORM" with various icons like gears, a brain, and a dollar sign. A sidebar on the right allows users to join a newsletter.

The screenshot shows a Databricks blog post. The title is "Deloitte's Guide to Declarative Data Pipelines With Delta Live Tables". It features a small image of two people, Mani Kandasamy and Vijay Balasubramanian, and the date October 19, 2022. The post is written in collaboration with Deloitte. On the right side, there are sections for "Value calculator" and "Business Value Report", along with a "Try for free" button.

Enables drag-and-drop like analytics pipeline to make informed business decisions without the need for complex coding



A Declarative Approach To Data Pipelines

Mutinex offers a marketing analytics & econometrics platform that helps marketers make better investment decisions faster.

Over the last year, Altis collaborated with Mutinex to establish an automated data ingestion and validation pipeline.

The screenshot shows two blog post thumbnails from the Mutinex website. The first is titled "WHY ORGANISATIONS NEED TO IMPLEMENT ASSET INTELLIGENCE" and the second is "WHAT IS THE LATEST ON MACHINE LEARNING IN THE MARKET". Both posts have "Blog Posts" and "Read more" buttons.

The screenshot shows the Einstein Platform website. The header features the Einstein logo and navigation links for Overview, Solutions, Hyperforce, Features, Automation, Related Products, Pricing, App Gallery, and Resources. The main content area has a section titled "Click, Not Code: The Benefits of Declarative Programming vs. Imperative Programming" with a lightning bolt icon. Below it is a diagram showing a cloud icon connected to a database and a checkmark. A sidebar on the right shows a "PRODUCT" section with a "See how to build apps on the Lightning Platform" link and a "LEARN MORE" button. There is also a "FREE TRIAL" button and a "Let's Chat" button.

- 3
- Makrynioti, N., & Vassalos, V. (2019). Declarative data analytics: A survey. *IEEE Transactions on Knowledge and Data Engineering*, 33(6), 2392-2411.
 - Salesforce. (2024). Click, Not Code: The Benefits of Declarative Programming vs. Imperative Programming retrieved from <https://www.salesforce.com/products/platform/best-practices/declarative-programming-vs-imperative-programming/> on 01/24/2026

Table of Contents

Exascale Computing
Datalog
Iterative Relational Algebra
Parallel Relational Algebra
Our Approaches
Topological Data Analysis for High-Dimensional Data
Future Research Direction
Conclusion

What is Exascale?

Ability of a computer system to perform
1 Exaflop or more

1 Exaflop = 1 quintillion flop (10^{18})
or 10^9 Gigaflop



Capacity?

Perform as many calculations as about
50 million laptops!

Speed?

A trillion times faster than a PC!

Rank	Site	Computer	Cores	HPL (Eflop/s)
1	DOE/SC/ORNL	Frontier	8,730,112	1.194
2	DOE/SC/ANL	Aurora (Half)	4,742,808	0.585
3	Microsoft Azure	Eagle	1,123,200	0.561
4	RIKEN	Fugaku	7,630,848	0.442
5	EuroHPC/CSC	LUMI	2,752,704	0.379

TOP500 HPL (High-Performance Linpack (HPL) benchmark of high-performance computing (HPC) from November 2023

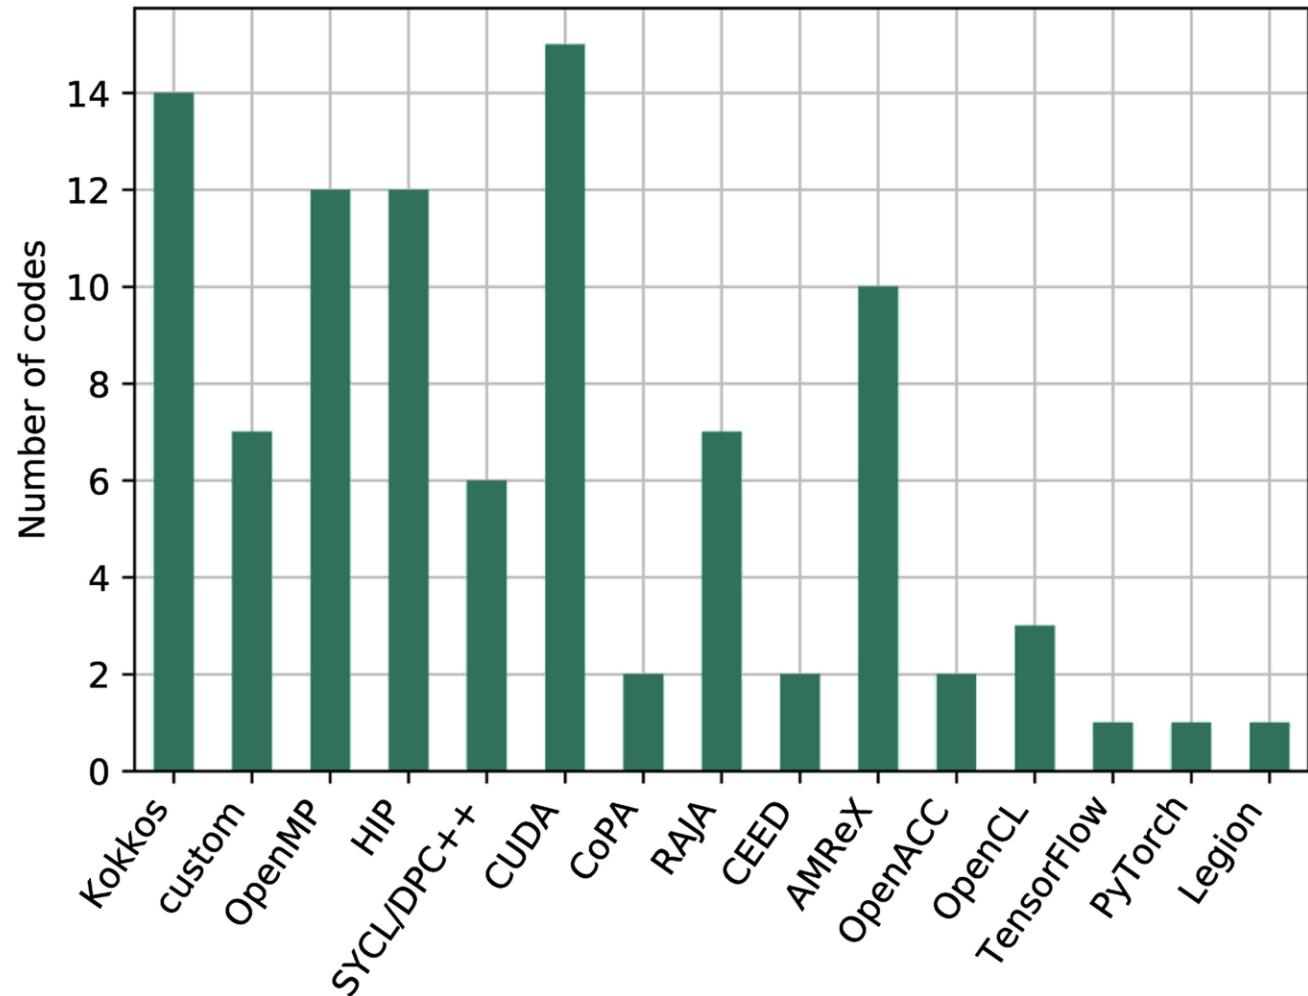
- Merritt, R. (2022, September 16). What is an exaflop? NVIDIA Blog. Retrieved April 9, 2023, from <https://blogs.nvidia.com/blog/2022/07/26/what-is-an-exaflop/>

- Strohmaier, E. et al. (2023, November). TOP500 Lists. HPL. Retrieved January 20, 2024, from <https://www.top500.org/lists/top500/2023/11/>

- Webstudios, G. (2021, August 31). Why countries are competing to build supercomputers? Griffon Webstudios. Retrieved April 9, 2023, from <https://griffonwebstudios.com/build-super-computers/>

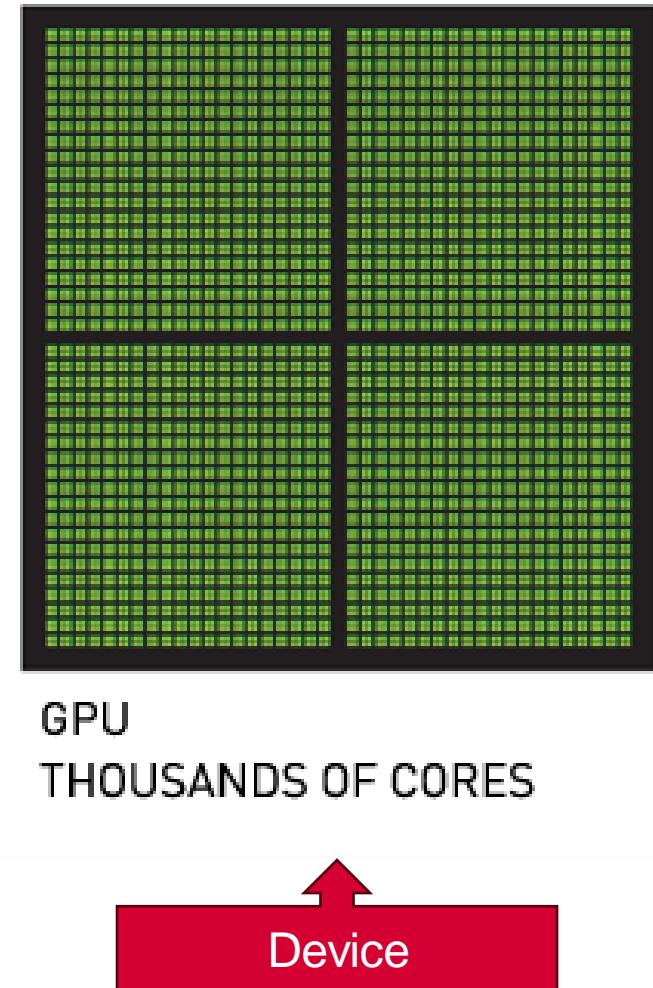
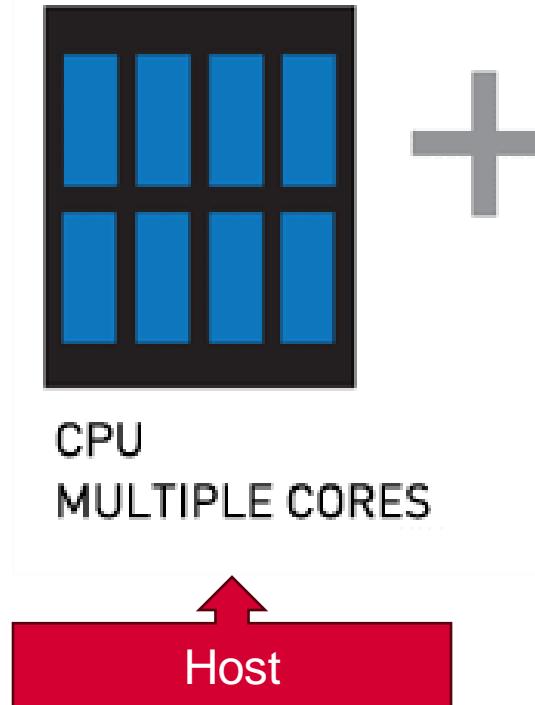
Heterogeneous Exascale Computing Models

- 22 combinations of GPU
- Most common: CUDA
- Aurora: SYCL (primary), HIP
- Frontier: HIP (primary), SYCL
- El Capitan: HIP (primary), SYCL



GPGPU

- General Purpose computing using GPU
- Work together with CPU
- GPU is designed for parallel processing
- Major manufacturers: Nvidia, AMD, Intel
- Power efficiency **TFlop** per Watt





GPGPU Advantages

Massive parallel processing:
Scientific simulations

Efficient large dataset handling:
Machine learning algorithms

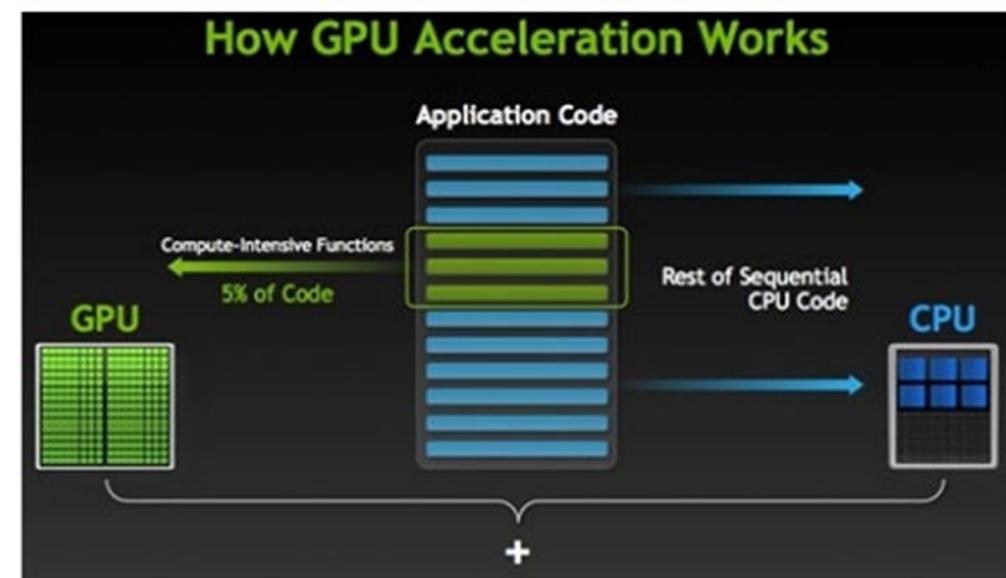
Real-time processing:
Gaming and streaming

Accelerated financial modeling:
Risk assessment and pricing

GPU Programming Model

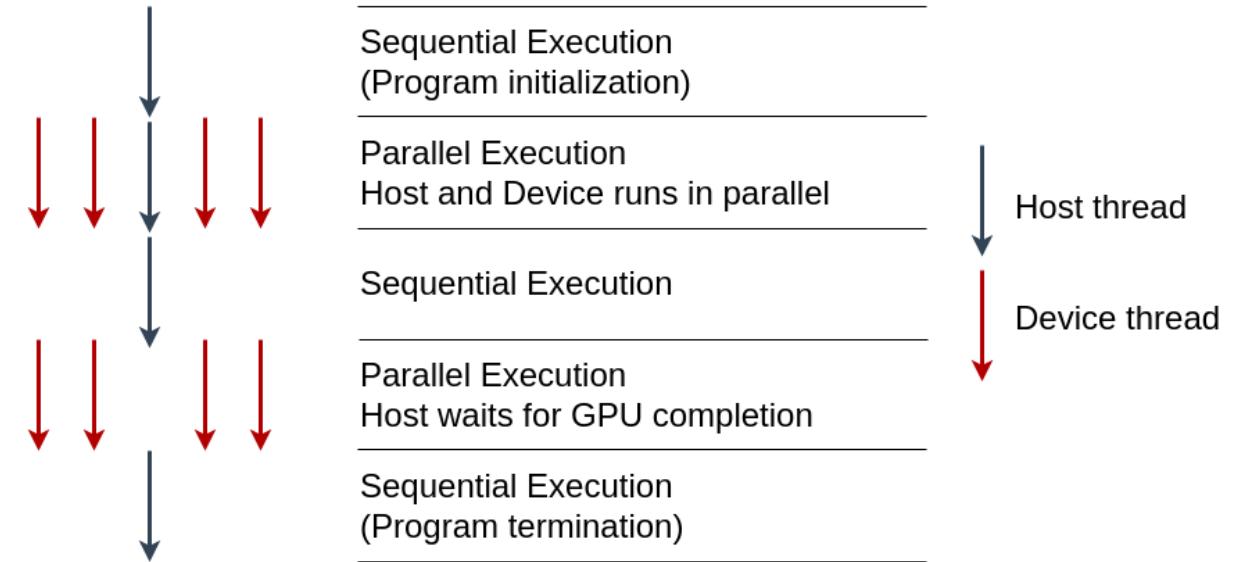
- **CUDA** - proprietary to Nvidia GPUs but most mature and established
- **HIP** - targets AMD GPUs, portable
- **SYCL** - open standard for cross-platform portability
- **DPC++** - Intel's implementation of SYCL
- **OneAPI** - Intel's unified programming model across devices

GPU and CPU communication

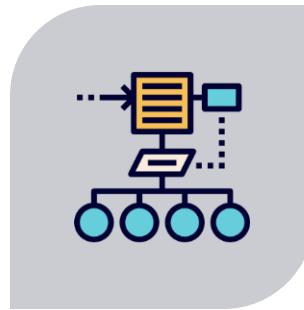


GPU Programming Model

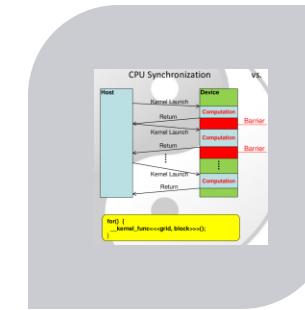
- Globally Sequential Locally Parallel programming pattern
- Invokes parallel **kernels** that execute across a set of threads
- Each thread executes an instance of the kernel



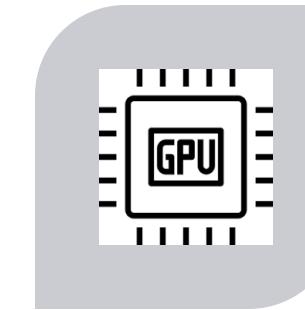
GPGPU Challenges



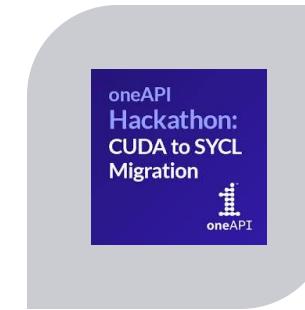
ALGORITHM ADAPTATION
SEQUENTIAL TO
PARALLEL



SYNCHRONIZATION
PUTTING BARRIER



MEMORY MANAGEMENT
H2D AND D2H DATA
TRANSFER



PORTABILITY
DIFFERENT GPU DEVICES

Table of Contents

Exascale Computing
Datalog
Iterative Relational Algebra
Parallel Relational Algebra
Our Approaches
Topological Data Analysis for High-Dimensional Data
Future Research Direction
Conclusion

Datalog: Bottom-Up Logic Programming Language

A lightweight logic-programming language for deductive-database systems

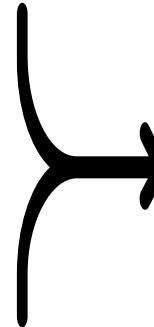


Running the Datalog program extends data from input database creating the output database with all data transitively derivable via the program rules

Datalog Example

Facts:

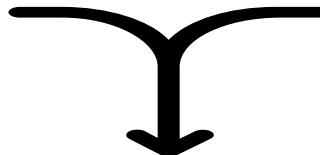
parents(x,y).
children(y,x).



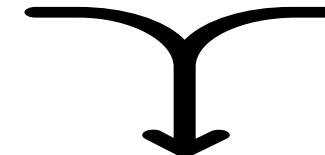
extensional

Rules:

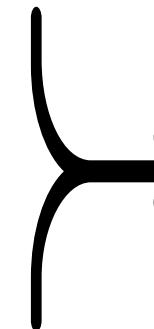
grandparent(x,y) :- parents(x,z), parents(z,y).



head

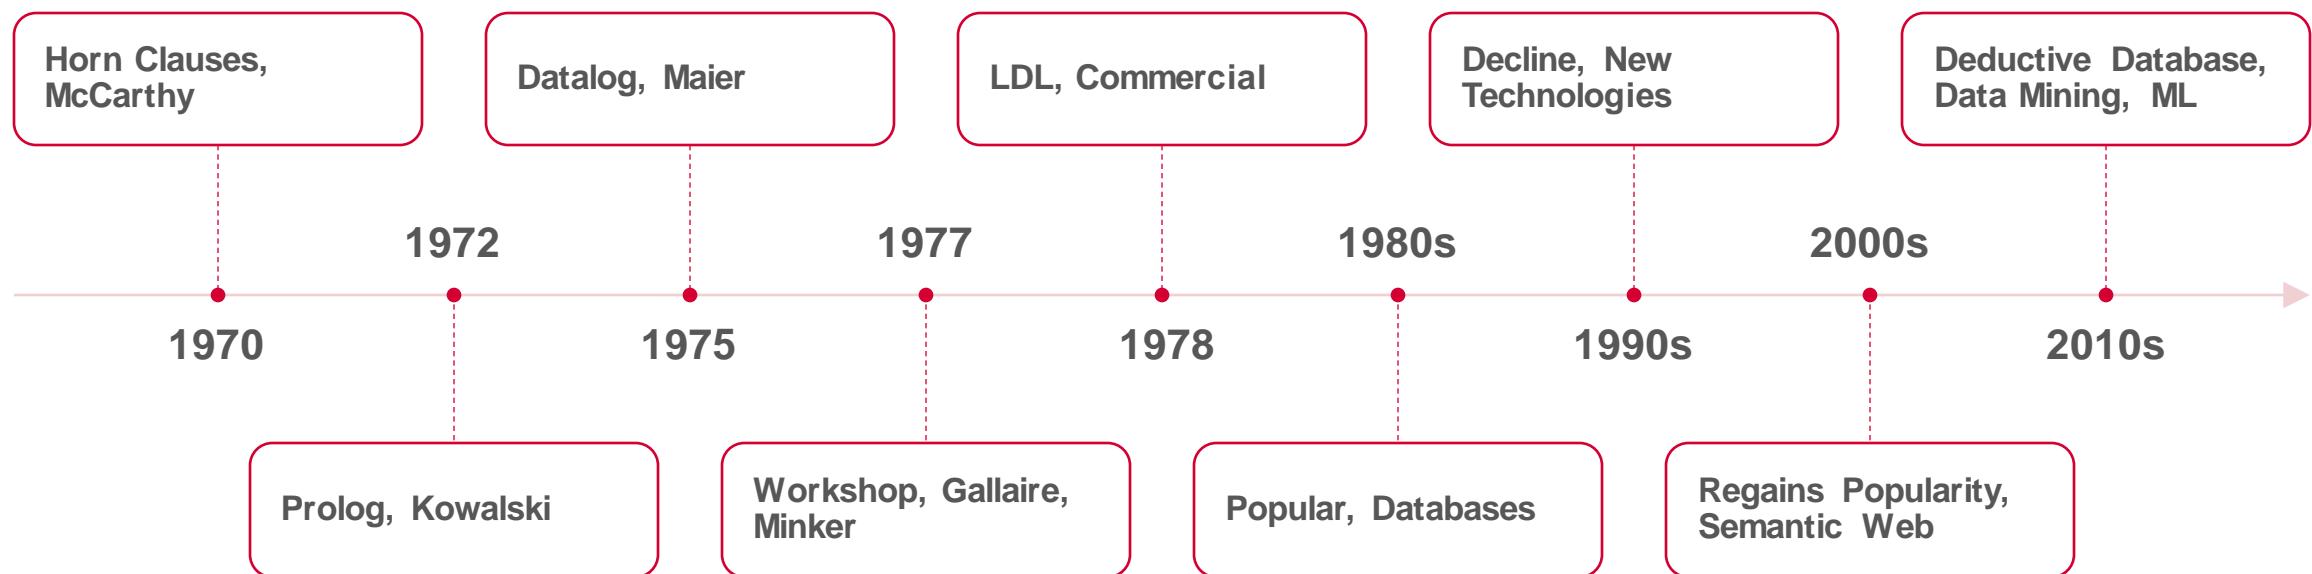


body



intensional

Datalog Timeline

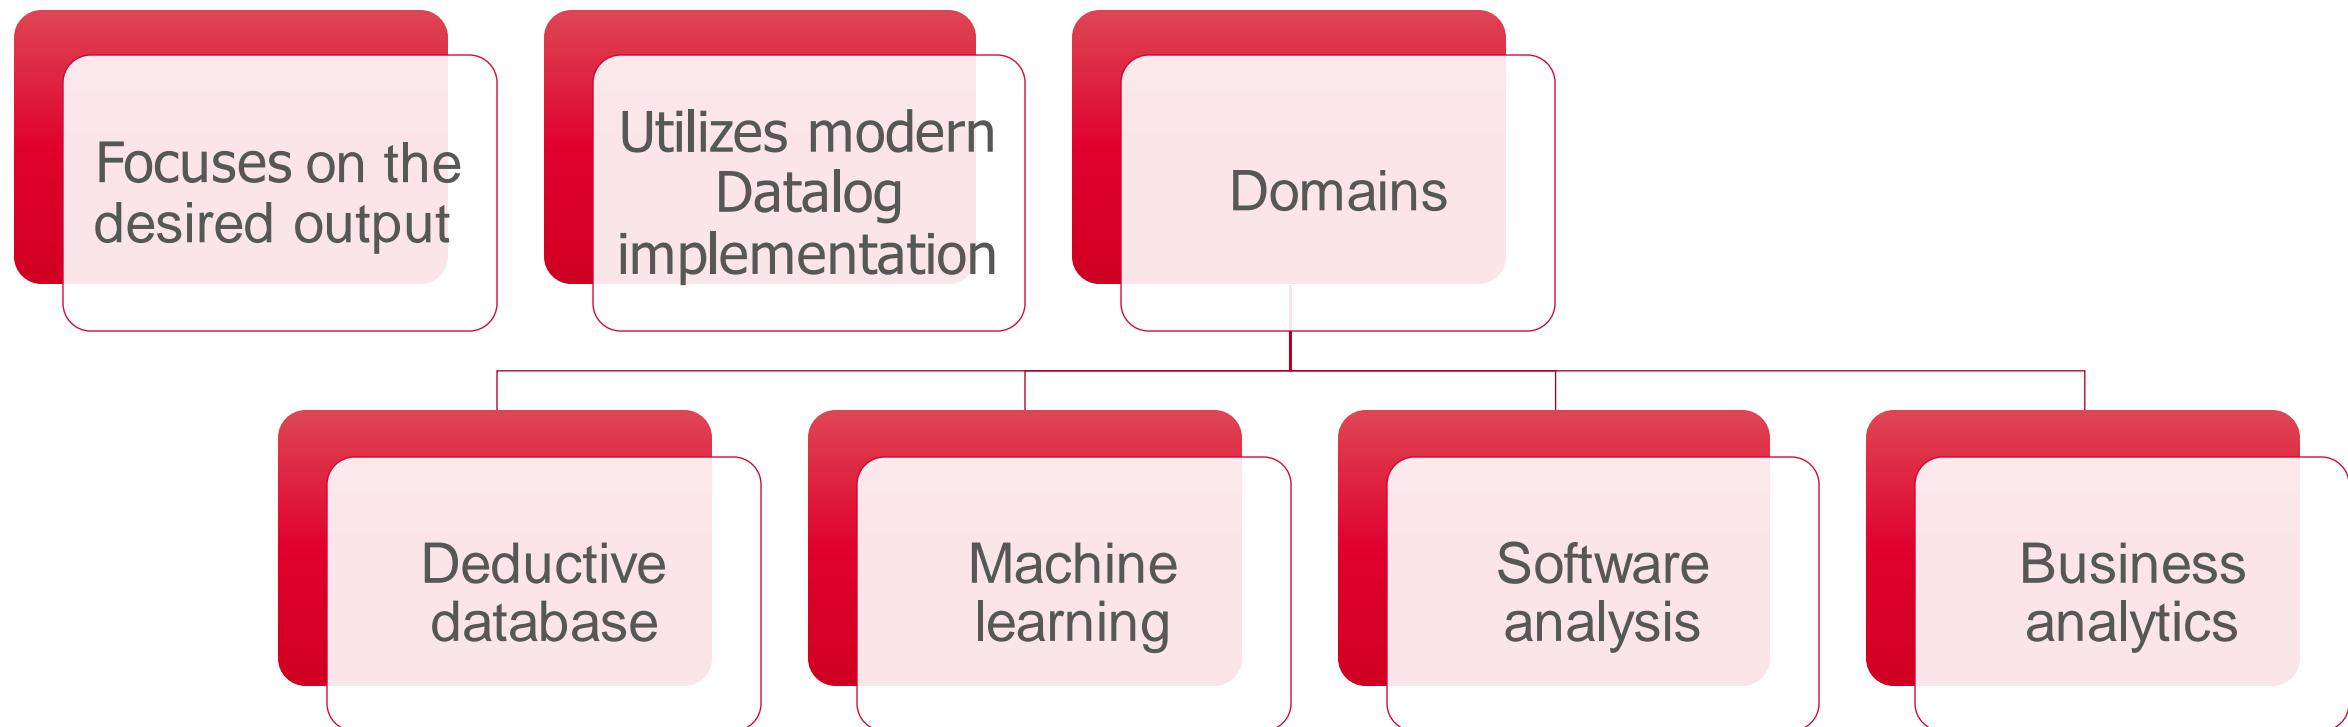


- Stefano Ceri, Georg Gottlob, Letizia Tanca, et al. What you always wanted to know about datalog (and never dared to ask). *IEEE transactions on knowledge and data engineering*, 1(1):146–166, 1989.

- David Maier, K Tuncay Tekle, Michael Kifer, and David S Warren. Datalog: concepts, history, and outlook. In *Declarative Logic Programming: Theory, Systems, and Applications*, pages 3–100. 2018.

- Shan Shan Huang, Todd Jeffrey Green, and Boon Thau Loo. Datalog and emerging applications: An interactive tutorial. In *Proceedings of the 2011 ACM SIGMOD International Conference on Management of Data, SIGMOD '11*, page 1213–1216, New York, NY, USA, 2011. Association for Computing Machinery.

Declarative Analytics Domains



- 16
- Martin Bravenboer and Yannis Smaragdakis. Strictly declarative specification of sophisticated points-to analyses. In Proceedings of the 24th ACM SIGPLAN conference on Object oriented programming systems languages and applications, pages 243–262, 2009.
 - Jiwon Seo, Stephen Guo, and Monica S Lam. Socialite: Datalog extensions for efficient socialnetwork analysis. In 2013 IEEE 29th International Conference on Data Engineering (ICDE), pages 278–289.IEEE, 2013

Classic Problems for Datalog

Transitive closure

Triangle counting

Finding maximal cliques

Finding frequent itemsets

Data mining

Assessment of Datalog Solvers

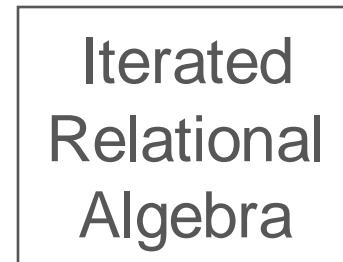
Solver	Technique	Data parallelism	Scalability	Limitations
LogicBlox	Shared memory	No	High	Limited number of threads
Soufflé	Multi core	No	Limited	Internal locking, single node
RadLog	Map Reduce	Yes	Limited	Inadequate for HPC
PRAM	Data parallel	Yes	High	Worse than Soufflé on AWS
SLOG	Data parallel	Yes	High	No GPU usage

- Moham Aref, Balder ten Cate, Todd J. Green, Benny Kimelfeld, Dan Olteanu, Emir Pasalic, Todd L. Veldhuizen, and Geoffrey Washburn. Design and implementation of the logicblox system. In Proceedings of the 2015 ACM SIGMOD International Conference on Management of Data, SIGMOD ’15, page1371–1382, New York, NY, USA, 2015. Association for Computing Machinery.
- Alexander Shkapsky, Mohan Yang, Matteo Interlandi, Hsuan Chiu, Tyson Condie, and Carlo Zaniolo. Big data analytics with datalog queries on spark. In Proceedings of the 2016 International Conference on Management of Data, pages 1135–1149, 2016.
- Herbert Jordan, Bernhard Scholz, and Pavle Subotić. Soufflé: On synthesis of program analyzers. In International Conference on Computer Aided Verification, pages 422–430. Springer, 2016.
- Thomas Gilray, Sidharth Kumar, and Kristopher Micinski. Compiling data-parallel datalog. In Proceedings of the 30th ACM SIGPLAN International Conference on Compiler Construction, CC 2021,page 23–35, New York, NY, USA, 2021. Association for Computing Machinery.
- Thomas Gilray, Arash Sahebolamri, Sidharth Kumar, and Kristopher Micinski. Higher-order, data parallel structured deduction. arXiv preprint arXiv:2211.11573, 2022.

Table of Contents

Exascale Computing
Datalog
Iterative Relational Algebra
Parallel Relational Algebra
Our Approaches
Topological Data Analysis for High-Dimensional Data
Future Research Direction
Conclusion

Bottom-Up Logic Programming with Datalog



Datalog rule for computing **Transitive Closure (TC)**

$$\begin{aligned} T(x, y) &\leftarrow G(x, y) . \\ T(x, z) &\leftarrow T(x, y), G(y, z) . \end{aligned}$$

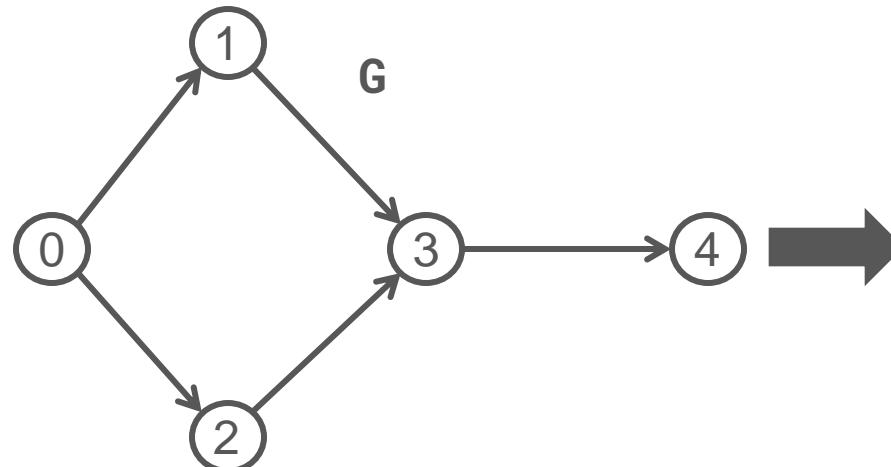

*Operationalized as a **fixed-point iteration** using F_G*

$$F_G(T) \triangleq G \cup \Pi_{1,2}(\rho_{0/1}(T) \bowtie_1 G)$$

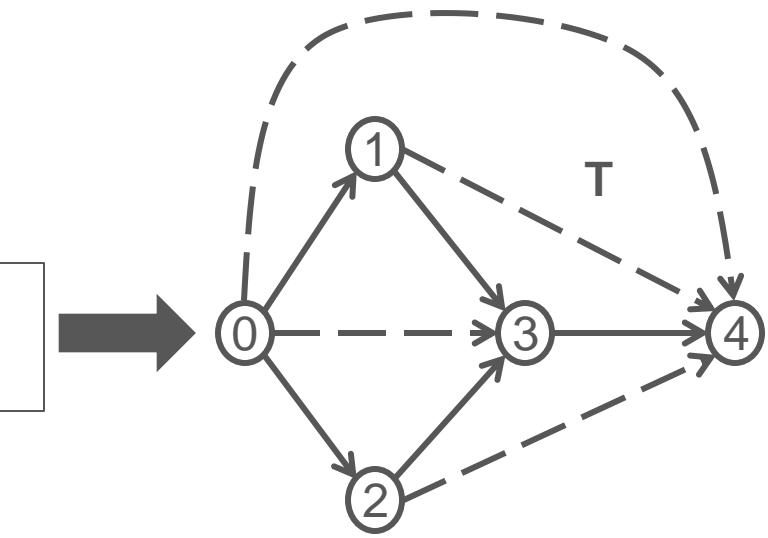
Relational algebra:



Transitive Closure: Logical Inference for Graphs



$T(x, y) \leftarrow G(x, y)$
 $T(x, z) \leftarrow T(x, y), G(y, z)$

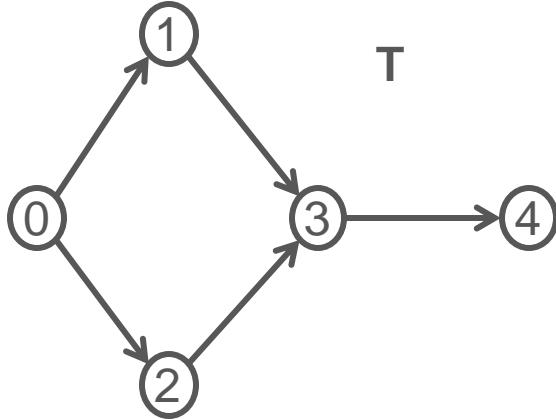


0	1
0	2
1	3
2	3
3	4

0	1
0	2
1	3
2	3
3	4
0	3
1	4
2	4
0	4

$$F_G(T) \triangleq G \cup \Pi_{1,2}(\rho_{0/1}(T) \bowtie_1 G)$$

Transitive Closure: Iterations



0	1
0	2
1	3
2	3
3	4

$$F_G(T) \triangleq G \cup \Pi_{1,2}(\rho_{0/1}(T) \bowtie_1 G)$$

Transitive Closure: Iterations

$\rho_{0/1}(T)$		G	
1	0	0	1
2	0	0	2
3	1	1	3
3	2	2	3
4	3	3	4

$$F_G(T) \triangleq G \cup \Pi_{1,2}(\rho_{0/1}(T) \bowtie_1 G)$$

Transitive Closure: Iterations



$$F_G(T) \triangleq G \cup \Pi_{1,2}(\rho_{0/1}(T) \bowtie_1 G)$$

Transitive Closure: Iterations



$$F_G(T) \triangleq G \cup \Pi_{1,2}(\rho_{0/1}(T) \bowtie_1 G)$$

Transitive Closure: Iterations

$\rho_{0/1}(T)$		G		$\rho_{0/1}(T) \bowtie G$	
1	0	0	1	1	3
2	0	0	2	2	3
3	1	1	3	3	4
3	2	2	3		
4	3	3	4		

$$F_G(T) \triangleq G \cup \Pi_{1,2}(\rho_{0/1}(T) \bowtie_1 G)$$

Transitive Closure: Iterations

$\rho_{0/1}(T)$		G		$\rho_{0/1}(T) \bowtie G$	
1	0	0	1	1	3
2	0	0	2	2	3
3	1	1	3	3	4
3	2	2	3	2	4
4	3	3	4		

The diagram illustrates the iterative computation of the transitive closure. It shows three stages: the initial matrix $\rho_{0/1}(T)$, the graph G , and the result of the first iteration $\rho_{0/1}(T) \bowtie G$. The matrices are shown as 4x2 tables. The \bowtie operation is represented by a large black X symbol between the second and third columns. The cells containing values 2 and 3 in the $\rho_{0/1}(T)$ matrix are highlighted with red borders, indicating they are being updated. The corresponding edges in the graph G are also highlighted with red borders.

$$F_G(T) \triangleq G \cup \Pi_{1,2}(\rho_{0/1}(T) \bowtie_1 G)$$

Transitive Closure: Iterations

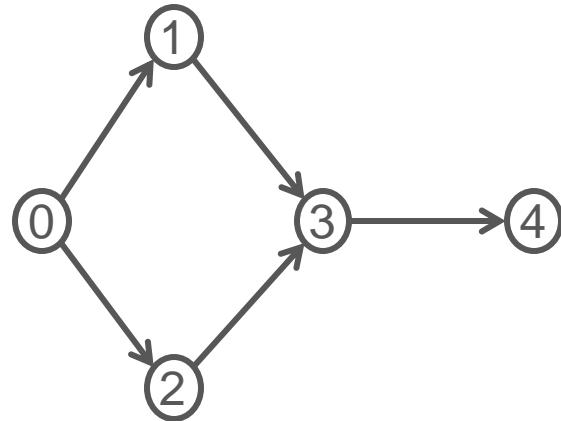
$\rho_{0/1}(T)$		G	
1	0	0	1
2	0	0	2
3	1	1	3
3	2	2	3
4	3	3	4



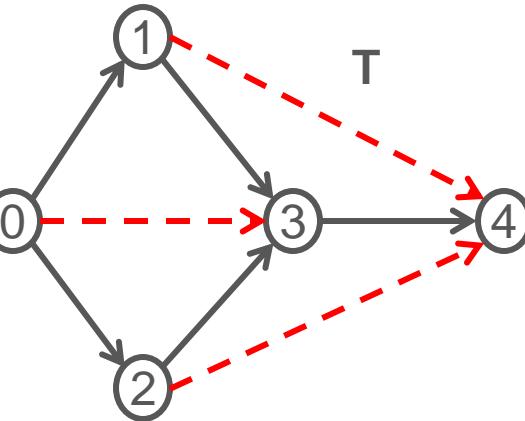
$\rho_{0/1}(T) \bowtie G$			$\Pi_{1,2}(\rho_{0/1}(T) \bowtie G)$	
1	0	3	0	3
2	0	3	1	4
3	1	4	2	4
3	2	4		

$$F_G(T) \triangleq G \cup \Pi_{1,2}(\rho_{0/1}(T) \bowtie_1 G)$$

Transitive Closure: Iterations 1



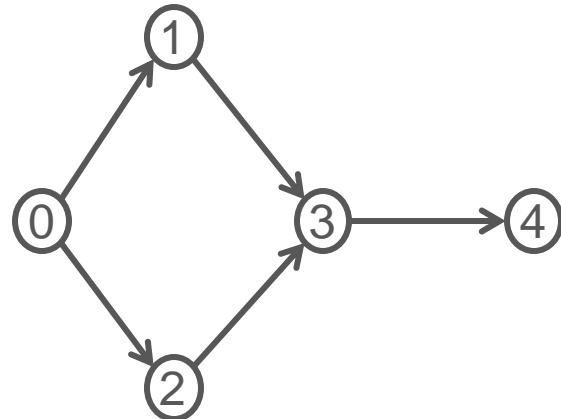
0	1
0	2
1	3
2	3
3	4



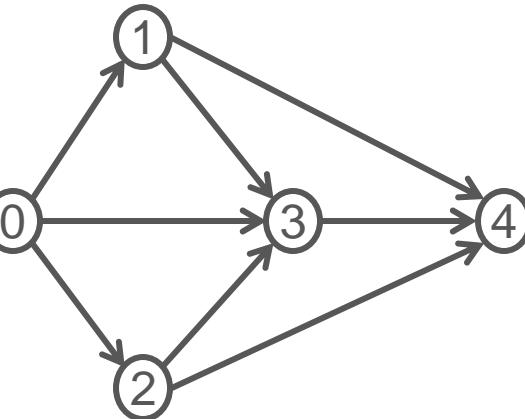
0	1
0	2
1	3
2	3
3	4
0	3
1	4
2	4

$$F_G(T) \triangleq G \cup \Pi_{1,2}(\rho_{0/1}(T) \bowtie_1 G)$$

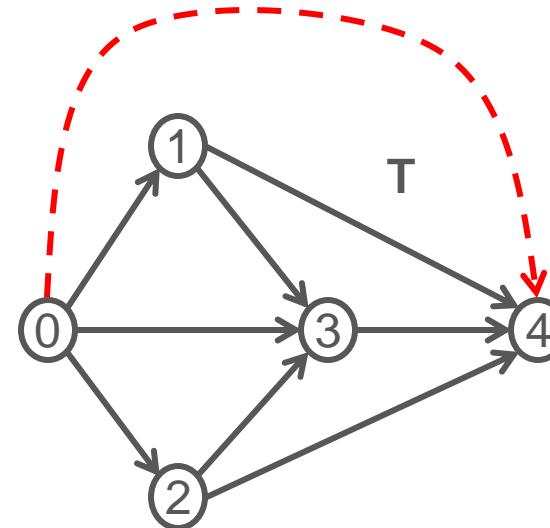
Transitive Closure: Iterations 2



0	1
0	2
1	3
2	3
3	4



0	1
0	2
1	3
2	3
3	4
0	3
1	4
2	4



0	1
0	2
1	3
2	3
3	4
0	3
1	4
2	4
0	4

$$F_G(T) \triangleq G \cup \Pi_{1,2}(\rho_{0/1}(T) \bowtie_1 G)$$

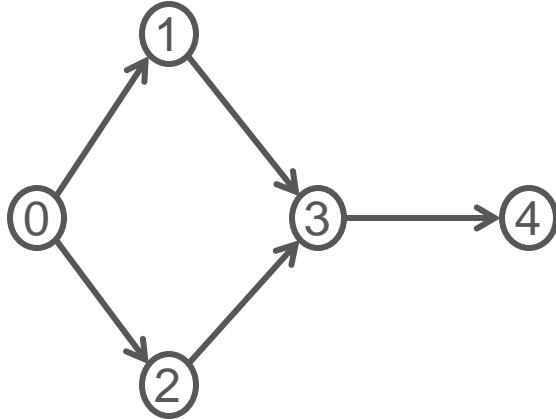
Transitive Closure: Iterations

3

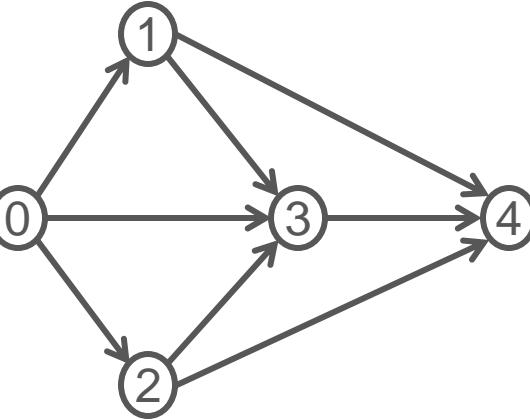
Union

Projection

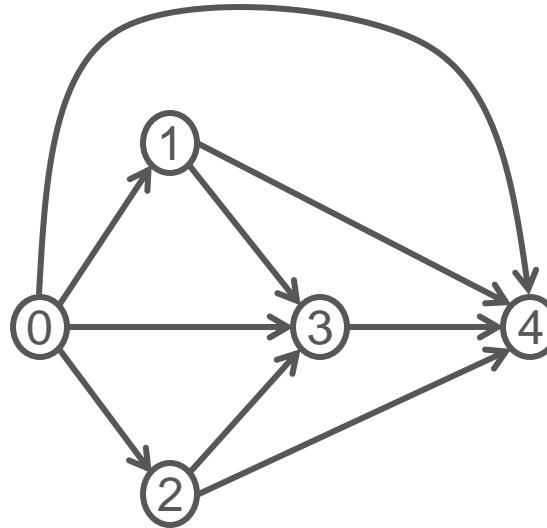
Join



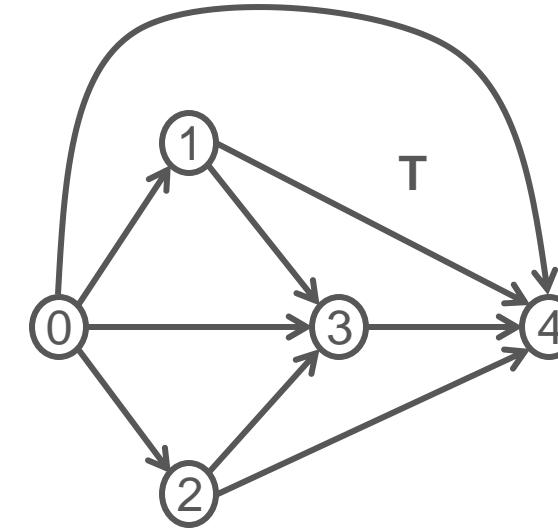
0	1
0	2
1	3
2	3
3	4



0	1
0	2
1	3
2	3
3	4
0	3
1	4
2	4



0	1
0	2
1	3
2	3
3	4
0	3
1	4
2	4
0	4



0	1
0	2
1	3
2	3
3	4
0	3
1	4
2	4
0	4

Table of Contents

Exascale Computing
Datalog
Iterative Relational Algebra
Parallel Relational Algebra
Our Approaches
Topological Data Analysis for High-Dimensional Data
Future Research Direction
Conclusion

How can we do join in parallel?

User (Outer Relation)

UserID	UserName	UserEmail
101	Alice	alice@example.com
102	Bob	bob@example.com
103	Eve	eve@example.com



Order (Inner Relation)

UserID	OrderTotal	Items
101	25.69	2
102	145.66	3
103	12.11	1
103	44.00	2

Parallel Join

User (Outer Relation)

UserID	UserName	UserEmail
101	Alice	alice@example.com
102	Bob	bob@example.com
103	Eve	eve@example.com

Order (Inner Relation)

UserID	OrderTotal	Items
101	25.69	2
102	145.66	3
103	12.11	1
103	44.00	2



Thread 1

Thread 2

Thread 3

All executing in parallel

Parallel Join

Parallel Join

Algorithms

Hardwares

Load balancing

Comparison of Join Algorithms

Algorithm	Pros	Cons	Dataset
Non indexed nested loop	Simple	High time complexity, low efficiency	Small
Indexed nested loop	Index for faster performance	High memory usage	Small to medium, indexed tables
Sort merge	Scalable, handles non-equality joins	Requires sorting, high memory usage	Medium
Hash	Scalable, handles equality joins	Requires partitioning, cannot do non-equality	Large
Hybrid (sort + hash)	Reduce communication cost, Handles skewed data	Complex, requires tuning	Large

- Bingsheng He, Ke Yang, Rui Fang, Mian Lu, Naga Govindaraju, Qiong Luo, and Pedro Sander. Relational joins on graphics processors. In Proceedings of the 2008 ACM SIGMOD international conference on Management of data, pages 511–524, 2008.
- Ran Rui, Hao Li, and Yi-Cheng Tu. Join algorithms on gpus: A revisit after seven years. In 2015 IEEE International Conference on Big Data (Big Data), pages 2541–2550. IEEE, 2015.
- Ran Rui, Hao Li, and Yi Cheng Tu. Efficient join algorithms for large database tables in a multi-gpu environment. Proceedings of the VLDB Endowment, 14:708–720, 2020.
- Daniel Zinn, Haicheng Wu, Jin Wang, Molham Aref, and Sudhakar Yalamanchili. General-purpose join algorithms for large graph triangle listing on heterogeneous systems. In Proceedings of the 9th Annual Workshop on General Purpose Processing Using Graphics Processing Unit, pages 12–21, 2016.
- Claude Barthels, Ingo Müller, Timo Schneider, Gustavo Alonso, and Torsten Hoefer. Distributed join algorithms on thousands of cores. Proc. VLDB Endow., 10(5):517–528, jan 2017
- Chengxin Guo, Hong Chen, Feng Zhang, and Cuiping Li. Parallel hybrid join algorithm on gpu. 2019 IEEE 21st International Conference on High Performance Computing and Communications; IEEE 17th International Conference on Smart City; IEEE 5th International Conference on Data Science and Systems (HPCC/SmartCity/DSS), pages 1572–1579, 2019

Parallel Join

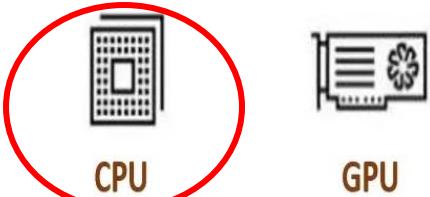
Parallel Join

Algorithms

Hardwares

Load balancing

Join on CPUs: Multi-core HJ and SMJ



- Barthels et al. compared HJ and SMJ on distributed multi-core systems using 4096 processor cores (MPI)
- Communication inefficiencies can significantly impede performance
- More cores not always means better performance, sometimes **degrade**
- Bounded by the **low computation** capacity of CPUs and **higher communication** costs

Join on GPUs: Benchmark



CPU



GPU

- Rui et al. assessed NINLJ, INLJ, SMJ, and HJ on modern GPU
- Modern GPUs can lead to **20X** speedup VS **7X** speedup of old GPUs
- Not suitable for HPC systems with multiple GPU environments
- New GPU architecture is introduced (Nvidia Hopper architecture)

Join on GPUs: WarpDrive



CPU



GPU

- Jünger et al. presented a single-node multi-GPU hashing for hashjoin
- Attained better memory coalescing
- Hashtable insertion rate:
 - 1.4B keys/sec (single GPU)
 - 4.3B keys/sec (4 GPUs)
- 32 bit keys only with no deduplication
- Incremental study: **WarpCore** supports 64 bit keys

40 • Daniel Jünger, Christian Hundt, and Bertil Schmidt. Warpdrive: Massively parallel hashing on multigpu nodes. In 2018 IEEE International Parallel and Distributed Processing Symposium (IPDPS), pages 441–450. IEEE, 2018.

• Daniel Jünger, Robin Kobus, André Müller, Christian Hundt, Kai Xu, Weiguo Liu, and Bertil Schmidt. Warpcore: A library for fast hash tables on gpus. In 2020 IEEE 27th International Conference on HighPerformance Computing, Data, and Analytics (HiPC), pages 11–20, 2020.

Join on GPUs: Secondary storage



CPU



GPU

- Zinn et al. presented multi-GPUs join with secondary storage support
- Focused on triangle counting problem
- Allow overfitting global GPU memory with **SSD** storage
- Leapfrog Triejoin (LFTJ) algorithm: a worst-case-optimal algorithm
- No mention of multi-node heterogeneous systems

Parallel Join

Parallel Join

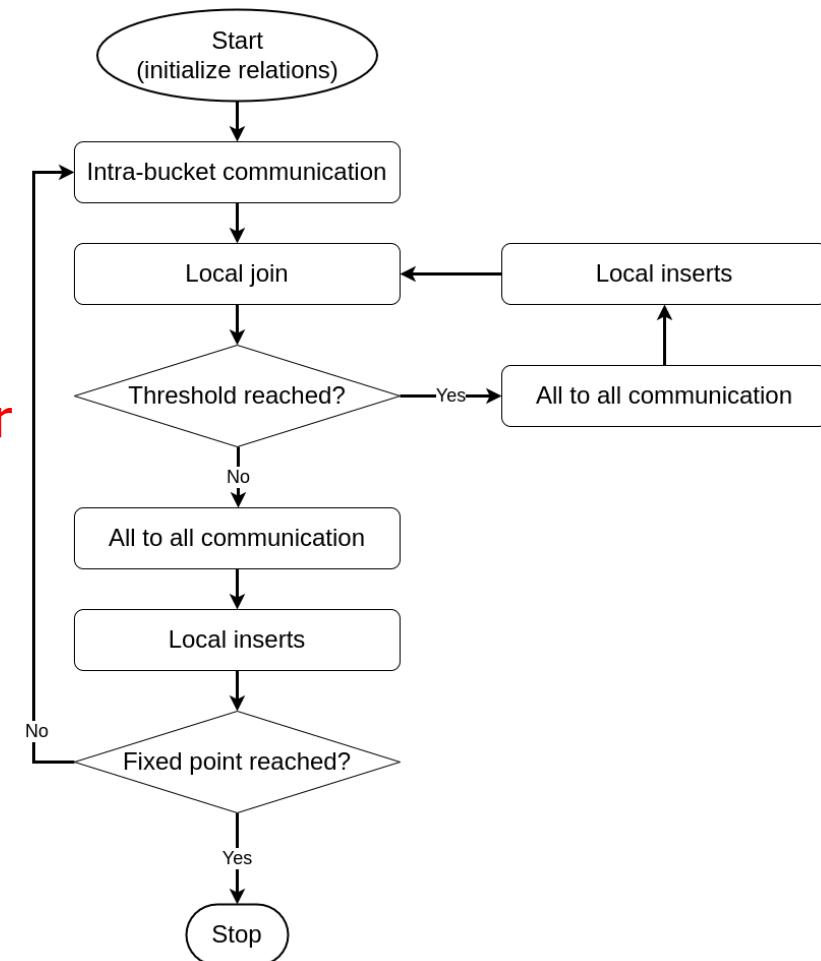
Algorithms

Hardwares

Load balancing

Load Balancing Multi-node Joins

- 2 types of load imbalance issues:
 - Spatial: across multiple computing nodes
 - Temporal: distributing newly discovered tuples
 - Solutions: **Bucket refinement, consolidation, Iteration rollover**
- Range of favorable scaling exists: 256 to 32K processes
- Performance deteriorates beyond that range



Handling temporal load imbalance during the join

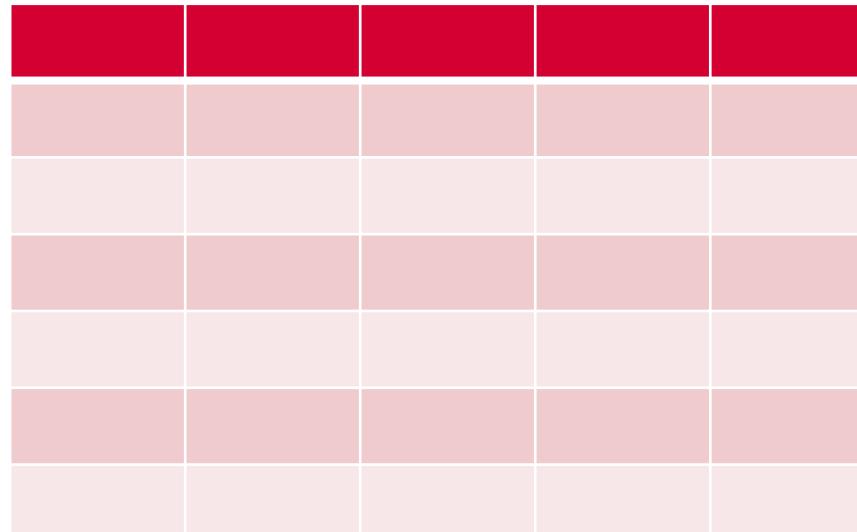
Table of Contents

Exascale Computing
Datalog
Iterative Relational Algebra
Parallel Relational Algebra
Our Approaches
Topological Data Analysis for High-Dimensional Data
Future Research Direction
Conclusion

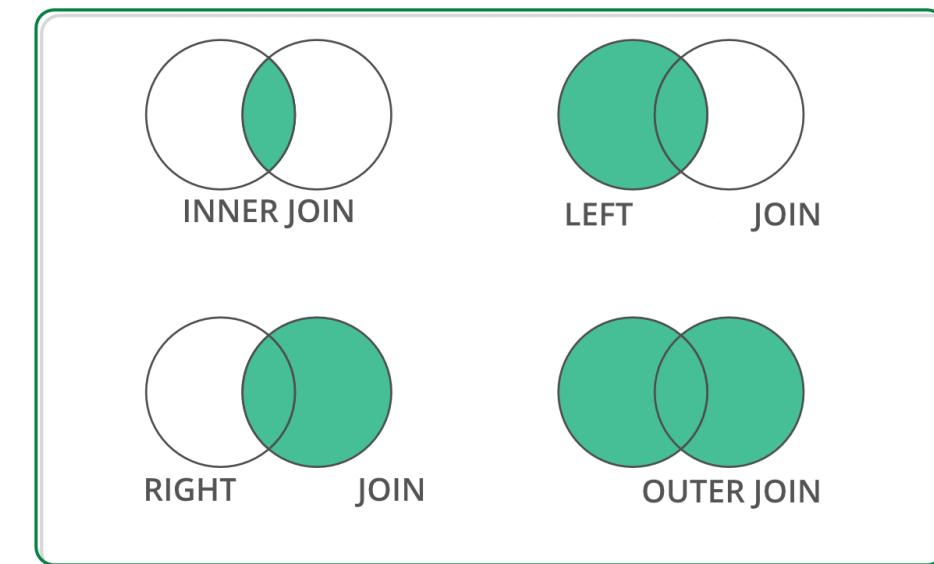
First approach: Off-the-shelf RA Data Structure

Off-the-shelf Data Structure for RA Operation

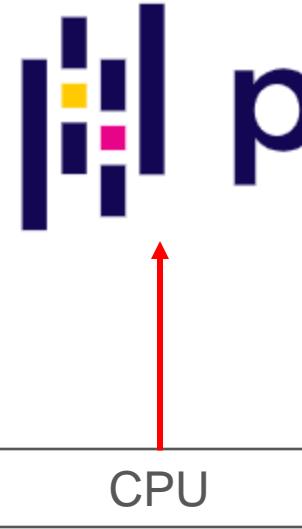
DataFrame: 2D labeled tabular data structure



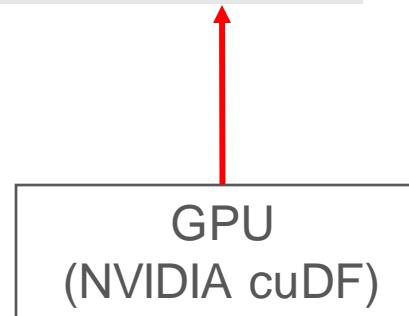
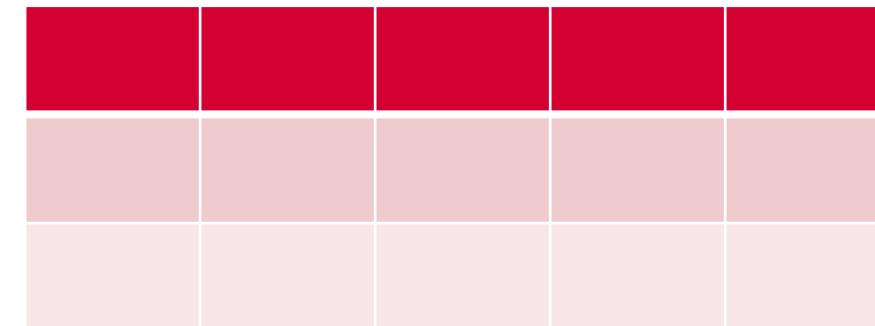
DataFrame has RA primitives APIs



Off-the-shelf Python Libraries



DataFrame: 2D labeled tabular data structure



Both supports join operation with similar APIs

- Reback, J., McKinney, W., Van Den Bossche, J., Augspurger, T., Cloud, P., Klein, A., ... & Seabold, S. (2020). *pandas-dev/pandas: Pandas 1.0. 5*. Zenodo.
- Chen, D. Y. (2017). *Pandas for everyone: Python data analysis*. Addison-Wesley Professional.
- Green, O., Du, Z., Patel, S., Xie, Z., Liu, H., & Bader, D. A. (2021, December). Anti-Section Transitive Closure. In *2021 IEEE 28th International Conference on High Performance Computing, Data, and Analytics (HiPC)* (pp. 192-201). IEEE.
- Fender, A., Rees, B., & Eaton, J. *RAPIDS cuGraph*. In *Massive Graph Analytics* (pp. 483-493). Chapman and Hall/CRC.

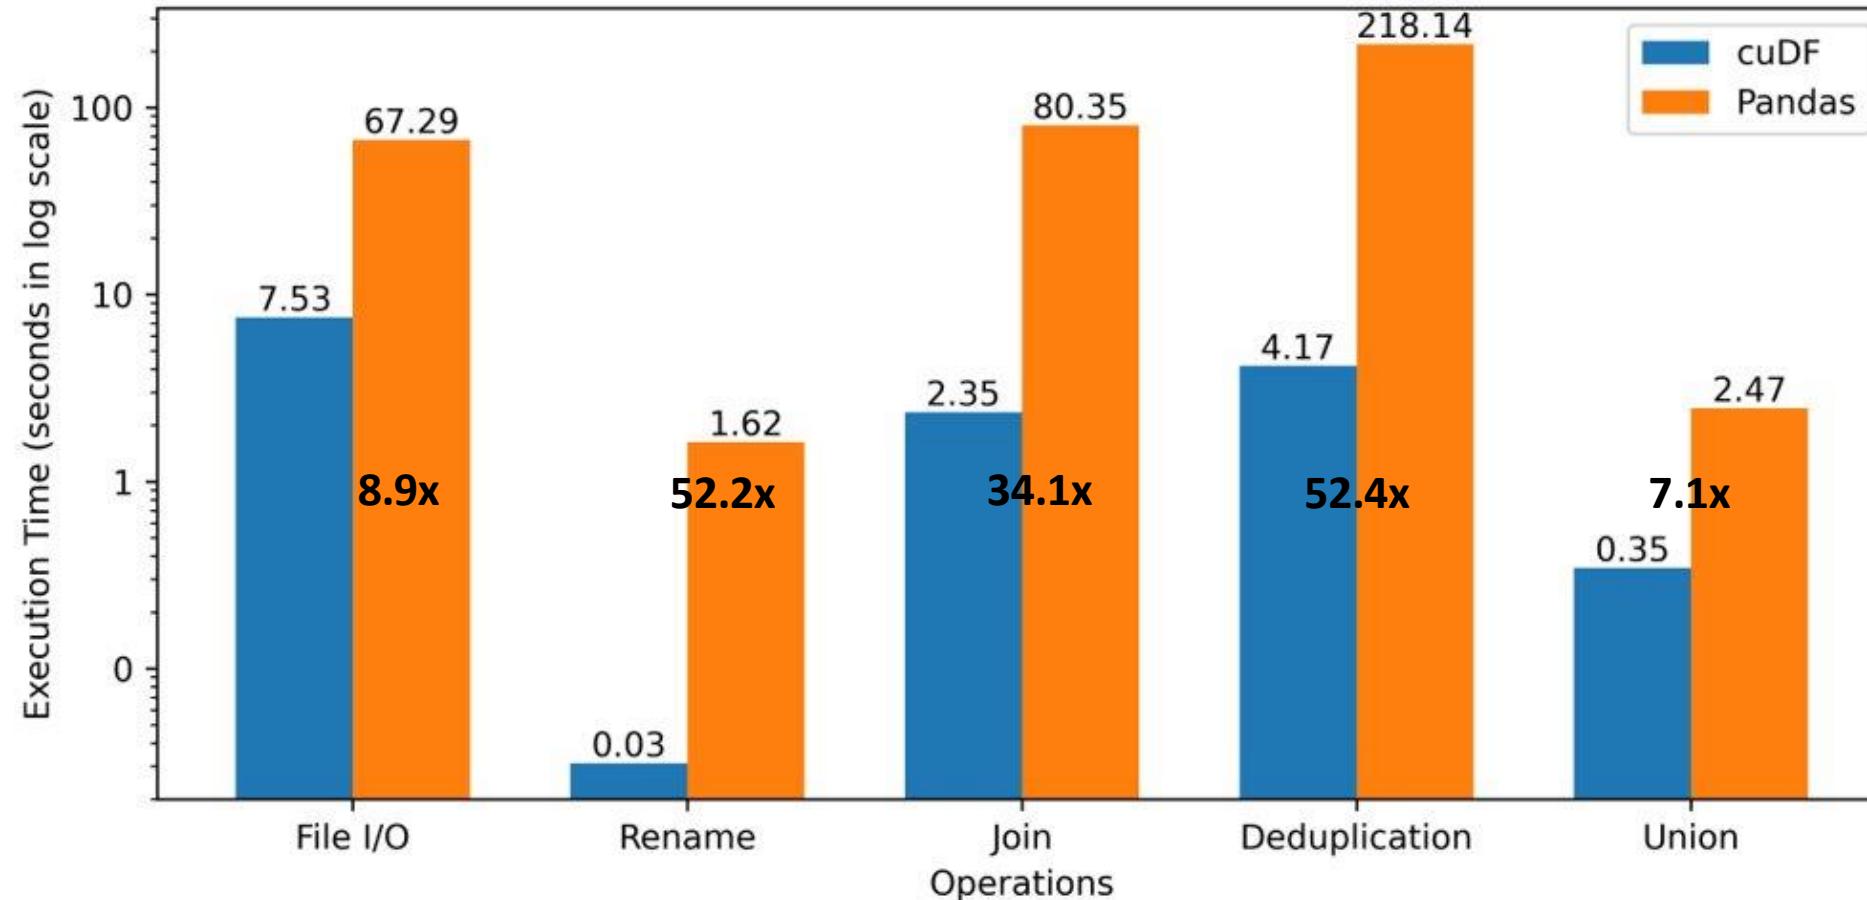
CPU (Pandas) and GPU (cuDF)

```
import pandas as pd
import cudf

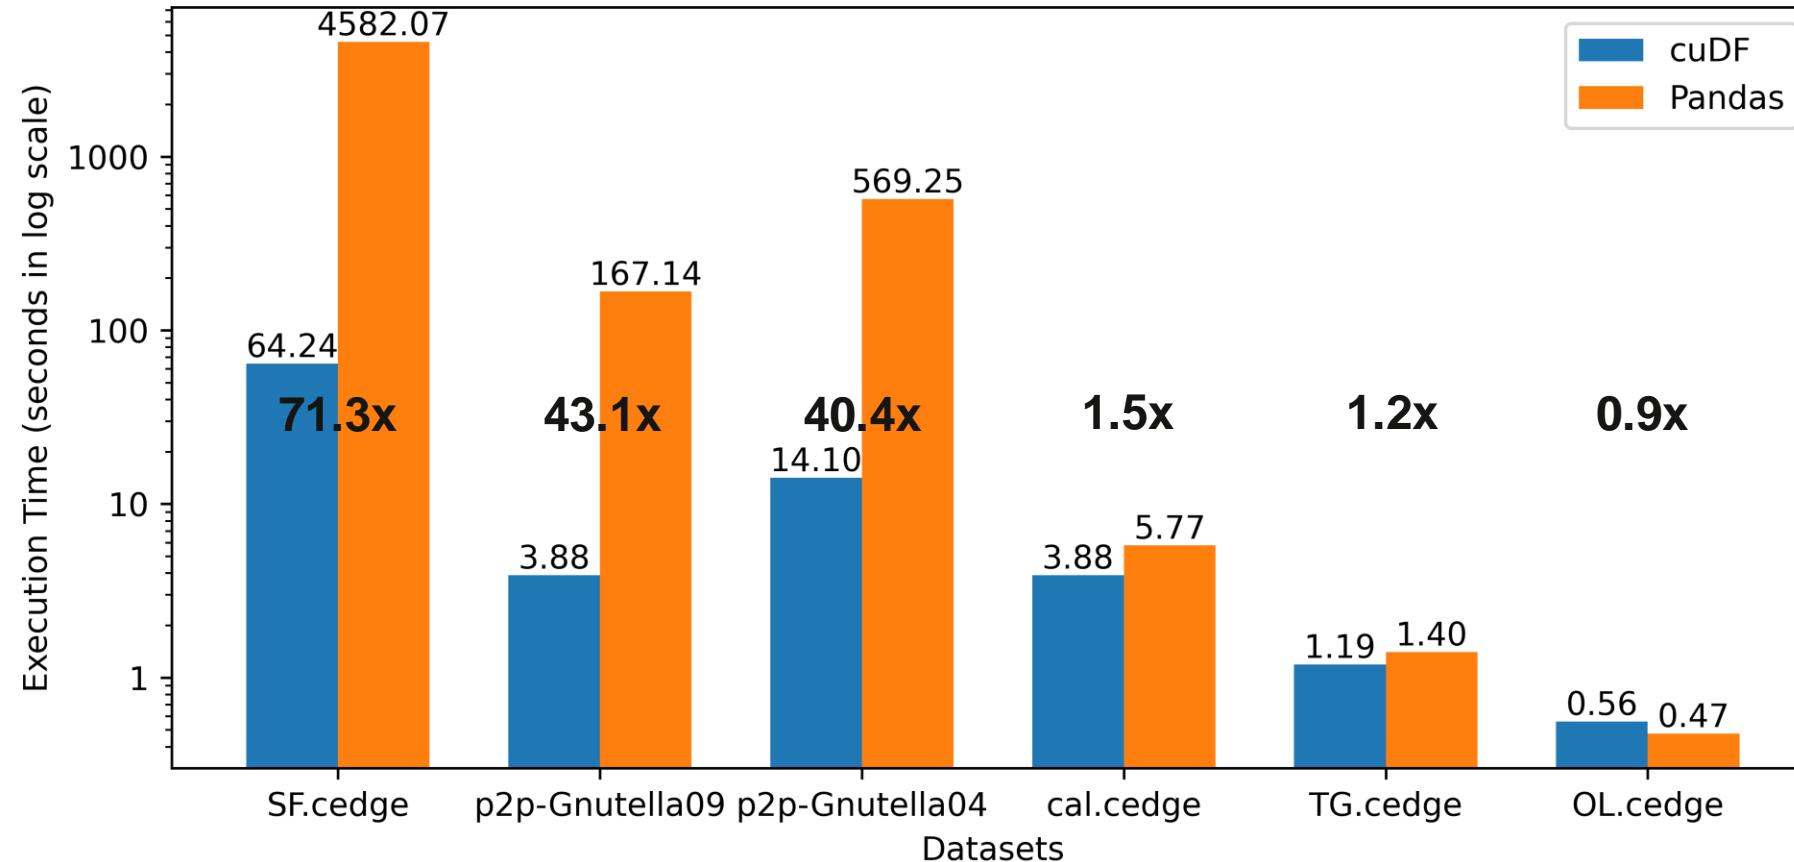
def get_read_csv(filename, method='cudf', n):
    column_names = ['column 1', 'column 2']
    if method == 'df':
        return pd.read_csv(filename, sep='\t', header=None,
                           names=column_names, nrows=n)
    return cudf.read_csv(filename, sep='\t', header=None,
                           names=column_names, nrows=n)

def get_join(relation_1, relation_2):
    column_names = ['column 1', 'column 2']
    return relation_1.merge(relation_2, on=column_names[0],
                           how="inner",
                           suffixes=('_relation_1', '_relation_2'))
```

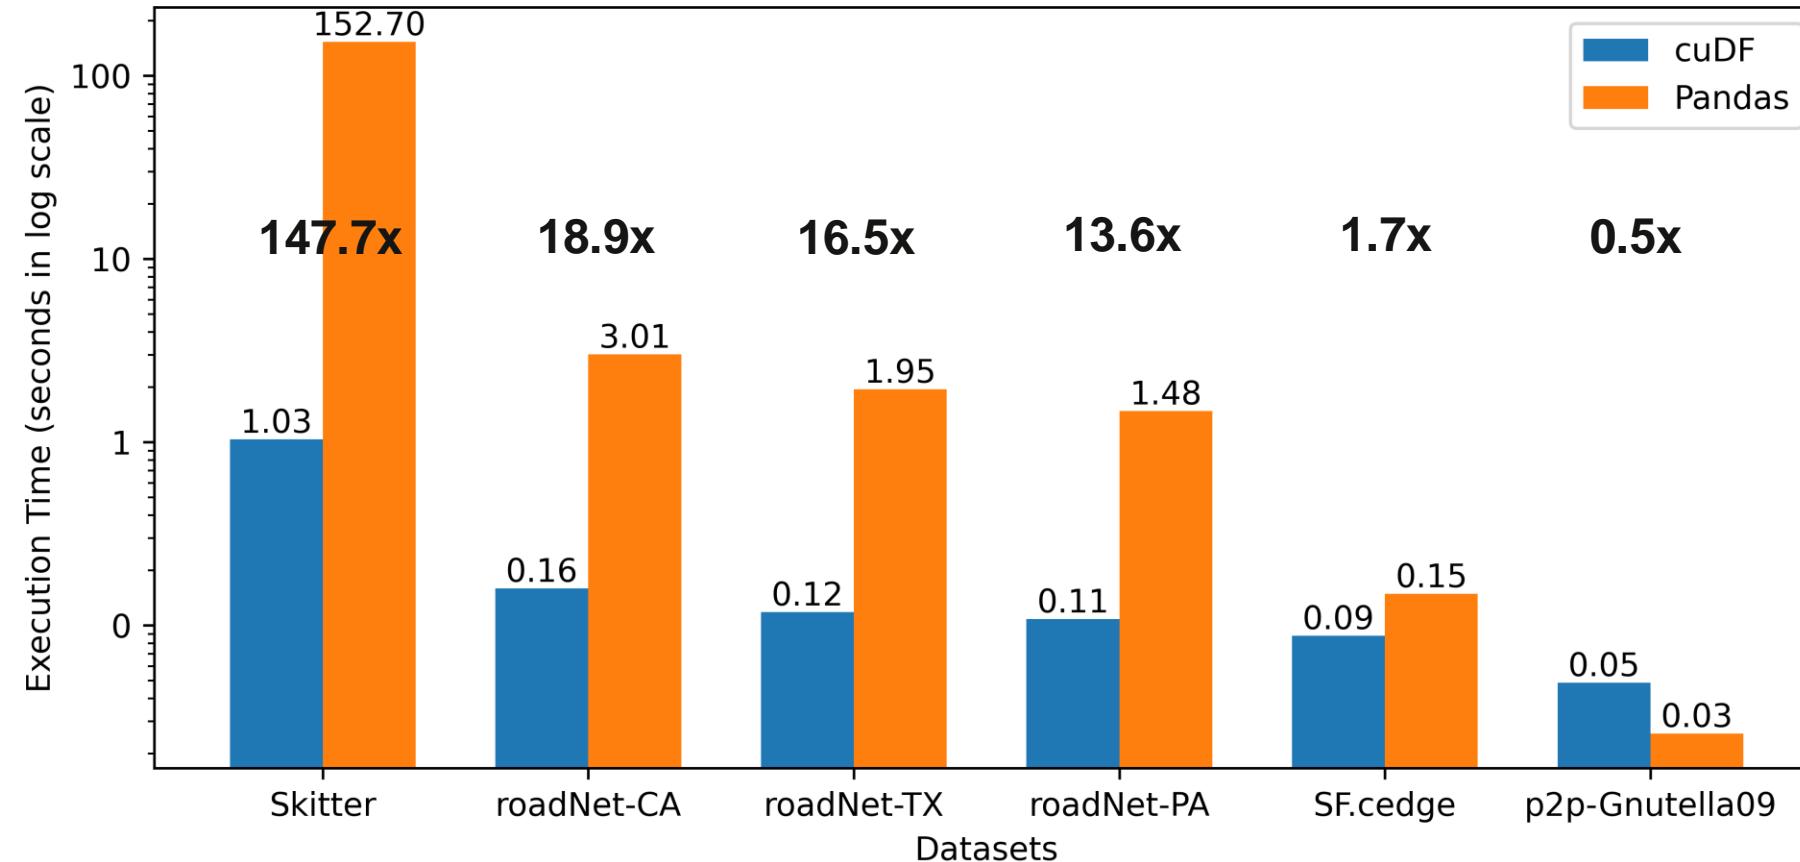
Performance Improvement: RA operations



Performance Improvement: Transitive Closure



Performance Improvement: Triangle Counting



DataFrame Based Datalog Applications

✓ Advantages

- ✓ Abstract memory management
- ✓ Abstract thread block configuration
- ✓ Same API signatures for CPU and GPU
- ✓ Easy-to-code interface

Improvement Opportunity for TC

Open-addressing based hashtable

Fuse join and projection

Sorted results for deduplication

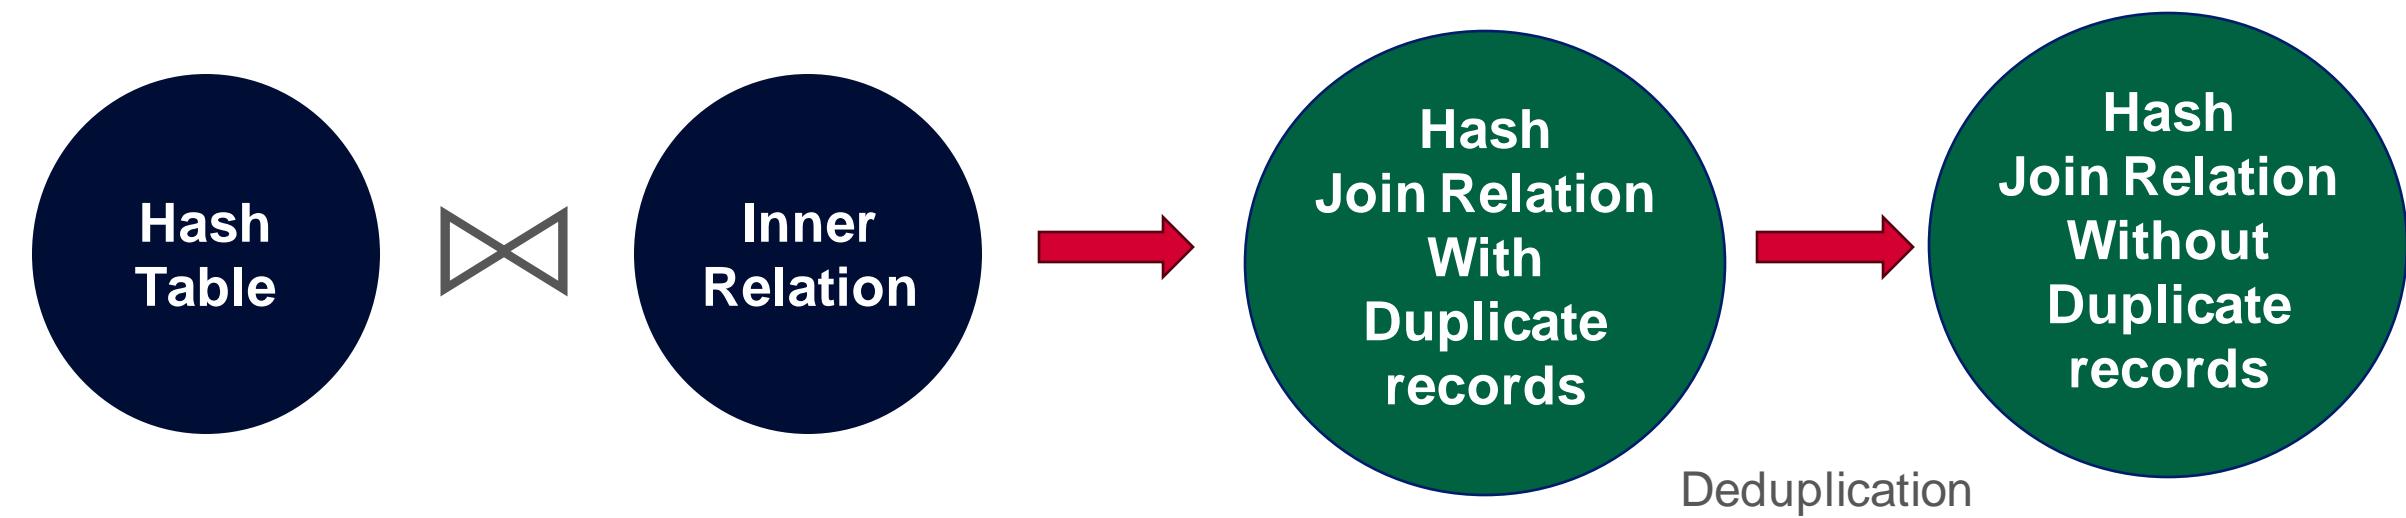
Pinned memory scheme

Intermediate memory clearance

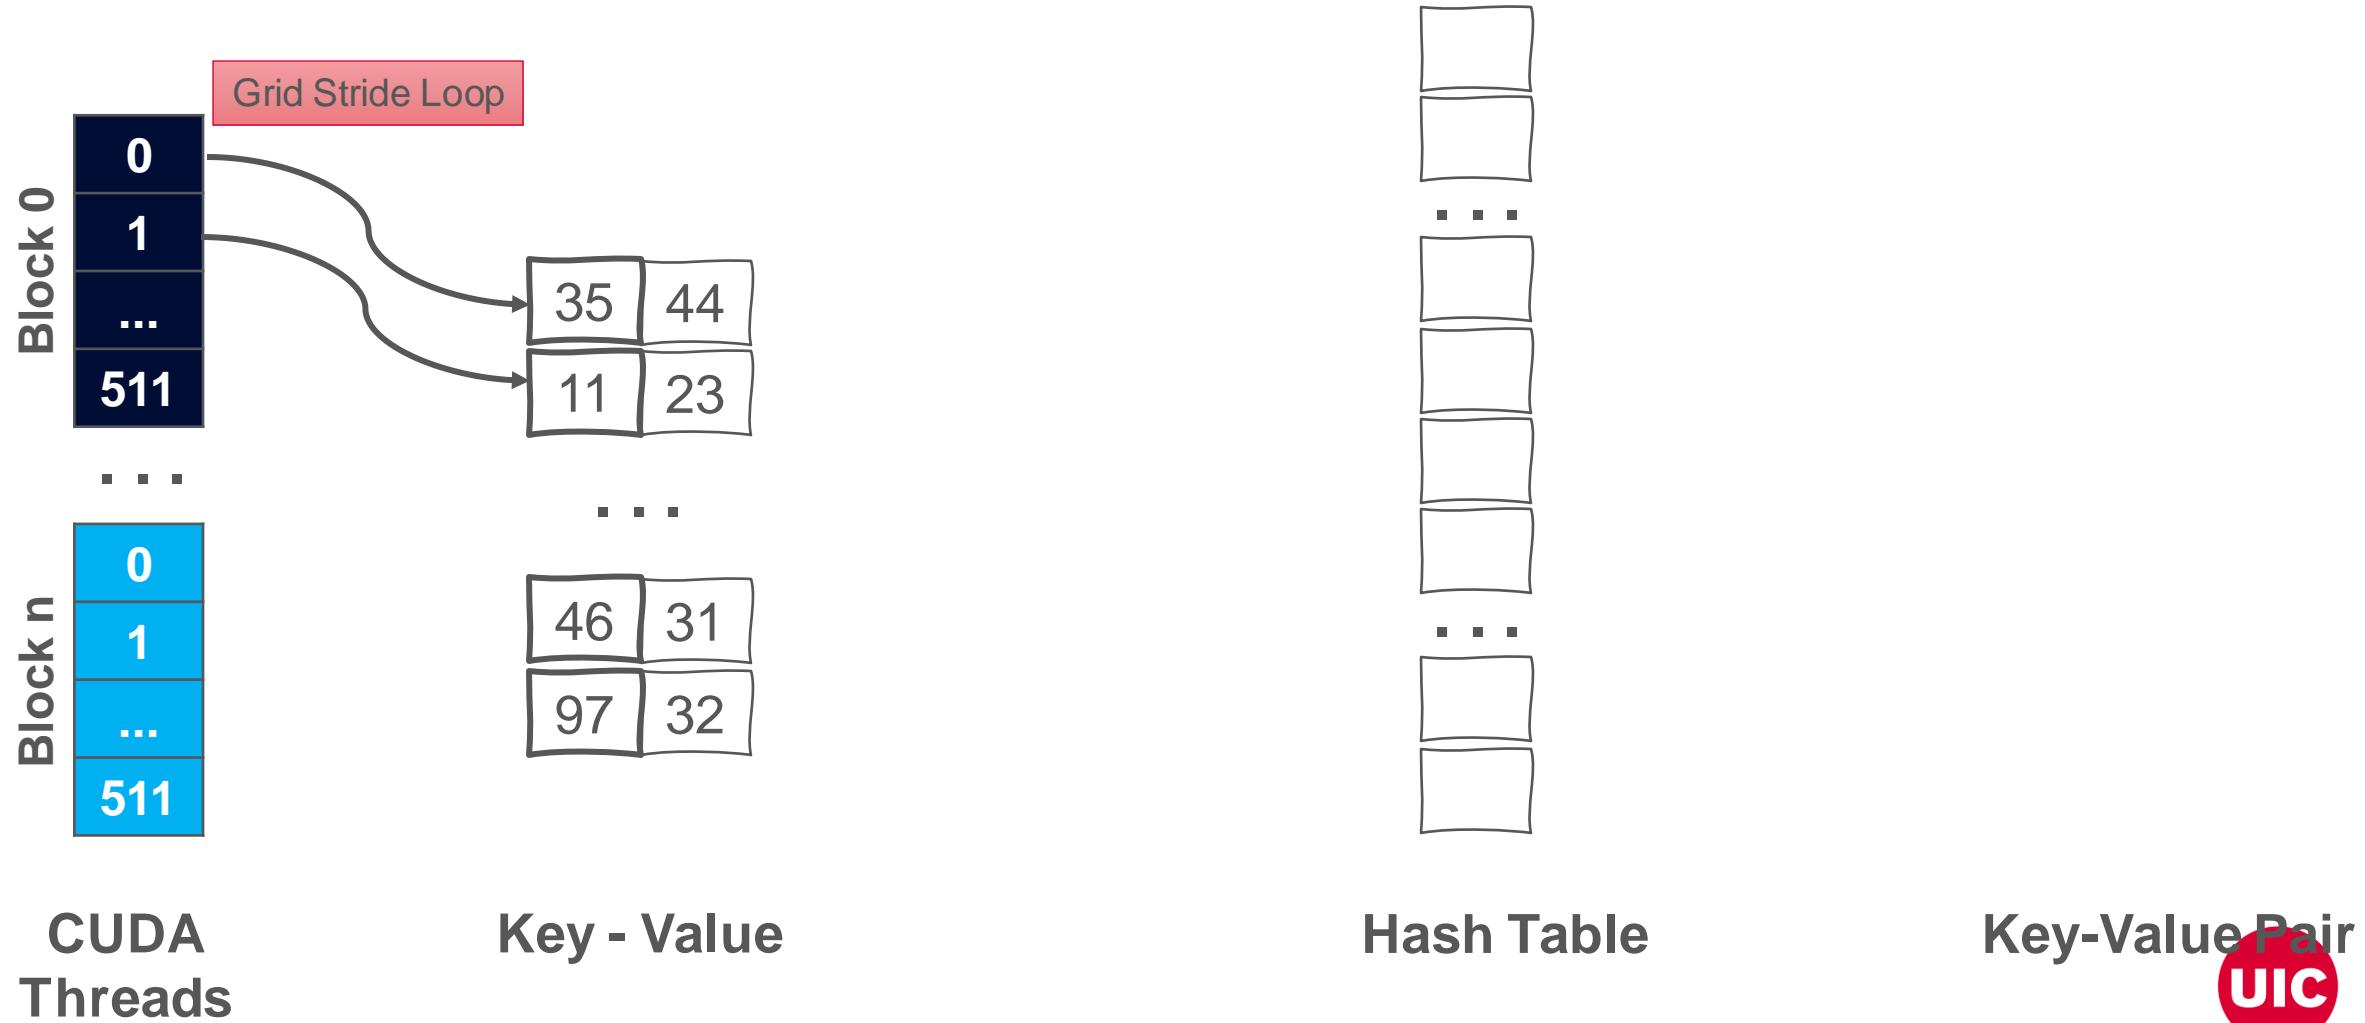
- A. R. Shovon, L. R. Dyken, O. Green, T. Gilray and S. Kumar, "Accelerating Datalog applications with cuDF," 2022 IEEE/ACM Workshop on Irregular Applications: Architectures and Algorithms (IA3), Dallas, TX, USA, 2022, pp. 41-45
- Green, O., Du, Z., Patel, S., Xie, Z., Liu, H., & Bader, D. A. (2021, December). Anti-Section Transitive Closure. In 2021 IEEE 28th International Conference on High Performance Computing, Data, and Analytics (HPC) (pp. 192-201). IEEE.

Second approach: Low level iterative hash join

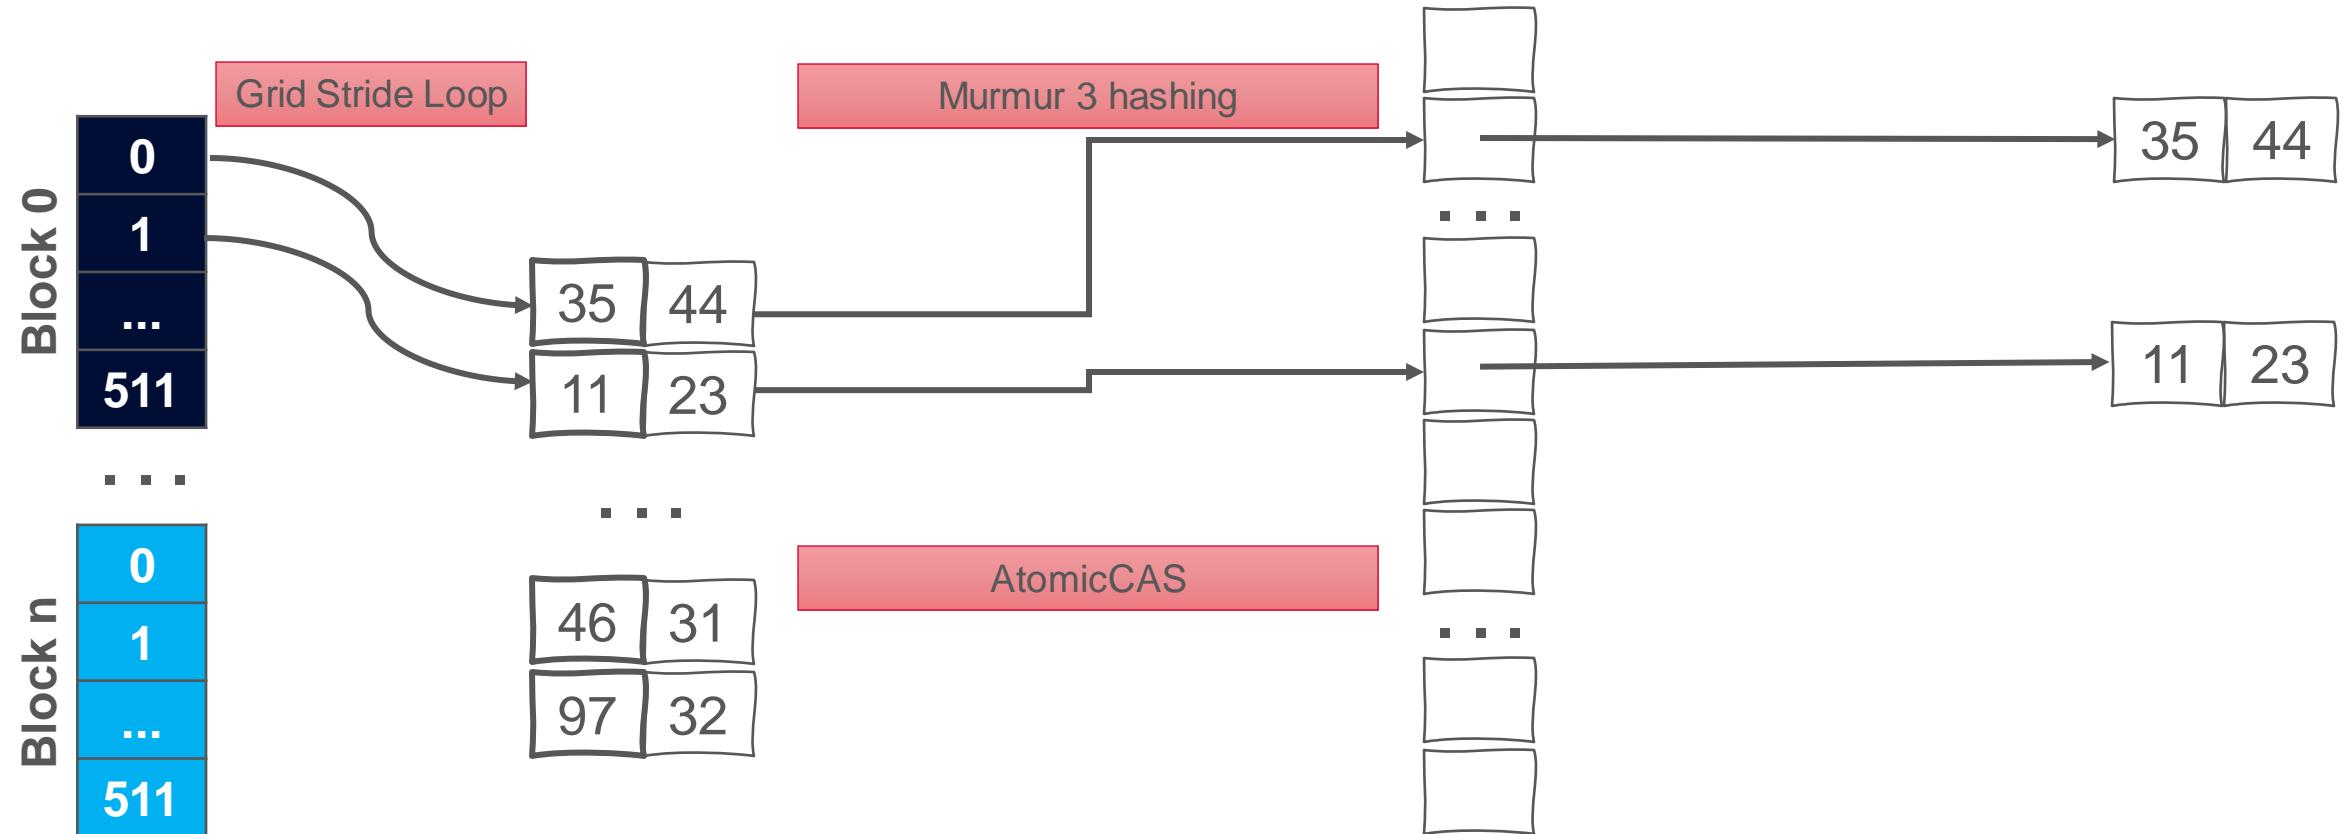
Hash Join Process



Hash Table (Open Addressing, Linear Probing)



Hash Table (Open Addressing, Linear Probing)



CUDA
Threads

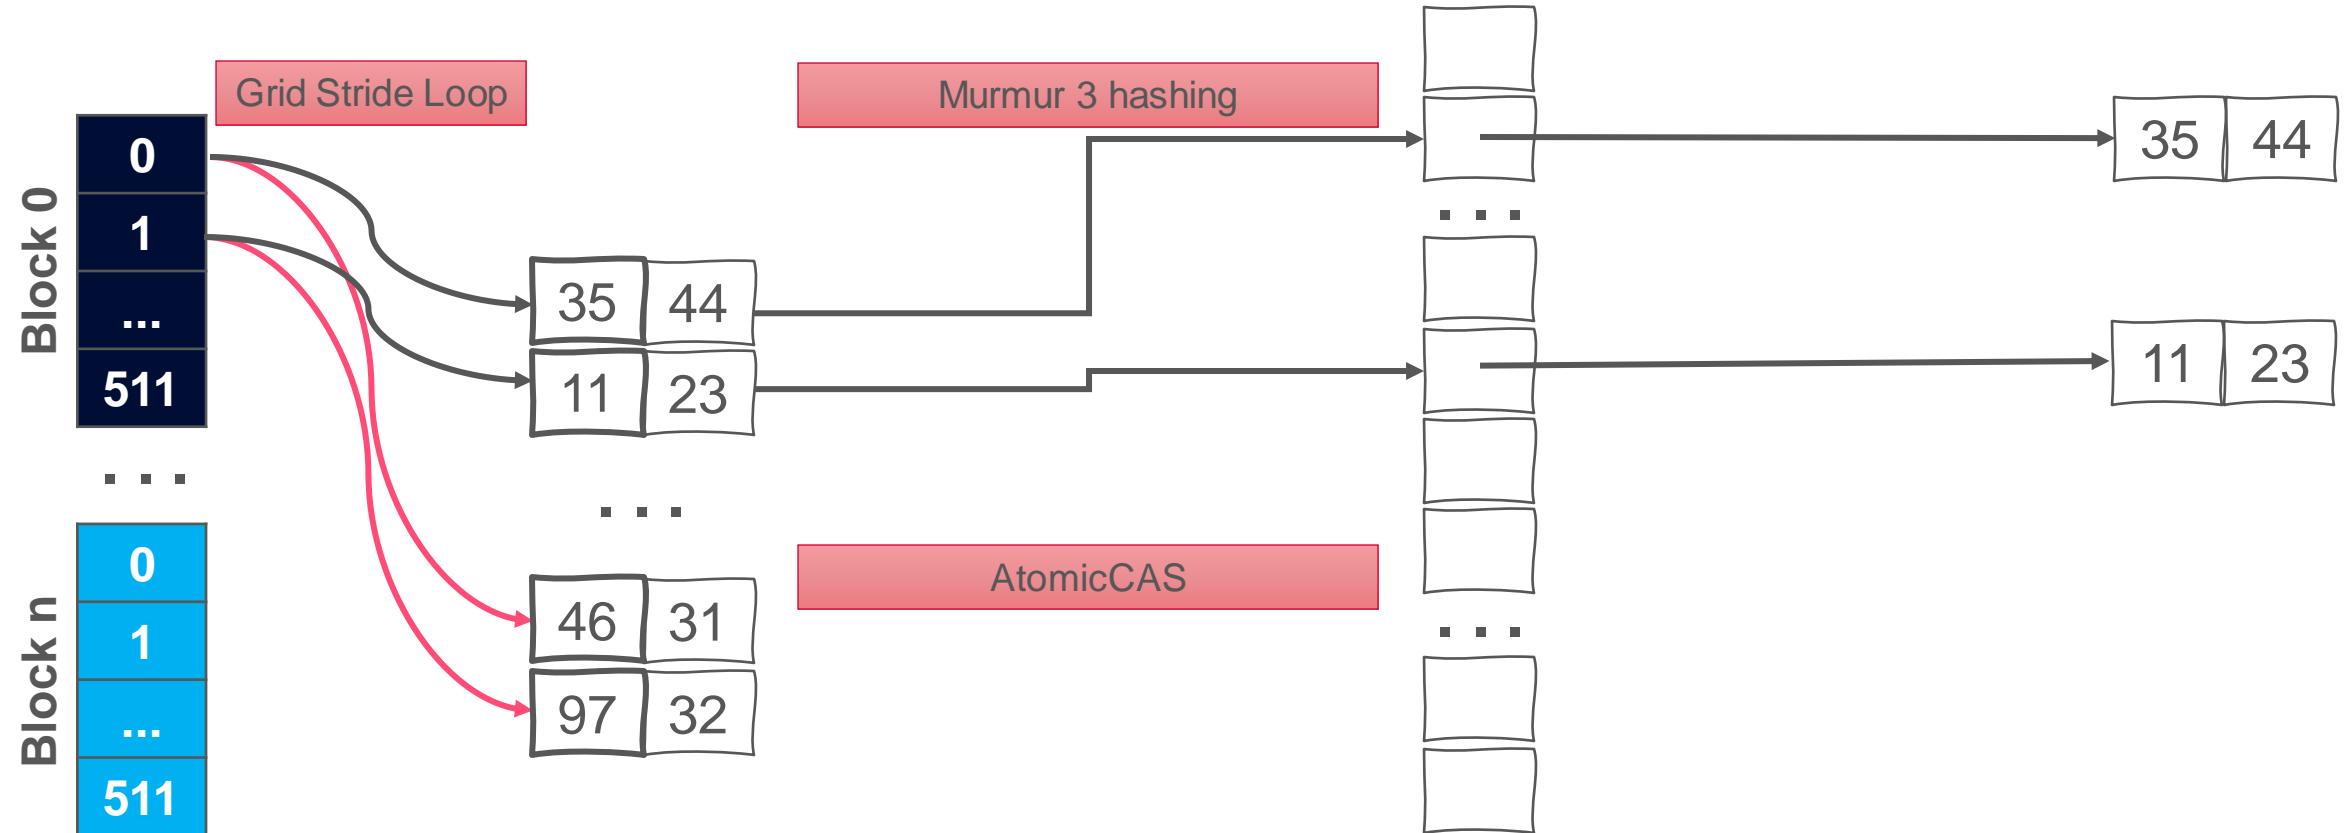
Key - Value

Hash Table

Key-Value Pair



Hash Table (Open Addressing, Linear Probing)



CUDA
Threads

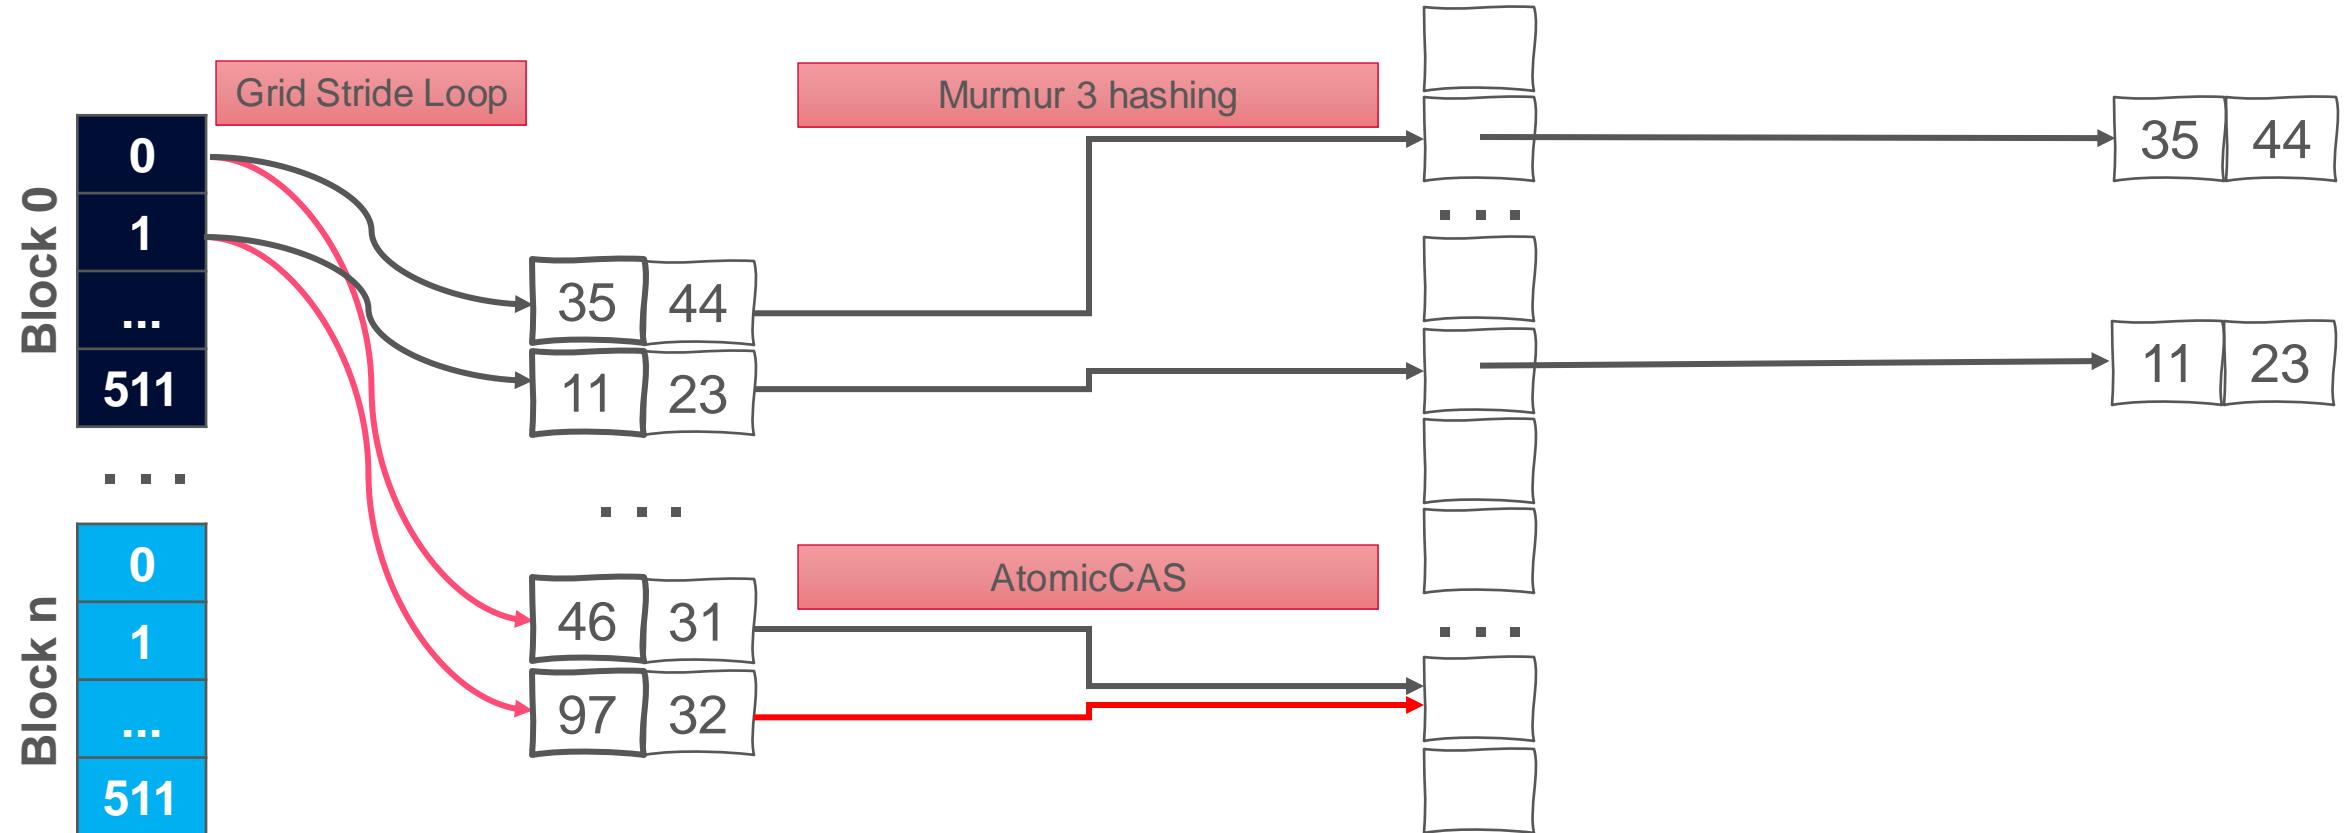
Key - Value

Hash Table

Key-Value Pair



Hash Table (Open Addressing, Linear Probing)



CUDA
Threads

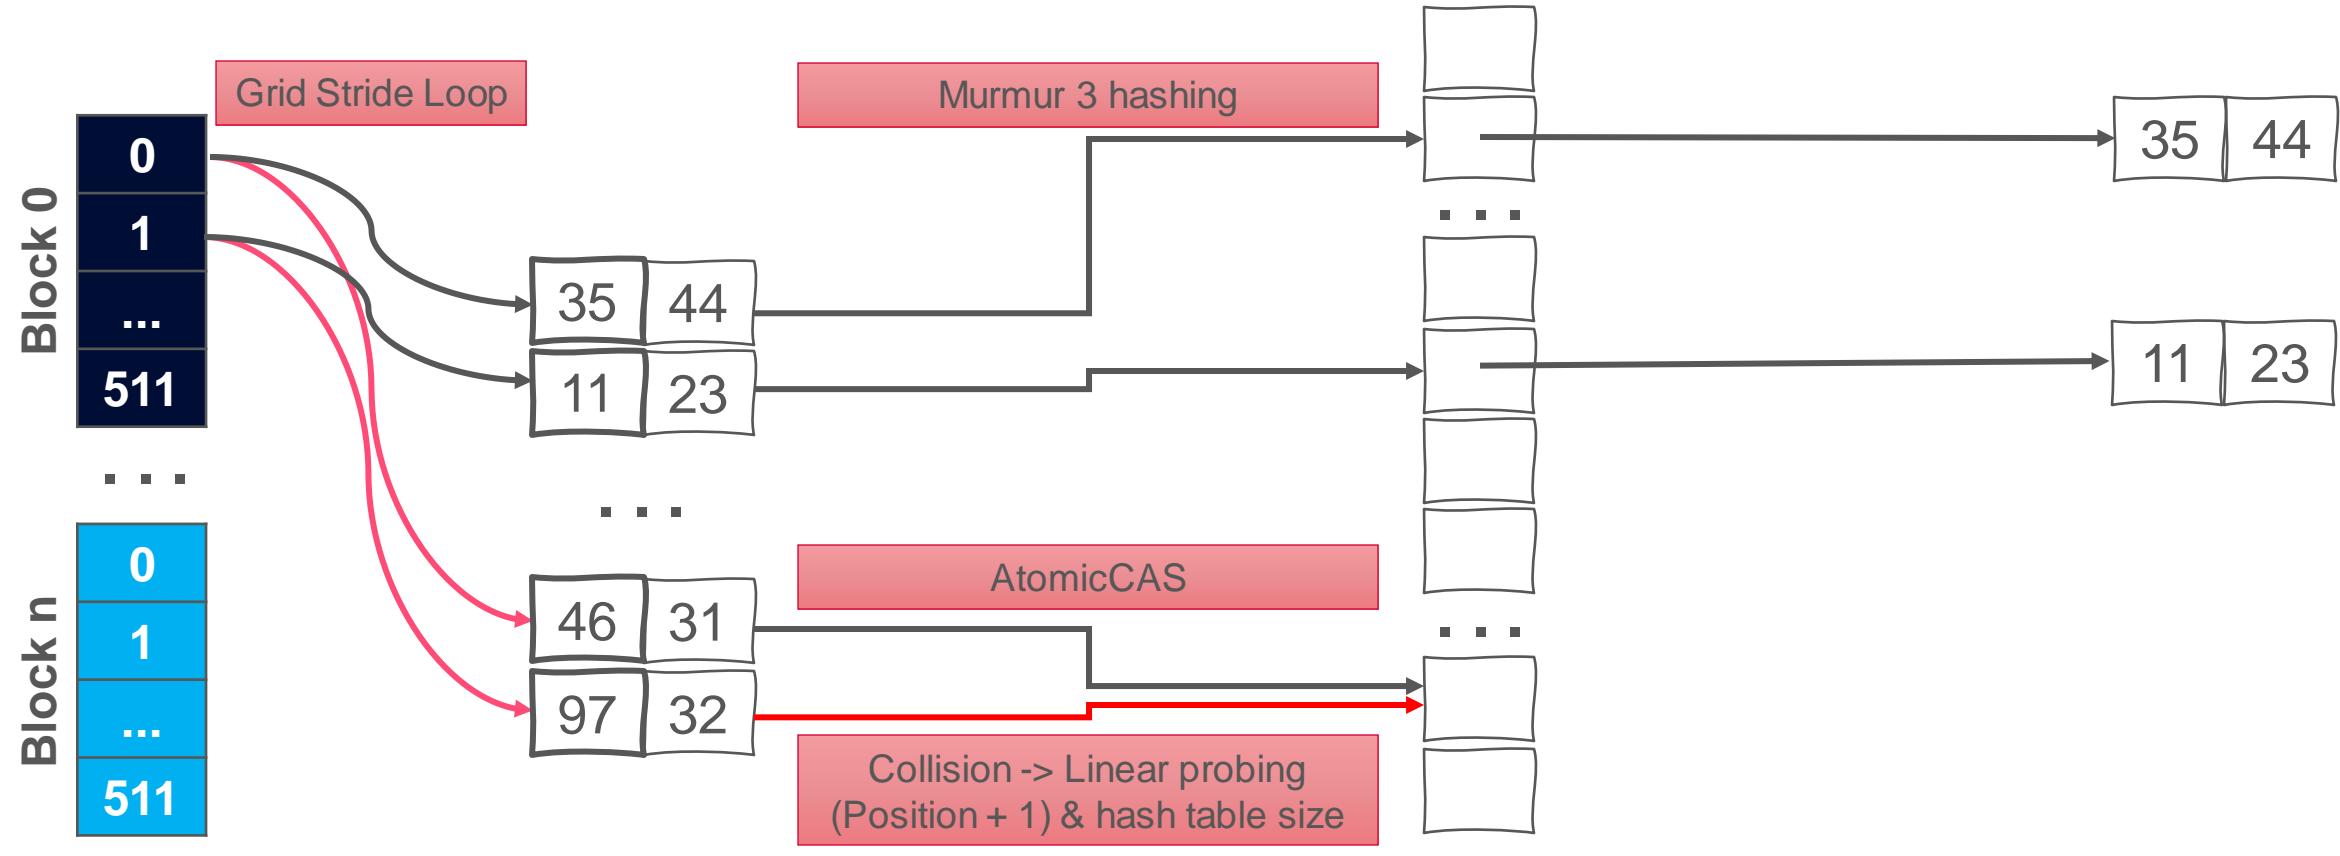
Key - Value

Hash Table

Key-Value Pair



Hash Table (Open Addressing, Linear Probing)



CUDA
Threads

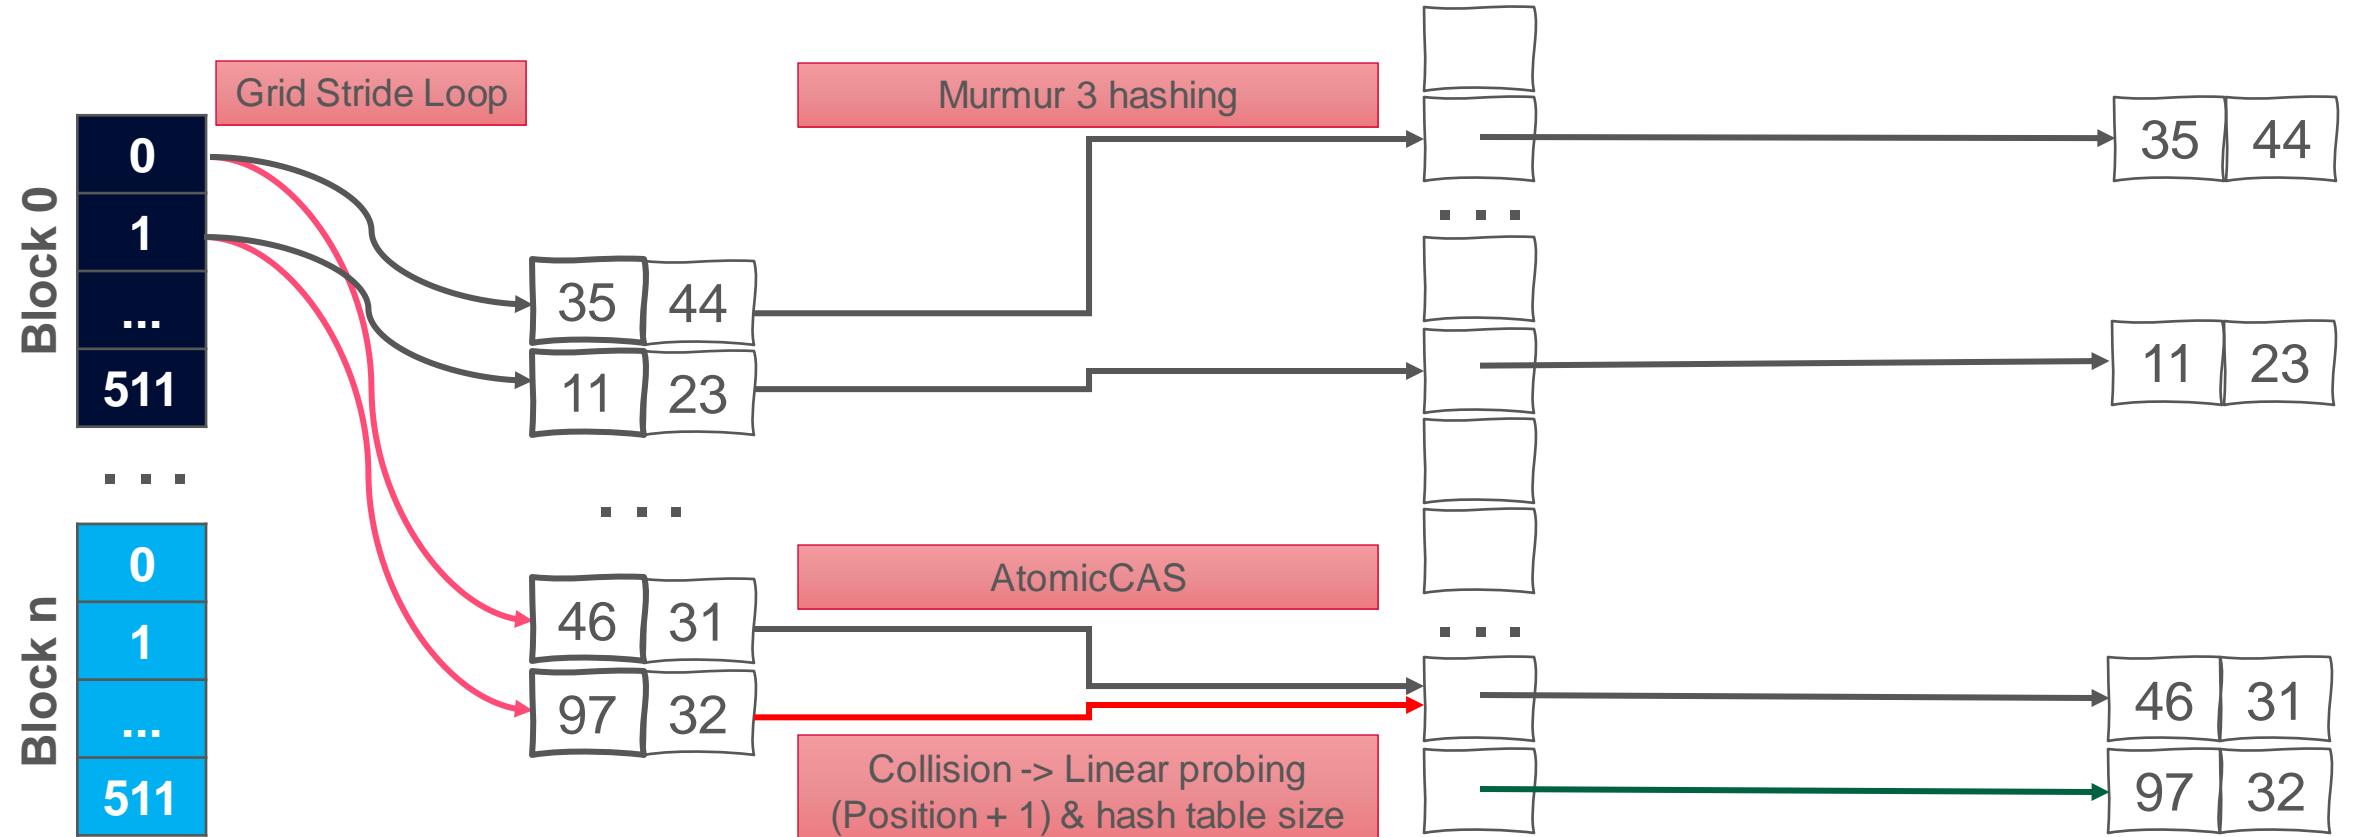
Key - Value

Hash Table

Key-Value Pair



Hash Table (Open Addressing, Linear Probing)



CUDA
Threads

Key - Value

Hash Table

Key-Value Pair



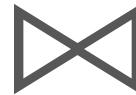
Performing Hash Join on GPU

Static Hash Table

Key	Value

Inner Relation

Key	Value



Calculate join size

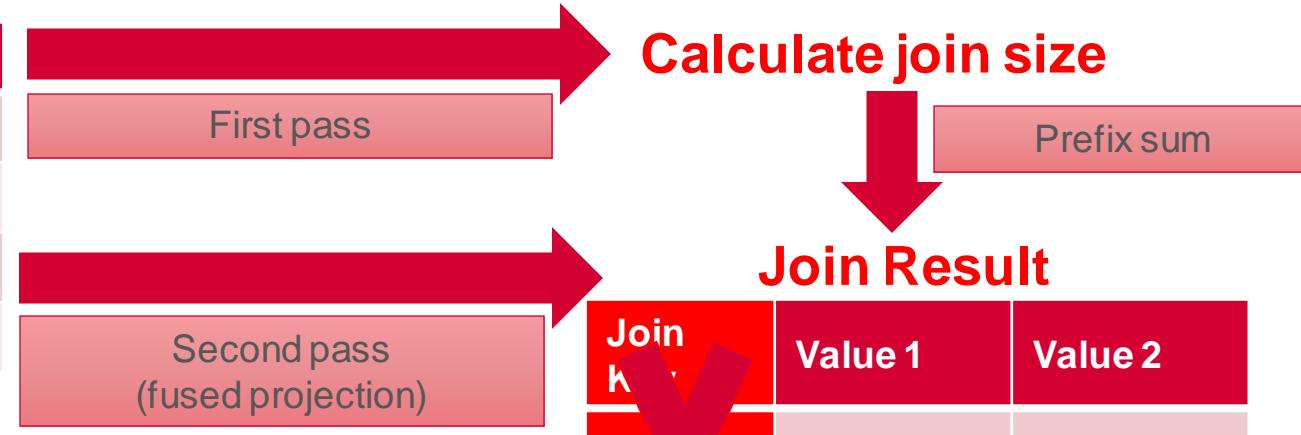
Performing Hash Join on GPU

Static Hash Table

Key	Value

Inner Relation

Key	Value



Performing Hash Join on GPU

Static Hash Table

Key	Value

Inner Relation

Key	Value



Calculate join size

First pass

Prefix sum

Join Result

Second pass
(fused projection)

Join key	Value 1	Value 2

Sort and Unique

Deduplicated Join Result

Key	Value

CUDA Advantages over cuDF

Fuse operations

Thread-block configuration

Memory management

Optimize data structure

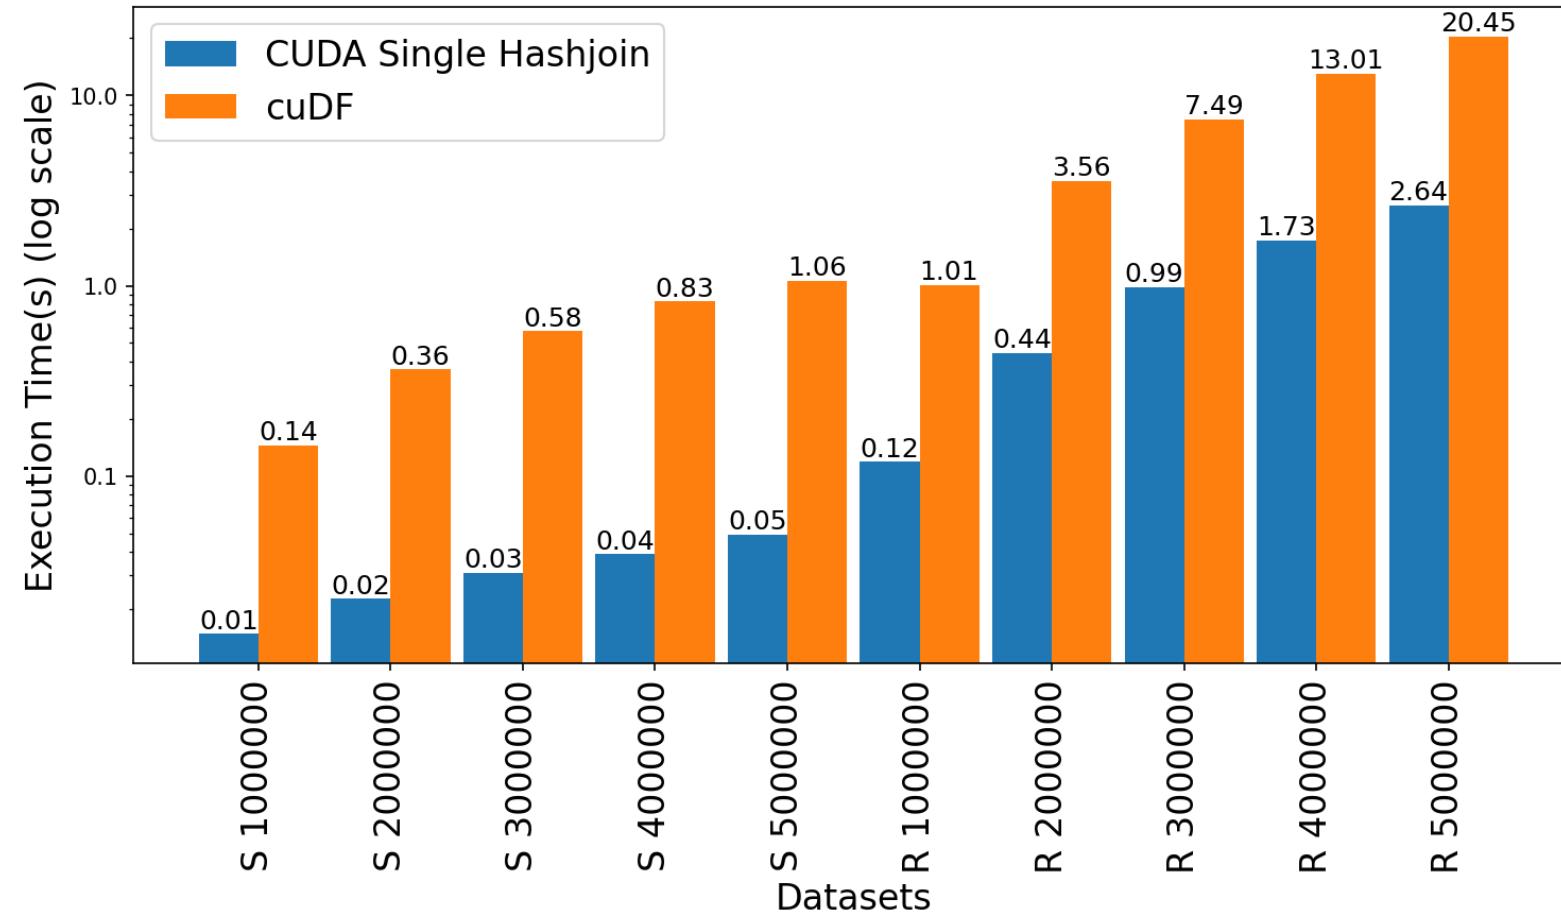
Hash Table Performance

Build rate:

- Random synthetic graph: 400 million keys/second
- String graph: 4 billion keys/second

Load factors are varied to ensure less memory overhead

Join Performance Comparison: CUDA vs cuDF



Transitive Closure: CUDA vs Soufflé vs cuDF

Dataset	Type	Rows	TC size	Iterations	CUDA Hashjoin(s)	Soufflé(s)	cuDF(s)
fe_ocean	U	409,593	1,669,750,513	247	138.237	536.233	Out of Memory
p2p-Gnutella31	D	147,892	884,179,859	31	Out of Memory	128.917	Out of memory
usroads	U	165,435	871,365,688	606	364.554	222.761	Out of Memory
fe_body	U	163,734	156,120,489	188	47.758	29.07	Out of Memory
loc-Brightkite	U	214,078	138,269,412	24	15.88	29.184	Out of Memory
SF.cedge	U	223,001	80,498,014	287	11.274	17.073	64.417
fe_sphere	U	49,152	78,557,912	188	13.159	20.008	80.077
CA-HepTh	D	51,971	74,619,885	18	4.318	15.206	26.115
p2p-Gnutella04	D	39,994	47,059,527	26	2.092	7.537	14.005
p2p-Gnutella09	D	26,013	21,402,960	20	0.72	3.094	3.906
wiki-Vote	D	103,689	11,947,132	10	1.137	3.172	6.841
cti	U	48,232	6,859,653	53	0.295	1.496	3.181
delaunay_n16	U	196,575	6,137,959	101	1.137	1.612	5.596
luxembourg_osm	U	119,666	5,022,084	426	1.322	2.548	8.194
ego-Facebook	U	88,234	2,508,102	17	0.544	0.606	3.719
cal.cedge	U	21,693	501,755	195	0.489	0.455	2.756
TG.cedge	U	23,874	481,121	58	0.198	0.219	0.857
wing	U	121,544	329,438	11	0.085	0.193	0.905
OL.cedge	U	7,035	146,120	64	0.148	0.181	0.523

Limitations



Limited to a single GPU that dictates scaling by available VRAM on the GPU

Memory overflow error for larger graphs

Open addressing based hash table causes memory overhead

Table of Contents

Exascale Computing
Datalog
Iterative Relational Algebra
Parallel Relational Algebra
Our Approaches
Topological Data Analysis for High-Dimensional Data
Future Research Direction
Conclusion

Topological Data Analysis



TDA analyzes high-dimensional data sets using algebraic topology



Main tool: persistent homology, adapts homology to point cloud



Extracts topological features and visualizes in lower-dimensional space



Analyzes complex systems: social networks, brain connectivity, internet

PH for High-Dimensional Data Analysis

Verify

PH of rs-fMRI can compactly represent topological features of FCNs

Demonstrate

FCN networks are statistically similar despite variations in TRs

Remove

Non-neural variability in multi-site case-control studies

- Ashley Suh, Mustafa Hajij, Bei Wang, Carlos Scheidegger, and Paul Rosen. Persistent homology guided force-directed graph layouts. *IEEE Transactions on Visualization and Computer Graphics*, 26(1):697–707, January 2020
- Tananun Songdechakriwut, Li Shen, and Moo Chung. Topological learning and its application to multimodal brain network integration. In *Medical Image Computing and Computer Assisted Intervention—MICCAI 2021: 24th International Conference*, Strasbourg, France, September 27–October 1, 2021, Proceedings, Part II 24, pages 166–176. Springer, 2021.

Table of Contents

Exascale Computing
Datalog
Iterative Relational Algebra
Parallel Relational Algebra
Our Approaches
Topological Data Analysis for High-Dimensional Data
Future Research Direction
Conclusion

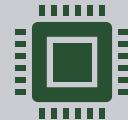
Future Research Direction

- Develop a multi-node multi-GPU backend for Datalog
- Design GPU benchmarking techniques for iterative RA
- Explore PH-based techniques on multi-site temporal brain FCNs
- Develop declarative analytics technique for high-dimensional datasets

Table of Contents

Exascale Computing
Datalog
Iterative Relational Algebra
Parallel Relational Algebra
Our Approaches
Topological Data Analysis for High-Dimensional Data
Future Research Direction
Conclusion

Conclusion



Reviewed Datalog applications on heterogeneous systems



Explored topological data analysis techniques for brain network



Identified research direction for declarative analytics on high-dimensional data

Publications

- Ahmedur Rahman Shovon, Landon Richard Dyken, Oded Green, Thomas Gilray, and Sidharth Kumar. Accelerating datalog applications with cudf. In 2022 IEEE/ACM 11th Workshop on Irregular Applications: Architectures and Algorithms (IA3), pages 41–45, 2022
- Ahmedur Rahman Shovon, Thomas Gilray, Kristopher Micinski, and Sidharth Kumar. Towards Iterative Relational Algebra on the GPU. In 2023 USENIX Annual Technical Conference (USENIX ATC 23), pp. 1009-1016.
- Yihao Sun, Ahmedur Rahman Shovon, Thomas Gilray, Kristopher Micinski, Sidharth Kumar. Modern Datalog on the GPU, In review.
- Sewell, A., Fan, K., Shovon, A. R., Dyken, L., Kumar, S., & Petruzza, S. (2024, January). Bruck Algorithm Performance Analysis for Multi-GPU All-to-All Communication. In Proceedings of the International Conference on High Performance Computing in Asia-Pacific Region (pp. 127-133).
- Sidharth Kumar, Ahmedur Rahman Shovon, and Gopikrishna Deshpande. The invariance of persistent homology of brain networks to data acquisition-related non-neural variability in resting state fmri. Human Brain Mapping, 44(13), 4637–4651
- Darshan Shimoga Chandrashekhar, Santhosh Kumar Karthikeyan, Praveen Kumar Korla, Henalben Patel, Ahmedur Rahman Shovon, Mohammad Athar, George J Netto, Zhaojun Qin, Sidharth Kumar, Upender Manne, et al. Ualcan: An update to the integrated cancer data analysis platform. Neoplasia, 25:18–27, 2022.
- Kashyap Balakavi, Rushitha Janga, Ahmedur Rahman Shovon, Don Dempsey, Elliot Lefkowitz, Sidarth Kumar. Scalable, interactive and hierarchical visualization of virus taxonomic data. In 2023 IEEE Workshop on Visual Analytics in Healthcare (VAHC), in conjunction with IEEE VIS 2023.

A grayscale photograph of the Chicago skyline, featuring the Willis Tower (formerly Sears Tower) as the central focus. In the foreground, there's a view of a university campus with modern buildings, trees, and students walking on paths.

Thank You

ashov@uic.edu





Appendix

Parallel Join: Algorithms

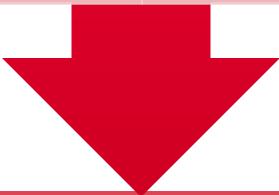
Popular algorithms

Non-Indexed Nested Loop Join (NINLJ)

Indexed Nested Loop Join (INLJ)

Sort-Merge Join (SMJ)

Radix Hash Join (HJ)



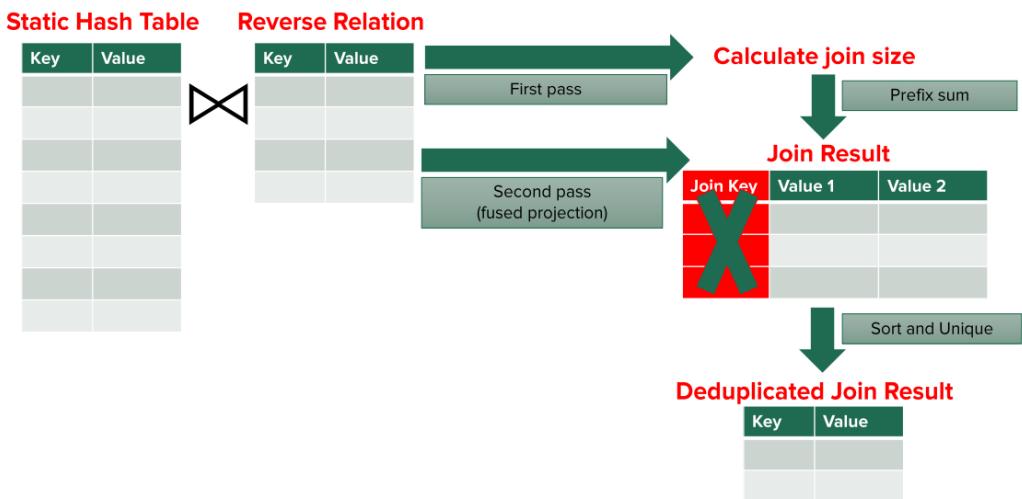
Sort-merge is suitable for small to medium-sized tables,
radix hash join is suitable for large tables

- Chengxin Guo, Hong Chen, Feng Zhang, and Cuiping Li. Parallel hybrid join algorithm on gpu. 2019 IEEE 21st International Conference on High Performance Computing and Communications; IEEE 17th International Conference on Smart City; IEEE 5th International Conference on Data Science and Systems (HPCC/SmartCity/DSS), pages 1572–1579, 2019.
- Hongzhi Wang, Ning Li, Zheng ke Wang, and Jianing Li. Gpu-based efficient join algorithms on hadoop. The Journal of Supercomputing, 77:292–321, 2020.

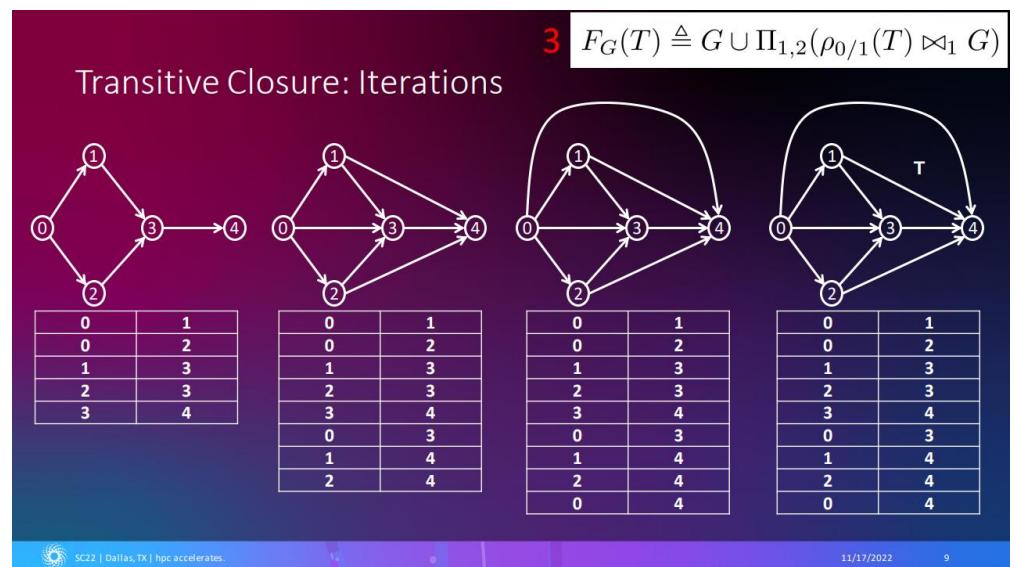
Publications

Shovon, A. R., Gilray, T., Micinski, K., & Kumar, S. (2023). Towards iterative relational algebra on the {GPU}. In 2023 USENIX Annual Technical Conference (USENIX ATC 23) (pp. 1009-1016).

Performing Hash Join on GPU



Shovon, A. R., Dyken, L. R., Green, O., Gilray, T., & Kumar, S. (2022, November). Accelerating Datalog applications with cuDF. In 2022 IEEE/ACM Workshop on Irregular Applications: Architectures and Algorithms (IA3) (pp. 41-45). IEEE.

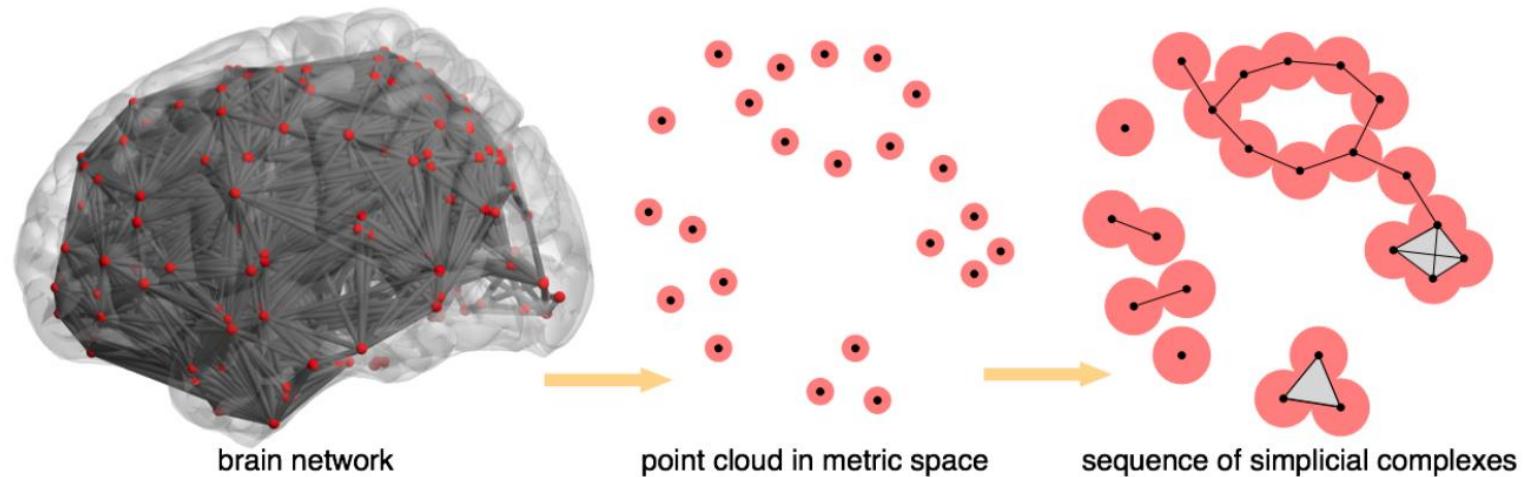


Preserve Brain Topology

Topology of brain networks should be **preserved** irrespective of sampling density

Metrics capturing topology may be **statistically similar** across sampling periods, reducing non-neural variability

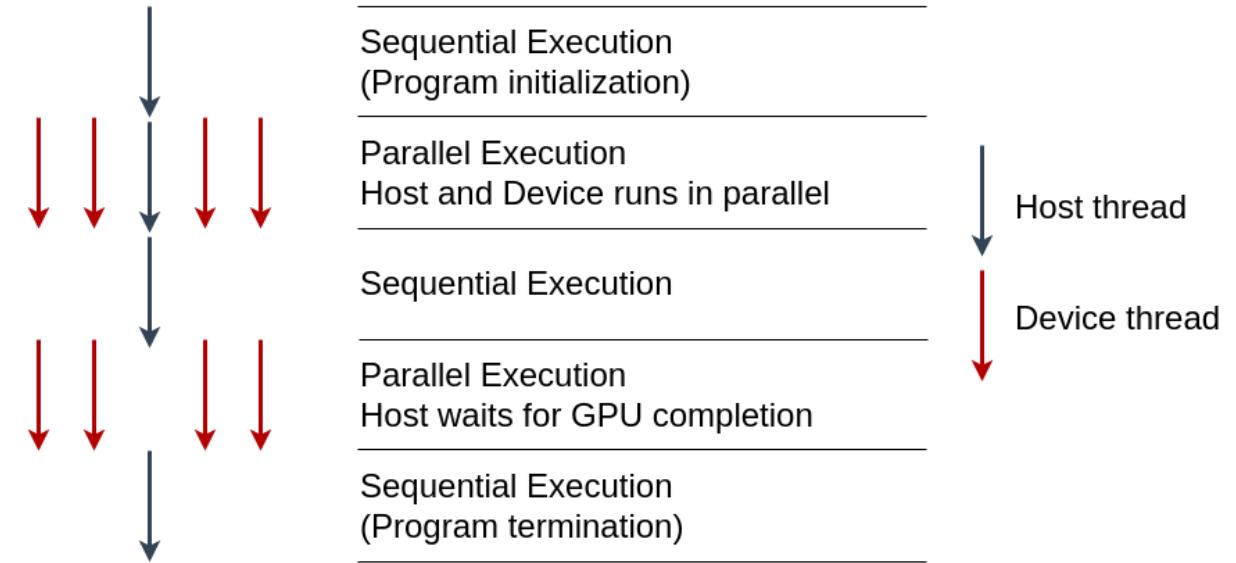
PH can be **robust** to changes in data acquisition parameters like sampling period



- A. Patania, F. Vaccarino, and G. Petri, "Topological analysis of data," EPJ Data Science, vol. 6, pp. 1–6, 2017.
- Bourakna, A.E.Y., Chung, M.K., Ombao, H. 2022 Topological data analysis for multivariate time series data arXiv:2204.13799

CUDA Programming Model

- Globally Sequential Locally Parallel programming pattern
- Invokes parallel **kernels** that execute across a set of threads
- CUDA spawns the threads from a hierarchy of **grid** (3D) and **block** (3D)
- Each thread executes an instance of the kernel
- Each thread has thread ID, program counter, registers, and per-thread private memory
- Supports **C/C++, Fortan, Python**



- Jason. Sanders. CUDA by example : an introduction to general-purpose GPU programming. Addison-Wesley, Upper Saddle River, NJ, 2011.
- Jianbin Fang, Chun Huang, Tao Tang, and Zheng Wang. Parallel programming models for heterogeneous many-cores: a comprehensive survey. CCF Transactions on High Performance Computing, 2(4):382–400, 2020.

GPGPU Challenges

Complex memory hierarchies

Different architectures from multicore CPUs

Convert sequential programming to parallel

- Gerassimos Barlas. Chapter 6 - gpu programming. In Gerassimos Barlas, editor, Multicore and GPU Programming, pages 391–526. Morgan Kaufmann, Boston, 2015.
- J. D. Owens, M. Houston, D. Luebke, S. Green, J. E. Stone and J. C. Phillips, "GPU Computing," in Proceedings of the IEEE, vol. 96, no. 5, pp. 879-899, May 2008, doi: 10.1109/JPROC.2008.917757.

Join on GPUs: Control Flow Analysis (CFA)



Parallel functional CFA encoded in Datalog utilizes RA as the foundation on GPU

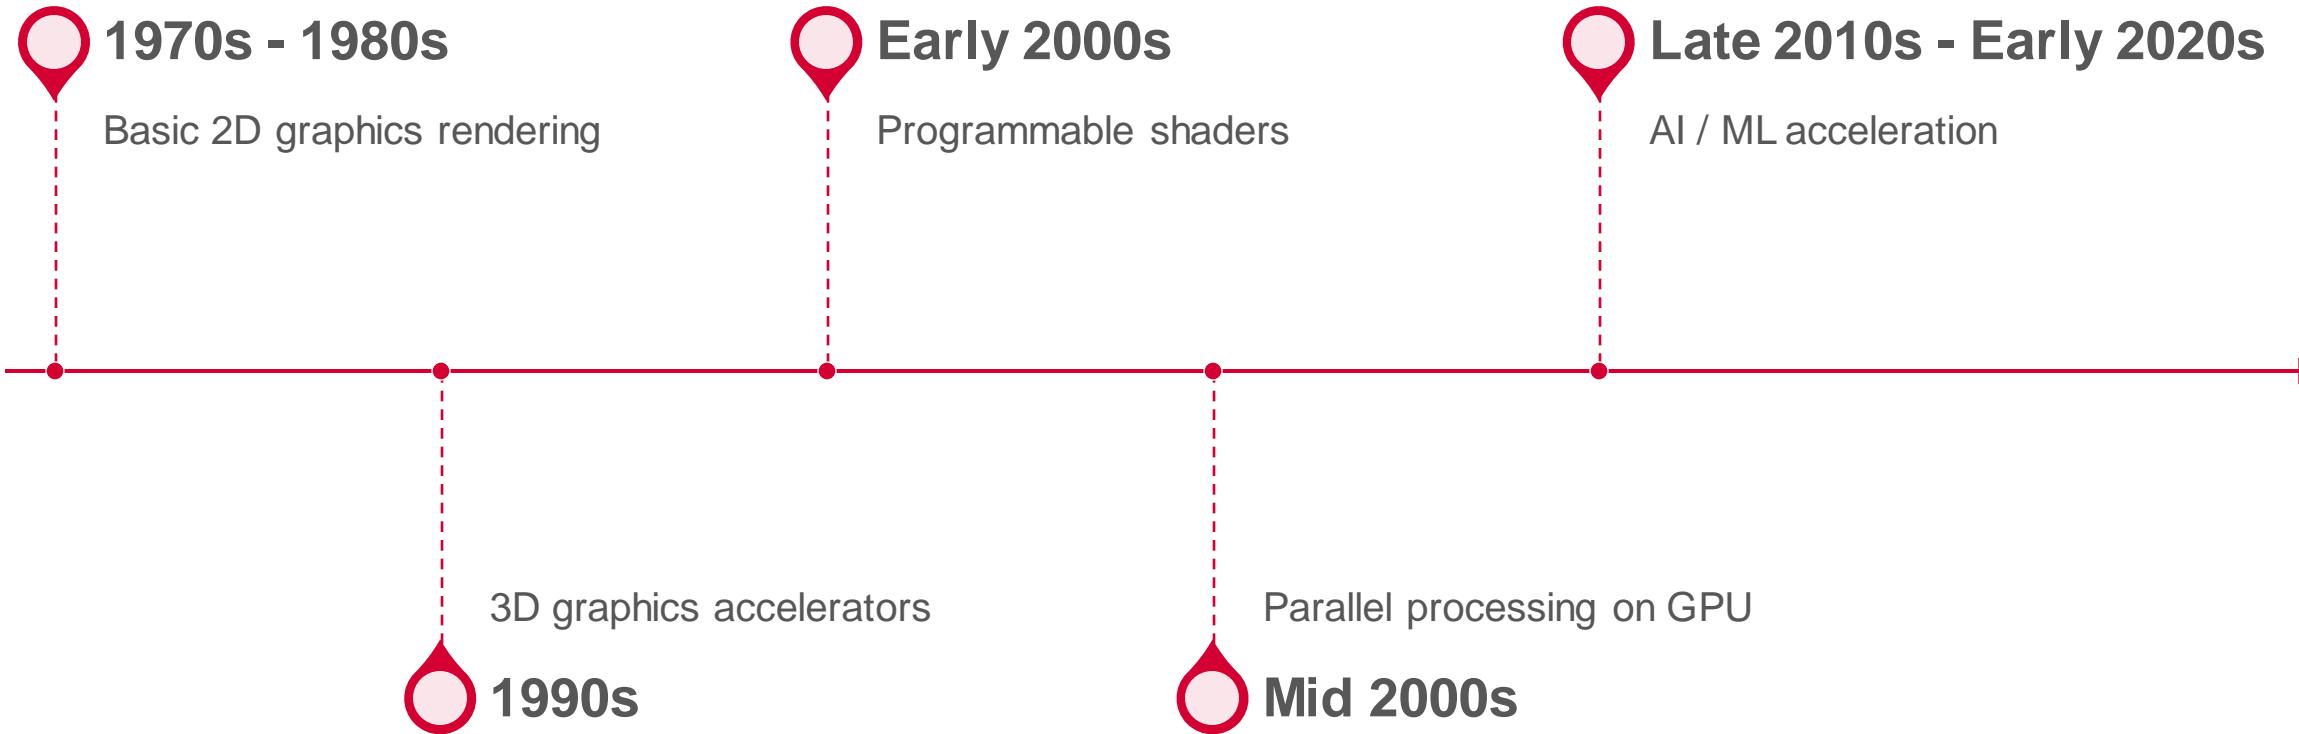


Applied GPU parallelism with multi-node multi-core HPC



Proposed partitioned global address space (PGAS) programming model

Advancements in GPU



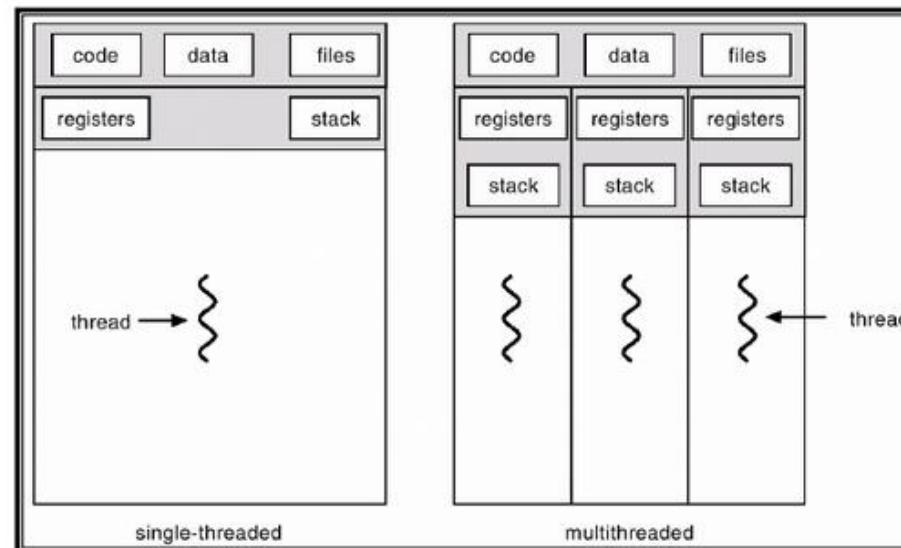
CPU vs GPU (A Sample Machine)

CPU

- 13th Gen Intel® Core™ i9-13900H
- **14** cores (6 P-cores + 8 E-cores)
- Total **20** threads
- Suitable for **serial** workloads
- Access to system memory (RAM)

GPU

- NVIDIA GeForce RTX 4070, 8 GB GDDR6
- **5888** CUDA cores
- Total **94,208** threads
- Suitable for **parallel** workloads
- Access to dedicated VRAM



Recap to Relational Data

- **Relation:** 2-dimensional structure
- **Attribute:** Represents characteristics
- **Row:** Represents unique record
- **Join (\bowtie):** Combines data from relations
- **Projection (Π):** Select specific columns

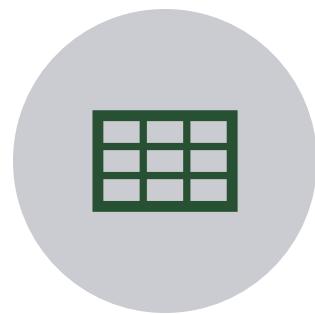
Relation

UserID	UserName	UserEmail
101	Alice	alice@example.com
102	Bob	bob@example.com

Attribute (Column)

Tuple (Row)

Why Join is Important?



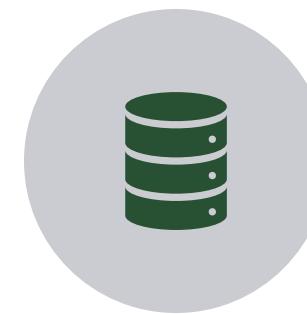
COMBINE DATA
FROM MULTIPLE
TABLES



FIND PATTERNS IN
DATA



CLEAN DATA



CREATE NEW DATA
SETS

Example of Natural Join

User

UserID	UserName	UserEmail
101	Alice	alice@example.com
102	Bob	bob@example.com
103	Eve	eve@example.com



Order

UserID	OrderTotal	Items
101	25.69	2
102	145.66	3
103	12.11	1
103	44.00	2

Example of Natural Join

User

UserID	UserName	UserEmail
101	Alice	alice@example.com
102	Bob	bob@example.com
103	Eve	eve@example.com

User



Order



Order

UserID	OrderTotal	Items
101	25.69	2
102	145.66	3
103	12.11	1
103	44.00	2

UserID	UserName	UserEmail	OrderTotal	Items
101	Alice	alice@example.com	25.69	2

Example of Natural Join

User

UserID	UserName	UserEmail
101	Alice	alice@example.com
102	Bob	bob@example.com
103	Eve	eve@example.com



Order

UserID	OrderTotal	Items
101	25.69	2
102	145.66	3
103	12.11	1
103	44.00	2

User



Order

UserID	UserName	UserEmail	OrderTotal	Items
101	Alice	alice@example.com	25.69	2
102	Bob	bob@example.com	145.66	3

Example of Natural Join

User

UserID	UserName	UserEmail
101	Alice	alice@example.com
102	Bob	bob@example.com
103	Eve	eve@example.com



Order

UserID	OrderTotal	Items
101	25.69	2
102	145.66	3
103	12.11	1
103	44.00	2

User



Order

UserID	UserName	UserEmail	OrderTotal	Items
101	Alice	alice@example.com	25.69	2
102	Bob	bob@example.com	145.66	3
103	Eve	eve@example.com	12.11	1

Example of Natural Join

User (Outer Relation)

UserID	UserName	UserEmail
101	Alice	alice@example.com
102	Bob	bob@example.com
103	Eve	eve@example.com



Order (Inner Relation)

UserID	OrderTotal	Items
101	25.69	2
102	145.66	3
103	12.11	1
103	44.00	2



UserID	UserName	UserEmail	OrderTotal	Items
101	Alice	alice@example.com	25.69	2
102	Bob	bob@example.com	145.66	3
103	Eve	eve@example.com	12.11	1
103	Eve	eve@example.com	44.00	2

Example of Natural Join

User (Outer Relation)

UserID	UserName	UserEmail
101	Alice	alice@example.com
102	Bob	bob@example.com
103	Eve	eve@example.com



Order (Inner Relation)

UserID	OrderTotal	Items
101	25.69	2
102	145.66	3
103	12.11	1
103	44.00	2

UserID	UserName	UserEmail	OrderTotal	Items
101	Alice	alice@example.com	25.69	2
102	Bob	bob@example.com	145.66	3
103	Eve	eve@example.com	12.11	1
103	Eve	eve@example.com	44.00	2

Duplicates on Join Result

User \bowtie Order

UserID	UserName	UserEmail	OrderTotal	Items
101	Alice	alice@example.com	25.69	2
102	Bob	bob@example.com	145.66	3
103	Eve	eve@example.com	12.11	1
103	Eve	eve@example.com	44.00	2

$\Pi_{(UserName, UserEmail)}(User \bowtie Order)$

UserName	UserEmail
Alice	alice@example.com
Bob	bob@example.com
Eve	eve@example.com
Eve	eve@example.com

Hybrid Join Algorithm

Guo et al. proposed PHYJ: SMJ with HJ join

Reduced host-to-device and device-to-host

Fused data communication with GPU execution

On a single GPU achieved up to **1.72X** speedup

Can handle skewed data

No information on multiple GPUs or distributed systems

- Chengxin Guo, Hong Chen, Feng Zhang, and Cuiping Li. Parallel hybrid join algorithm on gpu. 2019 IEEE 21st International Conference on High Performance Computing and Communications; IEEE 17th International Conference on Smart City; IEEE 5th International Conference on Data Science and Systems (HPCC/SmartCity/DSS), pages 1572–1579, 2019.
- Hongzhi Wang, Ning Li, Zheng ke Wang, and Jianing Li. Gpu-based efficient join algorithms on hadoop. *The Journal of Supercomputing*, 77:292–321, 2020.

Dataset: Stanford large network dataset collection

Graph	Type	Nodes	Edges
p2p-Gnutella09	Directed	8,114	26,013
p2p-Gnutella04	Directed	10,876	39,994
Skitter	Undirected	1,696,415	11,095,298
roadNet-CA	Undirected	1,965,206	2,766,607
roadNet-TX	Undirected	1,379,917	1,921,660
roadNet-PA	Undirected	1,088,092	1,541,898
SF.cedge	Undirected	1,74,955	2,23,001
cal.cedge	Undirected	21,048	21,693
TG.cedge	Undirected	18,263	23,874
OL.cedge	Undirected	6,105	7,035

- Leskovec, J., & Krevl, A. (2014). SNAP Datasets:Stanford large network dataset collection.

Join on CPUs

9
9

Processing large amounts of data is a challenging task

Radix hash joins and sort-merge joins have been widely discussed and compared in previous studies

Modern HPC architectures have led to advancements in distributed join algorithms

Most of the multi-core implementations use Message Passing Interface (MPI)

Fundamental part of fixed-point iterative algorithm such as Transitive Closure (TC)

Join on GPUs

GPUs have high memory bandwidth and hierarchy for faster data access

GPUs offer parallel architecture for faster join operation

Data divided into smaller subsets for parallel processing

GPU-based join operation improves performance and scalability

Join on GPUs: LogiQL

Wu et al. presents **Red Fox** high-performance accelerator core for **LogiQL** queries

Outperforms multi-threaded CPU-based implementations

Novel: multi-predicate join algorithm (worst-case optimal) on GPU

Issue: deduplication of tuples and maintaining join result in sorted order

- Haicheng Wu, Gregory Diamos, Tim Sheard, Molham Aref, Sean Baxter, Michael Garland, and Sudhakar Yalamanchili. Red fox: An execution environment for relational query processing on gpus. InProceedings of Annual IEEE/ACM International Symposium on Code Generation and Optimization, pages 44–54, 2014.
- Haicheng Wu. Acceleration and execution of relational queries using general purpose graphics processingunit (GPGPU). PhD thesis, Georgia Institute of Technology, 2015.

Join on GPUs: Relational Learning Framework

Expedites rule coverage on GPUs for healthcare records data

Outperforms **75%** of applications over multi-core CPU systems

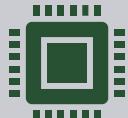
Duplicate tuples not efficiently managed and GPU memory overflows

- Carlos Alberto Martínez-Angeles, Haicheng Wu, Inês Dutra, Vítor Santos Costa, and Jorge Buenabá dChávez. Relational learning with gpus: Accelerating rule coverage. International Journal of Parallel Programming, 44(3):663–685, 2016

Parallel Join



What: Perform relational join operation simultaneously on a number of processors or machines



When: Useful when input data is enormous and the join is computationally costly



How: Divide the data into partitions and assign each partition to a different processor

Modern Implementations of Datalog

LogicBlox

Soufflé

RadLog

PRAM

SLOG

LogicBlox



Gradual advances in logic programming: Advancement over the years



LogiQL language: Datalog-like, null-free



LogicBlox system: Mature, enterprise-integrated



Massive financial benefits: \$300B in revenue for business analytics

- Molham Aref, Balder ten Cate, Todd J. Green, Benny Kimelfeld, Dan Olteanu, Emir Pasalic, Todd L. Veldhuizen, and Geoffrey Washburn. Design and implementation of the logicblox system. In Proceedings of the 2015 ACM SIGMOD International Conference on Management of Data, SIGMOD '15, page 1371–1382, New York, NY, USA, 2015. Association for Computing Machinery

Soufflé

A variant of Datalog for static analysis using OpenMP

State-of-the-art implementation for multi-core CPU systems with single-node

Translates Datalog programs to optimized C++ programs

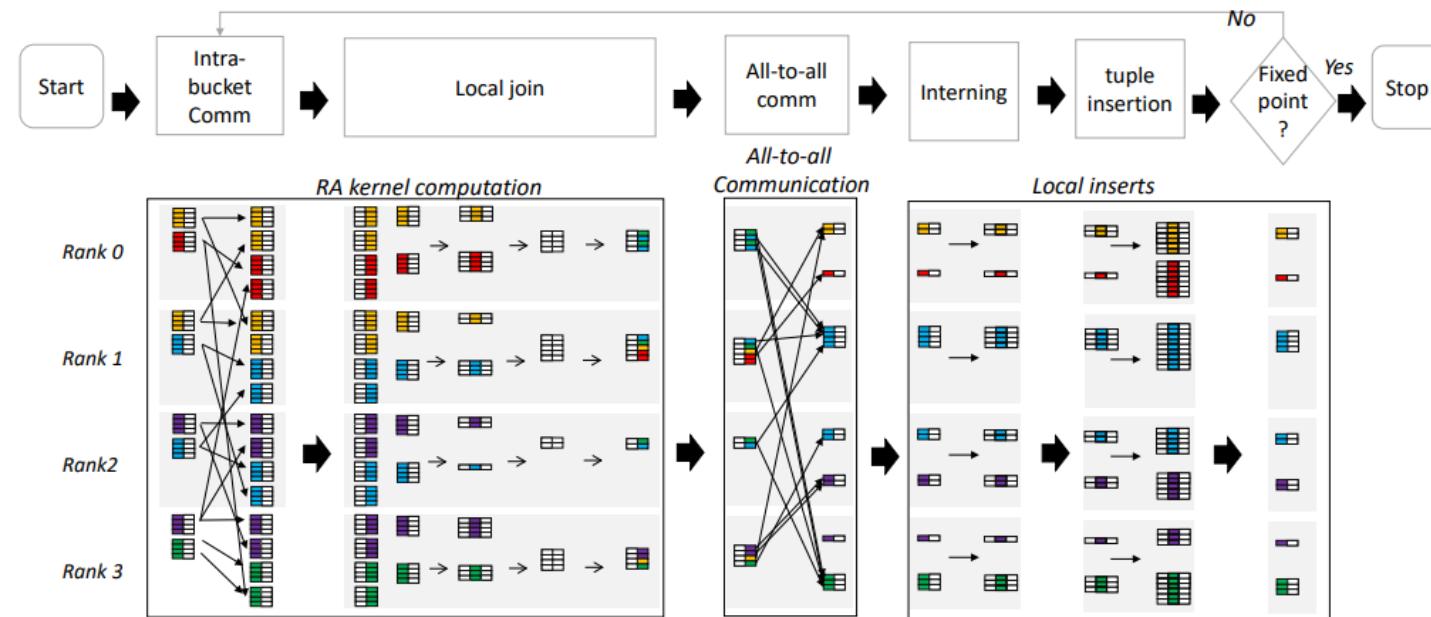
Supports limited number of threads for task-level parallelism

Cannot provide data parallelism

- Herbert Jordan, Bernhard Scholz, and Pavle Subotić. Soufflé: On synthesis of program analyzers. In International Conference on Computer Aided Verification, pages 422–430. Springer, 2016.
- Thomas Gilray, Sidharth Kumar, and Kristopher Micinski. Compiling data-parallel datalog. In Proceedings of the 30th ACM SIGPLAN International Conference on Compiler Construction, CC 2021, page 23–35, New York, NY, USA, 2021. Association for Computing Machinery.

SLOG

- High-performance implementation providing a compiler, REPL, and runtime
- Uses PRAM as relational algebra backend
- Outperforms Soufflé and RadLog for higher core counts
- Does not yet utilize GPU devices



- Thomas Gilray, Arash Sahebolamri, Sidharth Kumar, and Kristopher Micinski. Higher-order, dataparallel structured deduction. arXiv preprint arXiv:2211.11573, 2022.
- Thomas Gilray, Sidharth Kumar, and Kristopher Micinski. Compiling data-parallel datalog. In Proceedings of the 30th ACM SIGPLAN International Conference on Compiler Construction, pages 23–35, 2021.

Join on CPUs: Fault Tolerant Hash Joins

Nasibullin and Novikov developed a distributed fault-tolerant hash-join

Utilizes data replication and alive signals for fault detection and recovery

Can handle a single point of failure with a **keeper-worker** model

Does not outperform distributed hash join algorithms in a **failure-free** situation

Can be criticized for **data skewness** and lack of accelerator use

- Arsen Nasibullin and Boris Novikov. Fault tolerant distributed hash-join in relational databases. In Fifth Conference on Software Engineering and Information Management (SEIM-2020), pages 11–17, 2020.

Datalog on GPUs

Outperforms non-GPU systems for join and transitive closure computation

Use GPU kernels to eliminate unnecessary tuples, join, and project

Achieved up to 40x performance enhancement over single CPU system

Limitation in relation size and applicability to distributed platforms

- Carlos Alberto Martínez-Angeles, Inês Dutra, Vítor Santos Costa, and Jorge Buenabad-Chavez. A datalog engine for gpus. In Declarative Programming and Knowledge Management: DeclarativeProgramming Days, KDPD 2013, Unifying INAP, WFLP, and WLP, Kiel, Germany, September 11-13,2013, Revised Selected Papers 0, pages 152–168. Springer, 2014.
- Grigoris Antoniou, Sotiris Batsakis, Raghava Mutharaju, Jeff Z. Pan, Guilin Qi, Ilias Tachmazidis, Jacopo Urbani, and Zhangquan Zhou. A survey of large-scale reasoning on the web of data. TheKnowledge Engineering Review, 33, 2018.

Persistent Homology (PH)

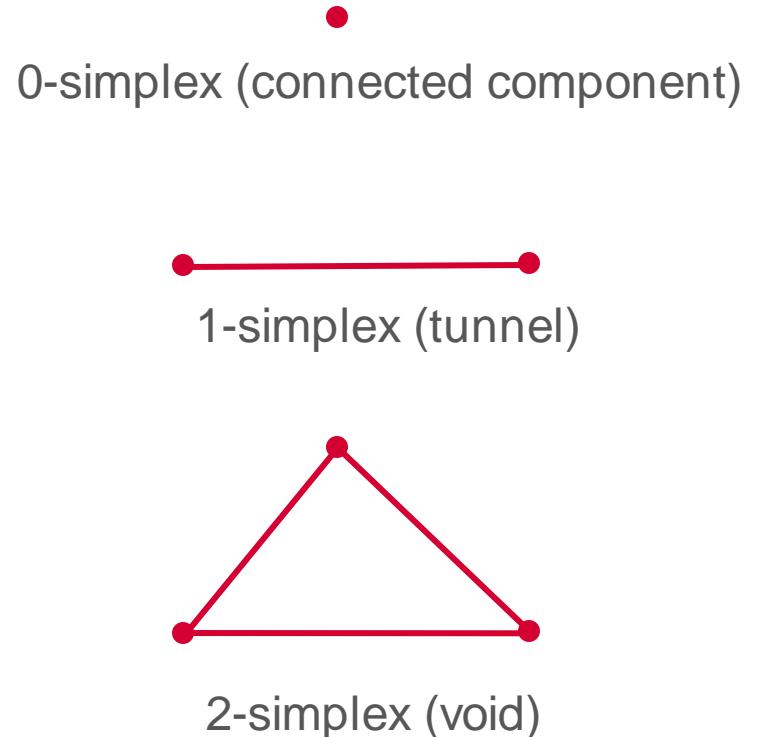
1
1
0

Investigates connections between different parts of complex networks using **simplicial complexes**

Simplicial complexes represent topological spaces as collections of simplices (points, line segments, triangles, etc.)

Detects **topological features** and measures the significance of relationships across multiple **thresholds**

Topological features refer to the **0-, 1-, 2-** dimensional homology groups of a metric space describing **connected components**, **tunnels**, and **voids**



- M. Rubinov and O. Sporns, "Complex network measures of brain connectivity: uses and interpretations," *Neuroimage*, vol. 52, no. 3, pp. 1059–1069, 2010.
- Aktas, M.E., Akbas, E. & Fatmaoui, A.E. Persistence homology of networks: methods and applications. *Appl Netw Sci* 4, 61 (2019). <https://doi.org/10.1007/s41109-019-0179-3>

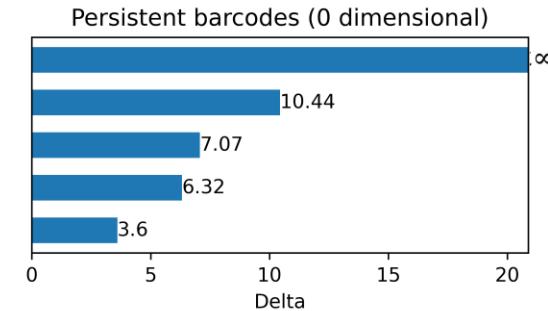
Barcodes and Persistent Diagram

PH can track when features are **created and destroyed** with varying scales

Quantifies **birth and death of features** in the graph according to their significance

Represented in the form of barcodes, a **fingerprint** for a graph

Persistent diagram: points in 2D plane obtained from barcodes, represent topological feature persistence with **x** as birth time and **y** as death time



Adjacency matrix

0.00	8.54	12.36	7.07	10.44
8.54	0.00	20.88	3.60	18.00
12.36	20.88	0.00	18.78	6.32
7.07	3.60	18.78	0.00	15.13
10.44	18.00	6.32	15.13	0.00

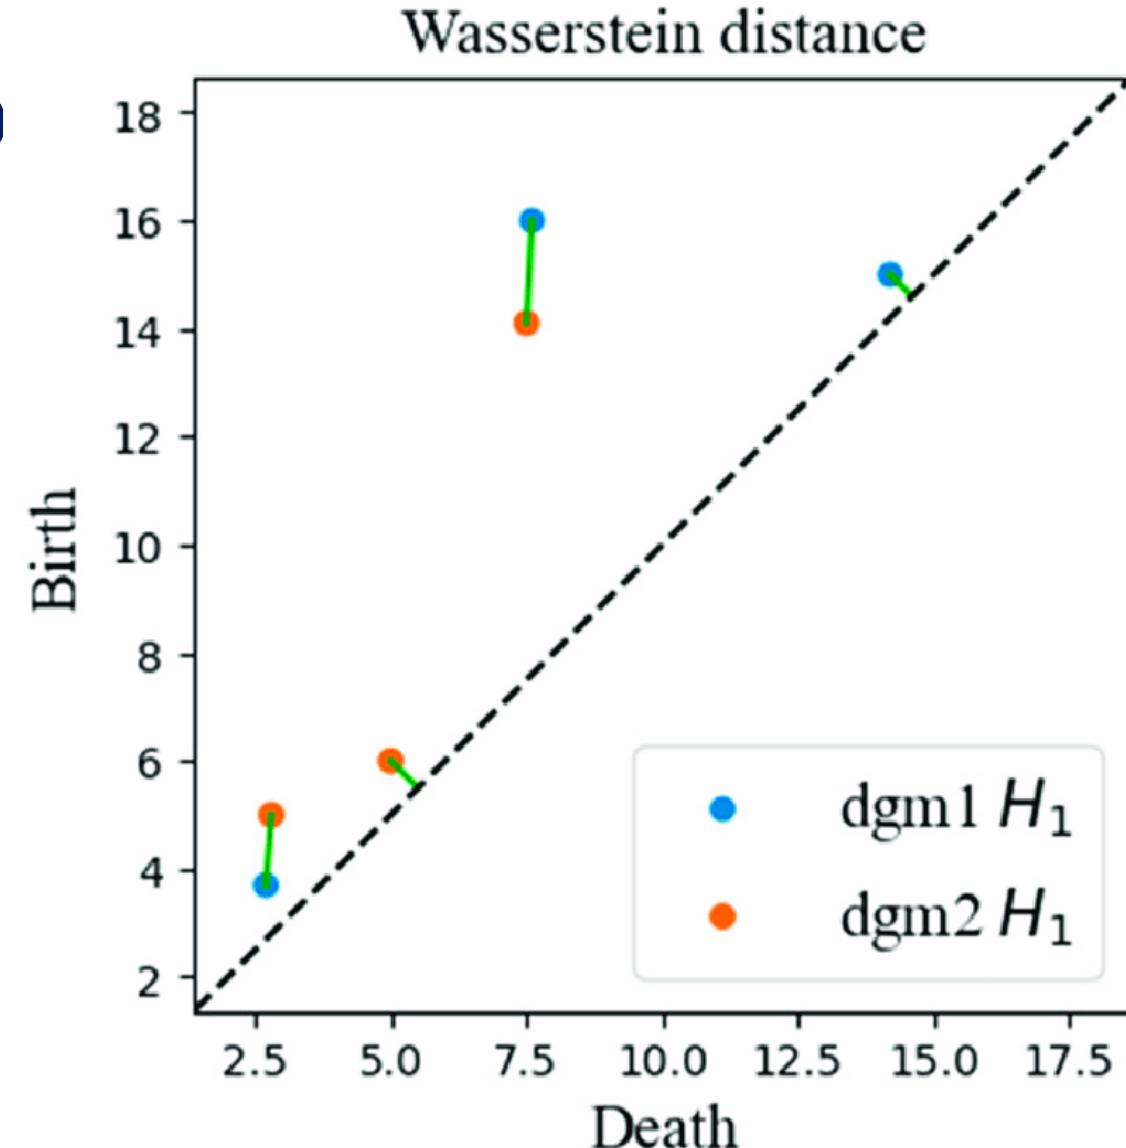
- Aktas, M.E., Akbas, E. & Fatmaoui, A.E. Persistence homology of networks: methods and applications. *Appl Netw Sci* 4, 61 (2019). <https://doi.org/10.1007/s41109-019-0179-3>

Wasserstein distance (WD)

Wasserstein distance (Earth Mover's Distance) measures distance between two probability distributions

Minimum amount of **work** required to transform u into v , **work = distribution weight \times distance**

Minimum value achieved by a **perfect matching** between the features of two Persistent Diagrams



- S. Vallender, "Calculation of the wasserstein distance between probability distributions on the line," Theory of Probability & Its Applications, vol. 18, no. 4, pp. 784–786, 1974.
- H. Edelsbrunner, "Persistent homology: theory and practice," 2013.
- Qingzhang, Shi. (2021). Time Series Clustering with Topological and Geometric Mixed Distance. Mathematics. 9. 10.3390/math9091046.

Experimental data and network construction

1
1
3

Structural T1-weighted and rs-fMRI data were obtained from the publicly available Enhanced Nathan Kline Institute Rockland Sample database (NKI-RS)

The MRI data from 3T Siemens Magnetom Tim Trio scanner

The rs-fMRI data were acquired from each subject using 3 acquisition protocols with 3 temporal sampling rates (**645ms, 1400ms, 2500ms**)

Total number of subjects: 316 with mean time series from **113 brain regions**

One weighted network per fMRI scan

- K. B. Noonan, S. Colcombe, R. Tobe, M. Mennes, M. Benedict, A. Moreno, L. Panek, S. Brown, S. Zavitz, Q. Li et al., "The nki-rockland sample: a model for accelerating the pace of discovery science in psychiatry," *Frontiers in neuroscience*, vol. 6, p. 152, 2012.

Embed FCNs in metric space

1
1
4

To map the association between two nodes u and v in the brain network, their Pearson's correlation coefficients $\text{corr}(u, v)$ are used

The objective is to transform this association into a distance metric where **stronger correlations** between nodes correspond to **shorter distances**

Mapping to all **316** subjects across **3** temporal frequencies:

$$d(u, v) = 1 - \text{corr}(u, v)$$

Included both positive and negative correlations

Generates symmetric adjacency matrix with size **113 x 113** using Pearson's correlation

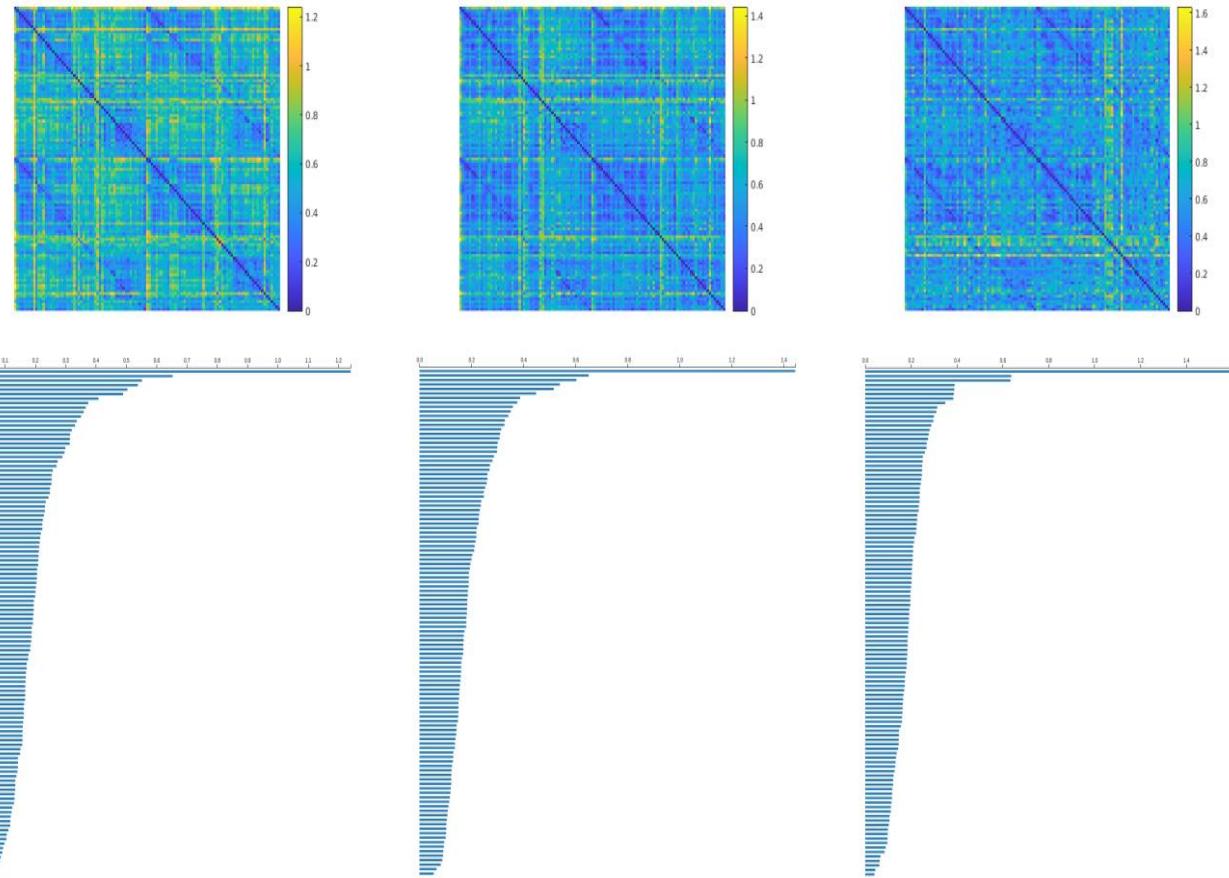
Total adjacency matrices: **948 (316 x 3)**

- B. Cassidy, C. Rae, and V. Solo, "Brain activity: Conditional dissimilarity and persistent homology," in 2015 IEEE 12th International Symposium on Biomedical Imaging (ISBI), 2015, pp. 1356–1359.
- S. M. Smith, K. L. Miller, G. Salimi-Khorshidi, M. Webster, C. F. Beckmann, T. E. Nichols, J. D. Ramsey, and M. W. Woolrich, "Network modelling methods for fmri," Neuroimage, vol. 54, no. 2, pp. 875–891, 2011.

Apply PH to metric space

Apply PH on the metric space to extract topological features

Compute 0-dimensional barcodes for all FCNs
(3×316)



FCN (top-row) and extracted topological feature (bottom-row) for subject 31 for 645ms (left), 1400ms (center), and 2500ms (right)

- The GUDHI Project, GUDHI User and Reference Manual. GUDHI Editorial Board, 2015. [Online]. Available: <http://gudhi.gforge.inria.fr/doc/latest/>
- M. Kerber, D. Morozov, and A. Nigmetov, "Geometry helps to compare persistence diagrams," Journal of Experimental Algorithmics (JEA), vol. 22, pp. 1–20, 2017.

Statistical analysis on PDs

1
1
6

Convert extracted barcodes to persistent diagram (PD)

Quantify the difference between two PDs using Wasserstein distance

Wasserstein metrics provides the foundation of the statistical inference framework

Performed 3 kinds of statistical analysis:

- Across cohort study
- Within cohort study
- Comparision with traditional FCN analysis

- S. Vallender, "Calculation of the wasserstein distance between probability distributions on the line," Theory of Probability & Its Applications, vol. 18, no. 4, pp. 784–786, 1974.
- H. Edelsbrunner, "Persistent homology: theory and practice," 2013.

Across Cohort Study

Perform comparison of PDs across cohorts using Wasserstein distance

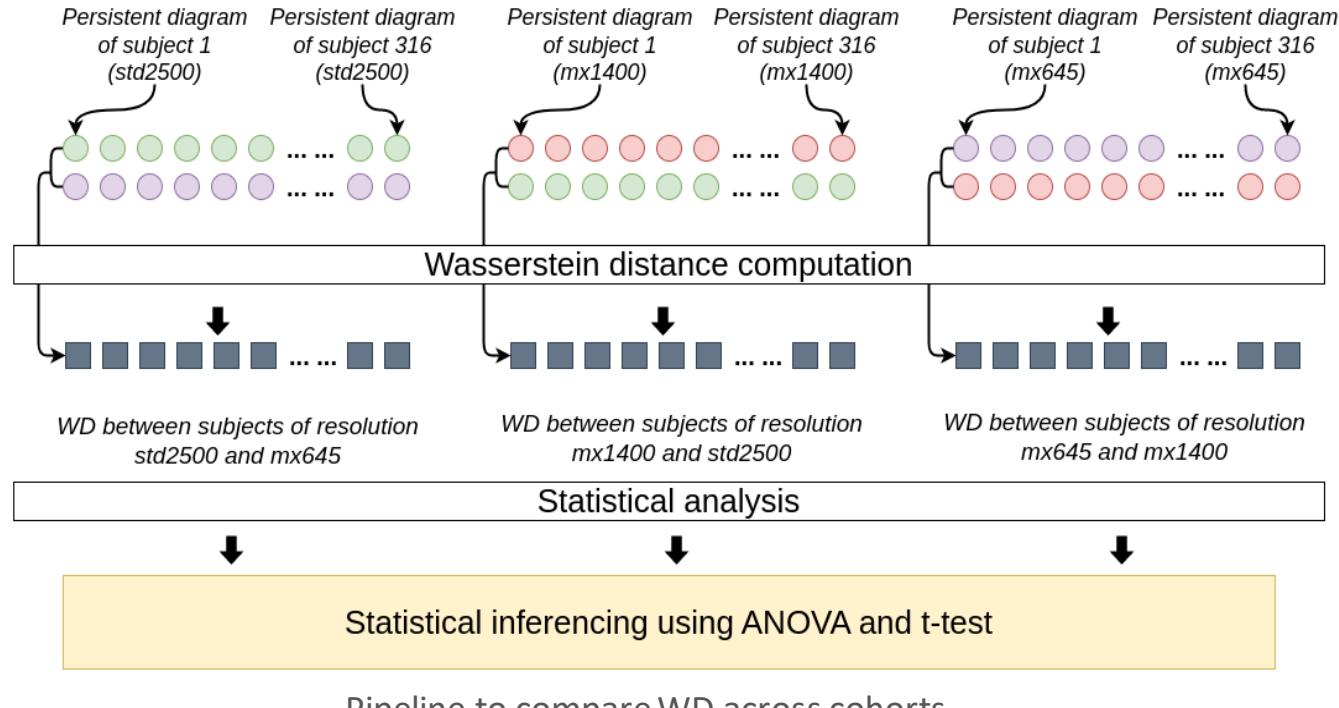
For each subject compute 3 WD

- 2500ms and 645ms
- 1400ms and 2500ms
- 645ms and 1400ms

Form 3 distributions each with 316 datapoints

Apply Analysis of Variance (ANOVA) and t-test on the distributions

High p-value (>0.05) will indicate the null hypothesis is true



• R. T. Shinohara, J. Oh, G. Nair, P. A. Calabresi, C. Davatzikos, J. Doshi, R. G. Henry, G. Kim, K. A. Linn, N. Papinutto et al., "Volumetric analysis from a harmonized multisite brain MRI study of a single subject with multiple sclerosis," American Journal of Neuroradiology, vol. 38, no. 8, pp. 1501–1509, 2017.

Within Cohort Study

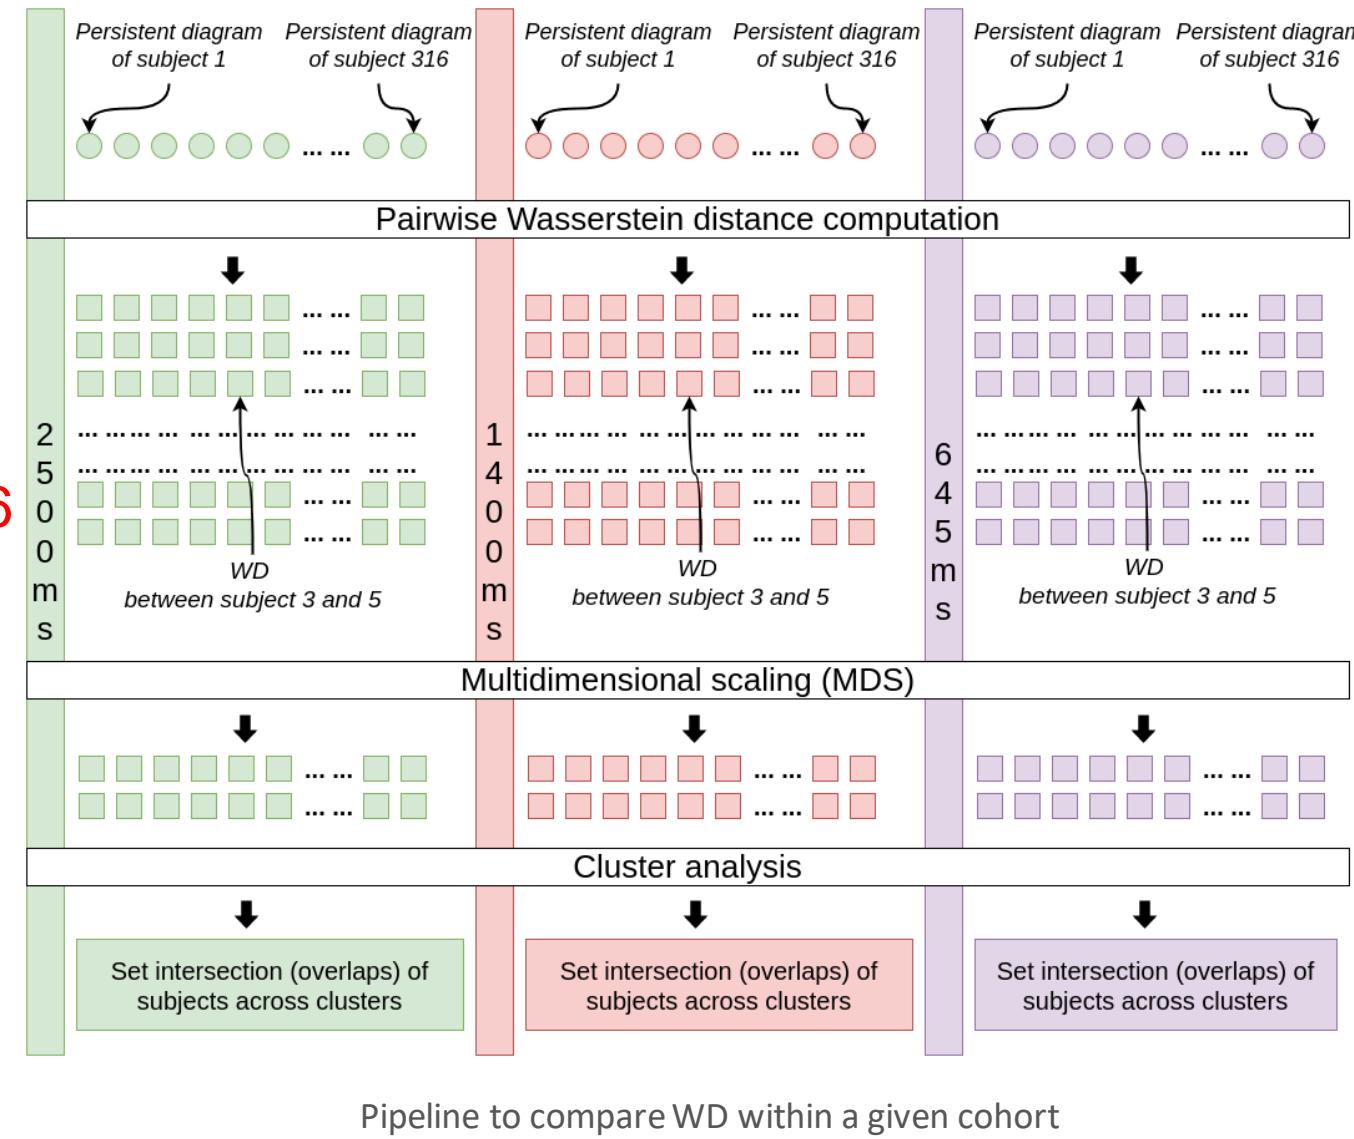
Perform comparison of pairwise WD between PDs of subjects within each cohort

Generate 3 adjacency matrix of size **316 x 316**

Reduce the dimension using 2-components MDS (matrix size **316 x 2**)

Apply K-Means clustering with Silhouette coefficient

Computes set intersection of the subjects across clusters



- Chung, M. K., Huang, S. G., Carroll, I. C., Calhoun, V. D., & Goldsmith, H. H. (2022). Dynamic Topological Data Analysis for Brain Networks via Wasserstein Graph Clustering. arXiv preprint arXiv:2201.00087.

Comparison with traditional FCN analysis

1
1
9

Use raw FCN instead of using PH features

Perform ANOVA test across all subjects for each node of the raw FCN

Compute 1 p-value per entry of the FCN

As FCN size is 113×113 , generates **113 x 113 p-values**

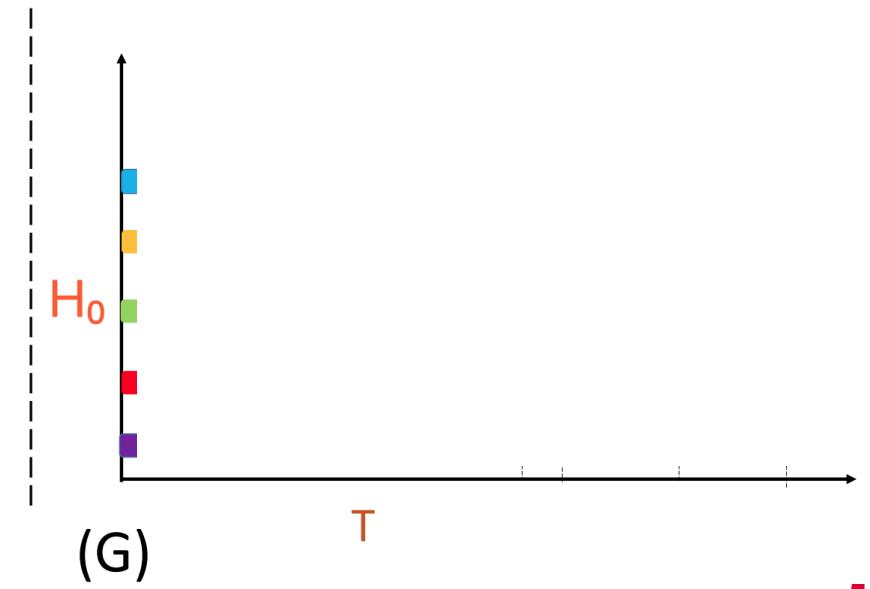
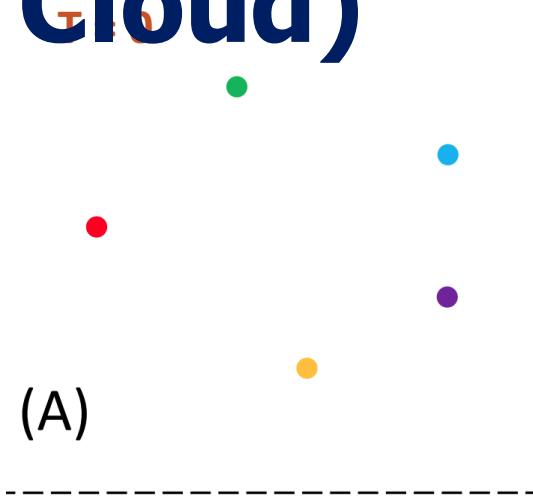
Threshold the p-values with significance level of $\alpha = 0.05$

- Assign 1 if p-value is less than the threshold, 0 otherwise

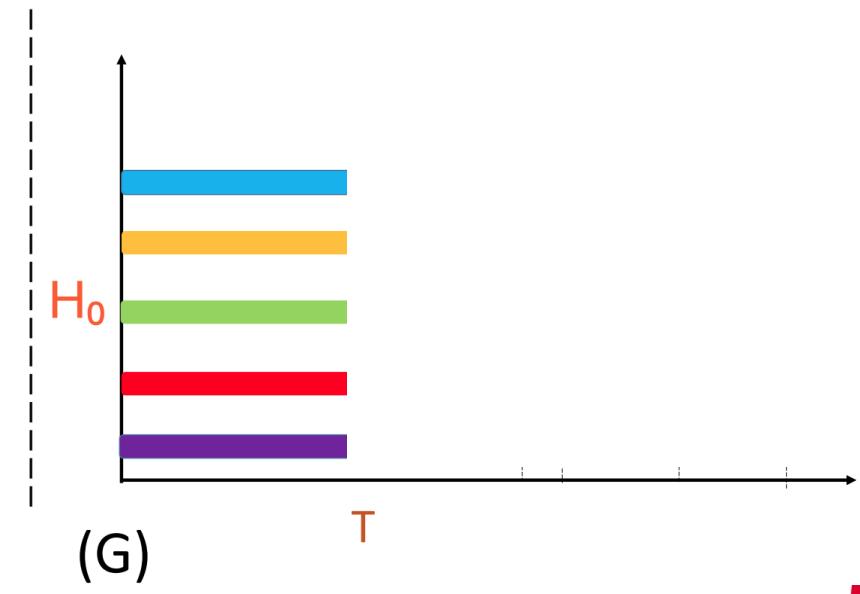
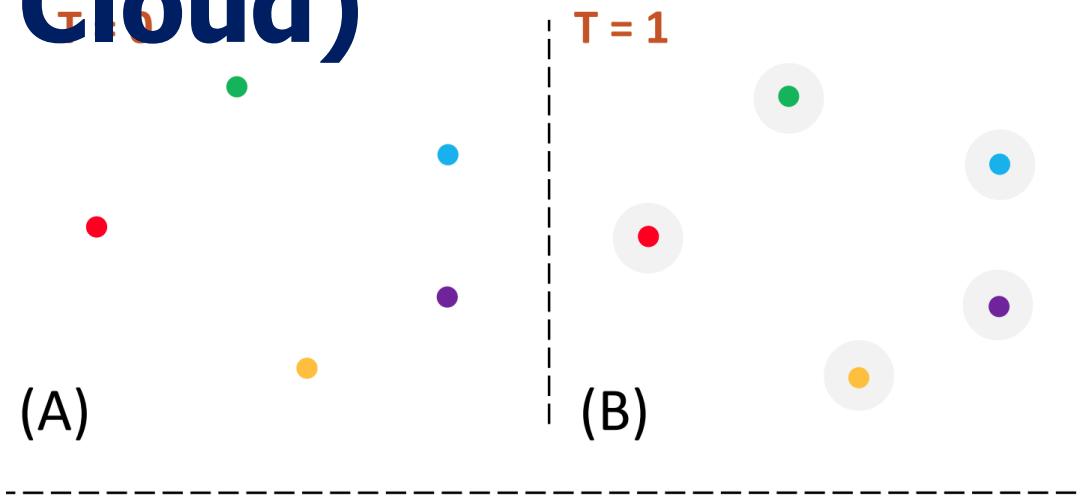
If majority of the value are 1, then it will indicate that FCN acquired from different sampling periods are different

• Feng, W., Liu, G., Zeng, K., Zeng, M., & Liu, Y. (2022). A review of methods for classification and recognition of ASD using fMRI data. Journal of neuroscience methods, 368, 109456.

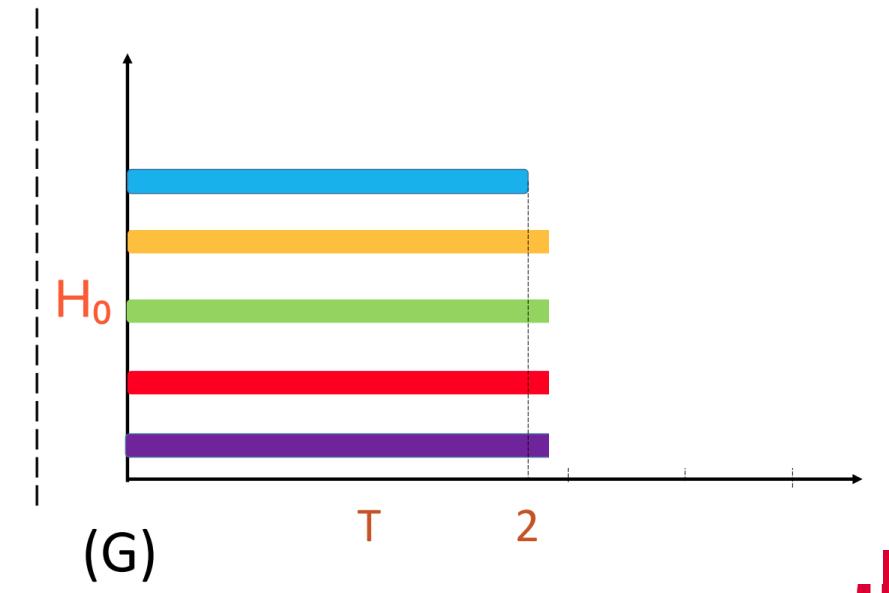
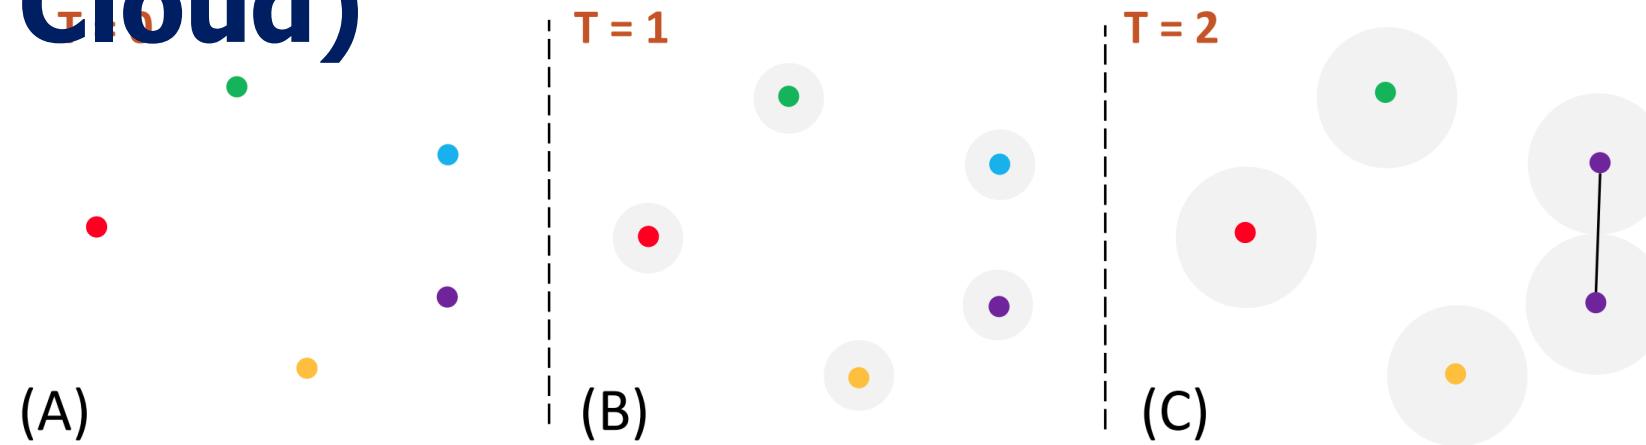
Computing Persistent Homology (Point Cloud)



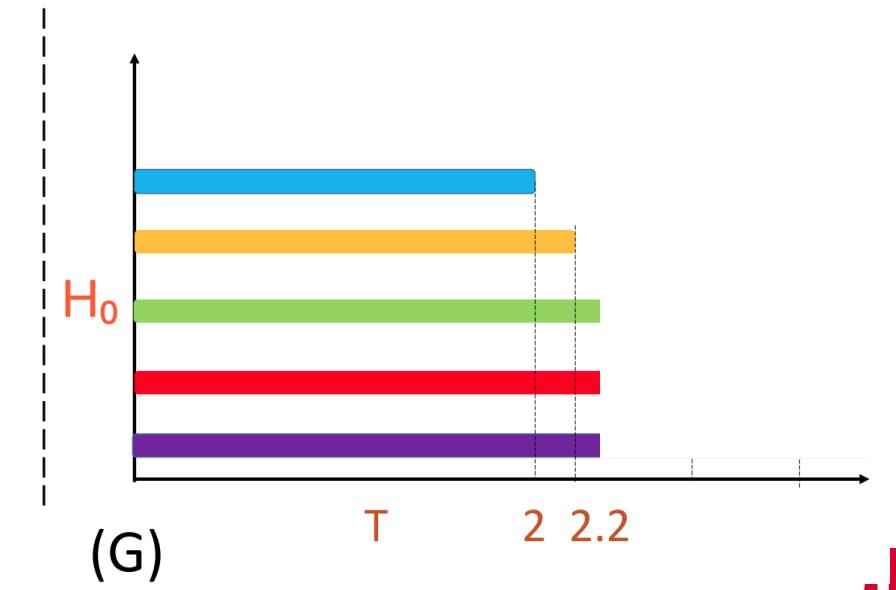
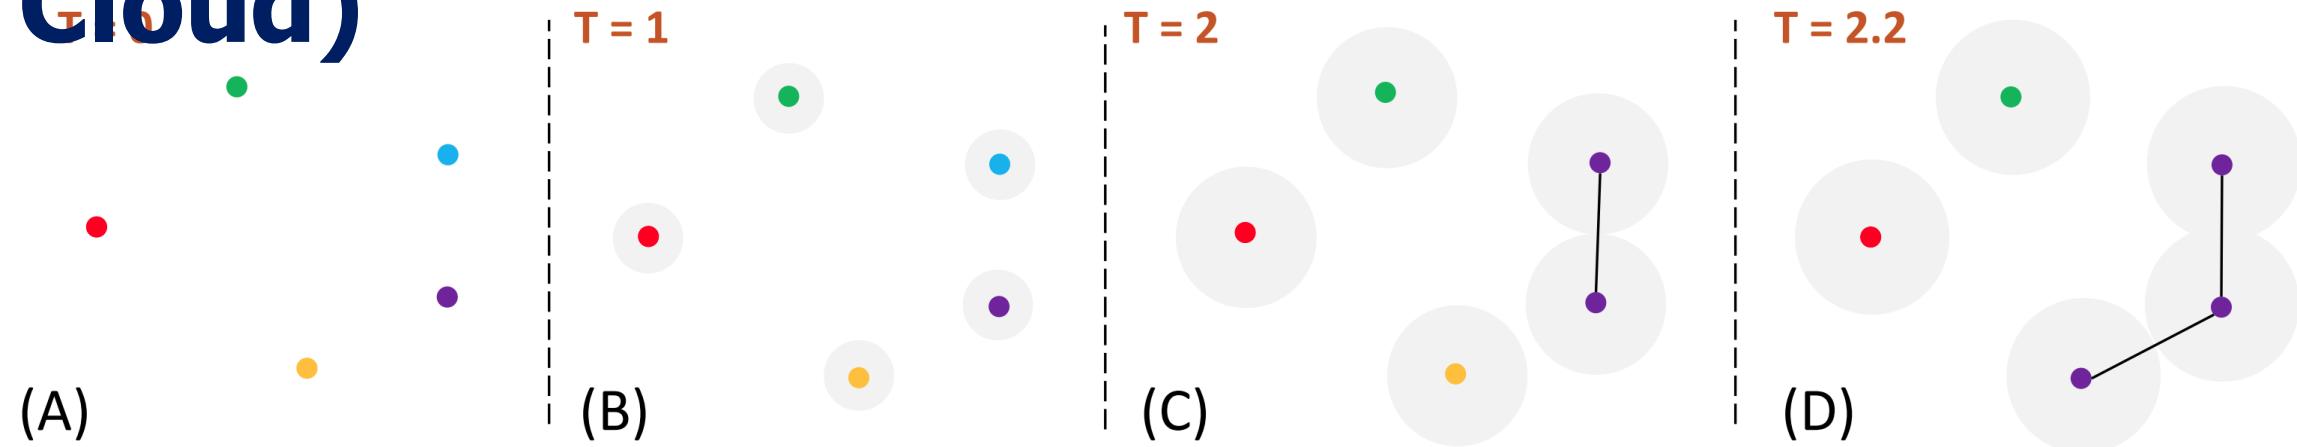
Computing Persistent Homology (Point Cloud)



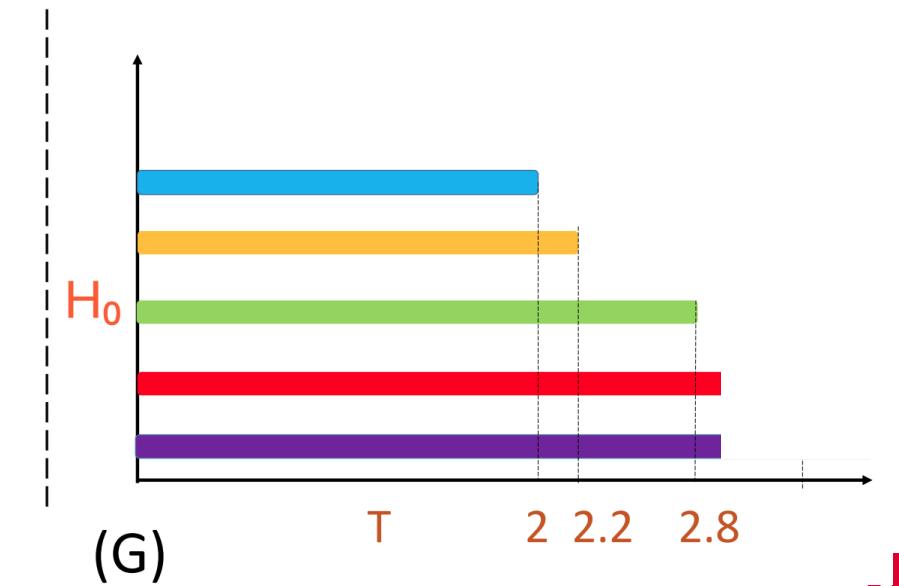
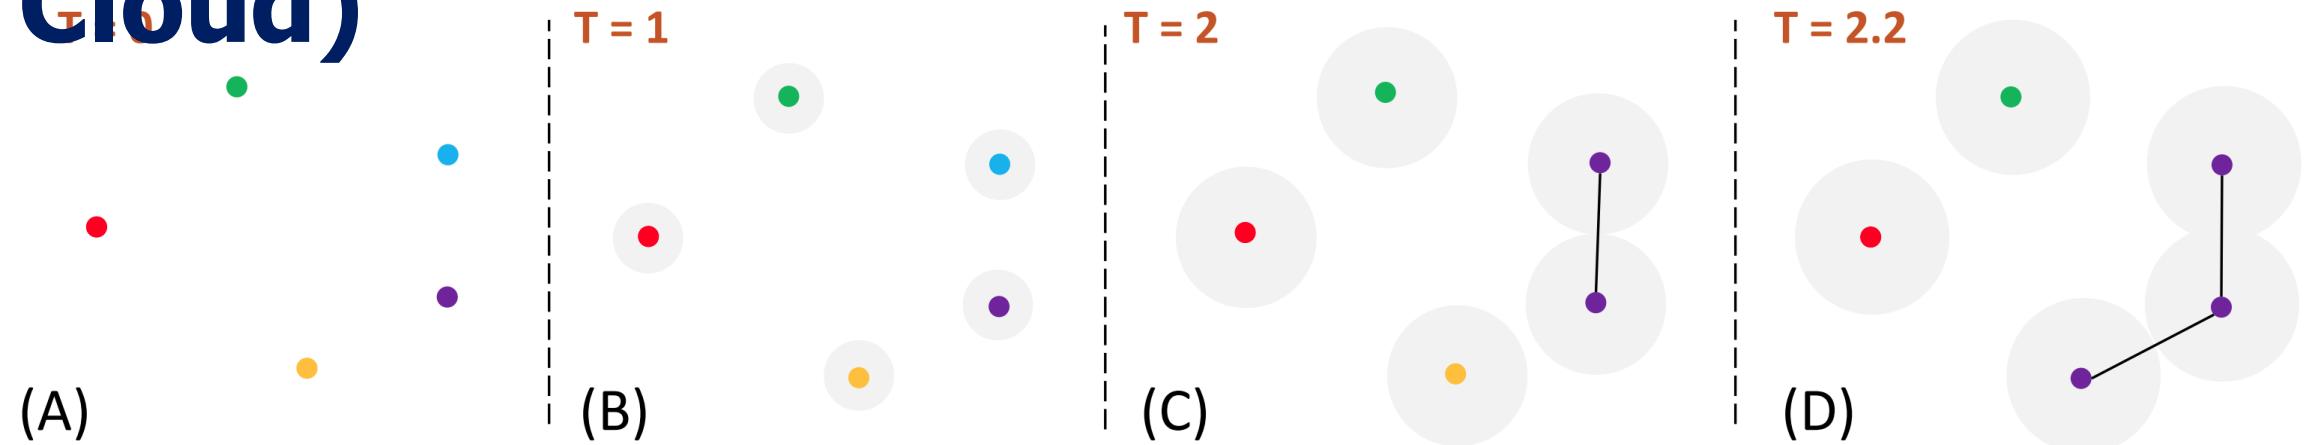
Computing Persistent Homology (Point Cloud)



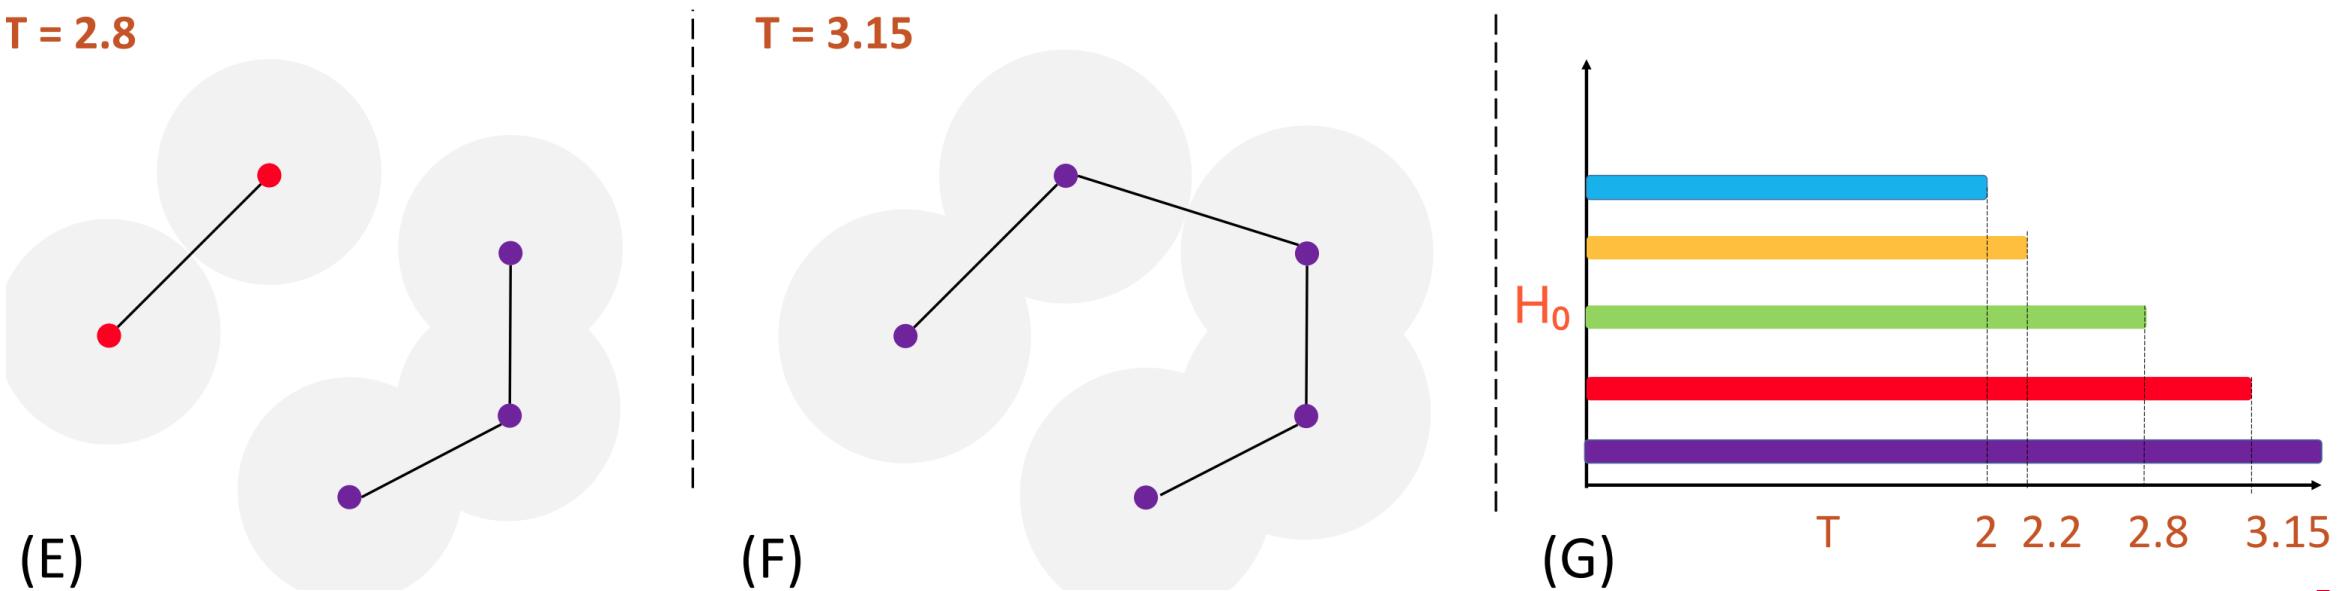
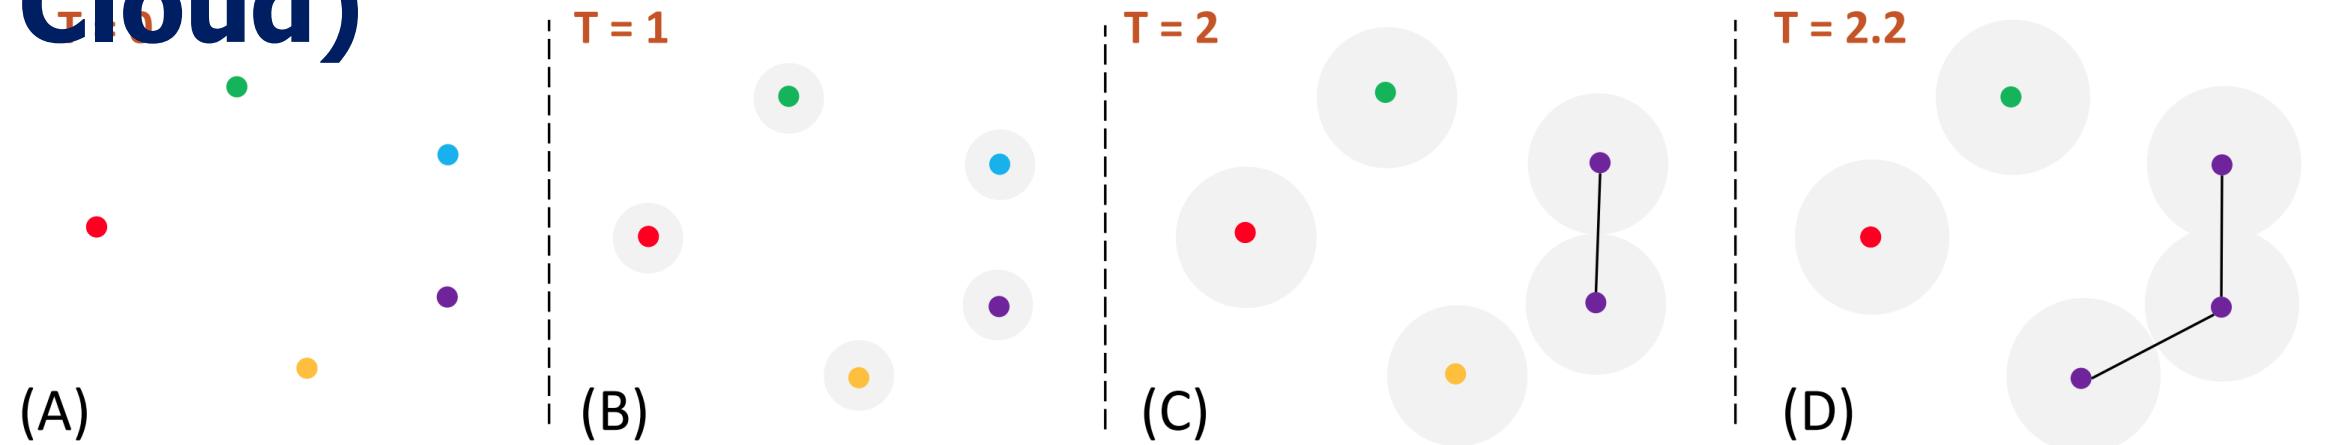
Computing Persistent Homology (Point Cloud)



Computing Persistent Homology (Point Cloud)



Computing Persistent Homology (Point Cloud)



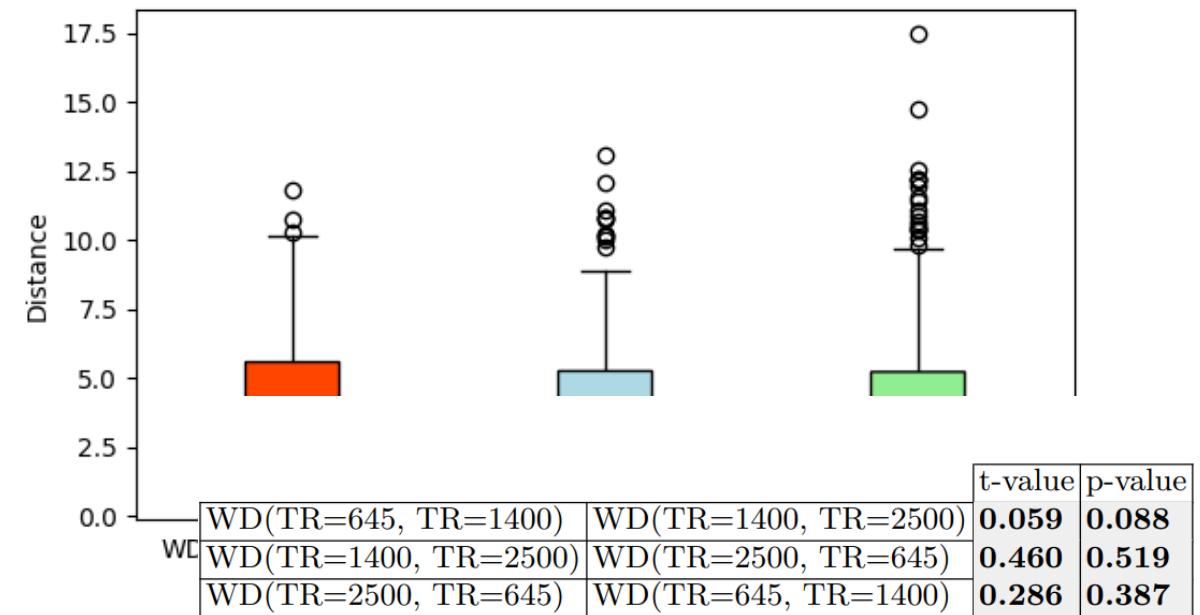
Across Cohort Study Results

1
2
6

Boxplots are aligned indicating similarity of the three distributions

ANOVA test yields p-value of 0.29

T-values and p-values are greater than the significance level of $\alpha = 0.05$



Boxplot and t-test results for across study, significance level of $\alpha = 0.05$

Functional Connectivity Networks Analysis

FCNs use Pearson's correlation to measure time series association

Multi-site rs-fMRI studies recruit faster, larger samples

Sensitive to data acquisition parameters (sampling rates)

Non-neural variability negates advantages of larger sample sizes

- B. B. Biswal, M. Mennes, X.-N. Zuo, S. Gohel, C. Kelly, S. M. Smith, C. F. Beckmann, J. S. Adelstein, R. L. Buckner, S. Colcombe et al., “Toward discovery science of human brain function,” *Proceedings of the national academy of sciences*, vol. 107, no. 10, pp. 4734–4739, 2010.
- C. Dansereau, Y. Benhajali, C. Risterucci, E. M. Pich, P. Orban, D. Arnold, and P. Bellec, “Statistical power and prediction accuracy in multisite resting-state fmri connectivity,” *Neuroimage*, vol. 149, pp. 220–232, 2017.

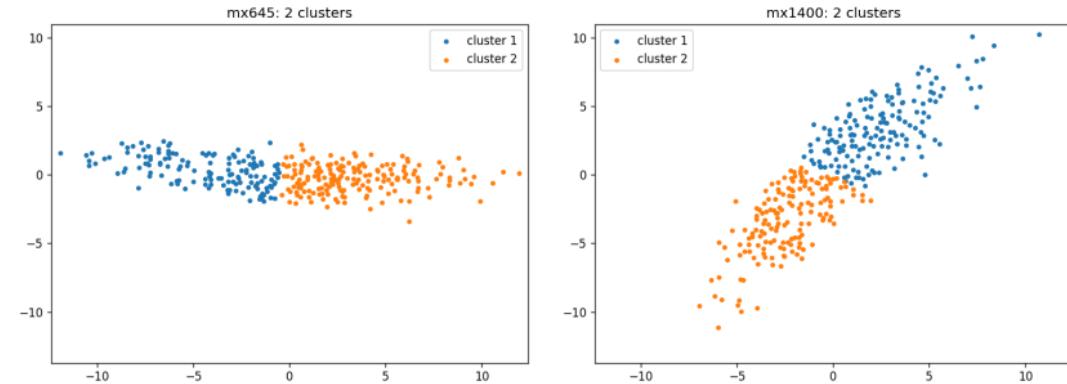
Within Cohort Study Results

Apply K-Means clustering with Silhouette score on 2-components MDS, 3 matrices (316 X 2)

198 out of 316 subjects reside in same cluster group in all 3 TRs

251 out of 316 subjects reside in same pairwise cluster intersection

Compared against random FCNs which gave p-value of 0



(a) Clustering result on MDS data for TR=645ms. (b) Clustering result on MDS data for TR=1400ms.



(c) Clustering result on MDS data for MDS plots of three data cohorts

Parallel Join (Continue)

Design

Consider partition,
load balancing,
communication

Implement

Challenging due to
the uncertain
output size

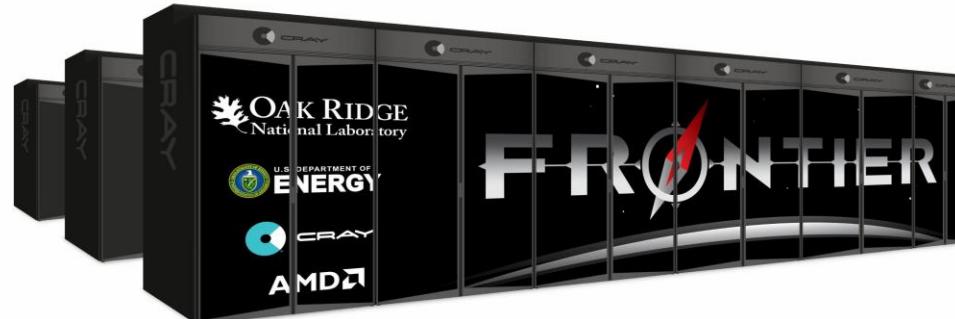
Optimize

Efficient joins
requires sorting or
indexing

Exascale Supercomputers



Aurora (ANL)



Frontier (ORNL)



El Capitan (LLNL)

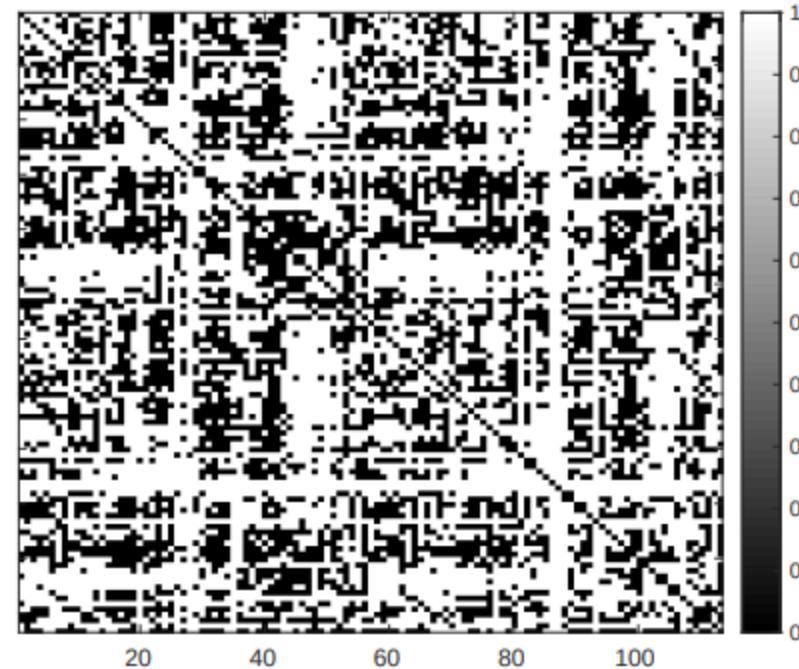
- Evans, T. M., Siegel, A., Draeger, E. W., Deslippe, J., Francois, M. M., Germann, T. C., ... & Martin, D. F. (2022). A survey of software implementations used by application codes in the Exascale Computing Project. *The International Journal of High Performance Computing Applications*, 36(1), 5-12.
- D. Kothe, S. Lee and I. Qualters, "Exascale Computing in the United States," in *Computing in Science & Engineering*, vol. 21, no. 1, pp. 17-29, 1 Jan.-Feb. 2019, doi: 10.1109/MCSE.2018.2875366.

Comparison with traditional FCN analysis results

Majority pixels are black corresponding to a p-value < significance level of $\alpha = 0.05$

7731 connections show FCNs are different for 3 TRs while **5037** connections are not

Traditional FCN analysis fails to identify connections from same subjects with different TRs



Results after ANOVA test on raw FCNs of 316 subjects, white pixels indicate no matching of the same subject across different TRs

Table of Contents

Exascale Computing
Datalog
Declarative Analytics using Datalog Applications
Iterative Relational Algebra
Parallel Join
Our Approaches
Topological Data Analysis for High-Dimensional Data
Future Research Direction
Conclusion